

No.	Title	Page
1	Getting to know Scalability	1
2	SE ZG583, Scalable Services Lecture No. 1	2
3	Introduction to Performance, Consistency and availability	3
4	CAP Theorem	4
5	CAP Theorem	5
6	Eventual Vs Strong Consistency	6
7	Performance	7
8	Availability of a system	8
9	What is scalability?	9
10	Need for scalable architectures	10
11	Monolithic Architecture	11
12	Advantages of Monolith	12
13	Disadvantages of Monolith	13
14	Principles of Scalability	14
15	Guidelines for Building Highly Scalable Systems	15
16	Architecture's scalability requirements	16
17	Challenges for Scalability	17
18	Case Study	18
19	YouTube case Study	19
20	Scalability in YouTube's Architecture	20
21	Video Processing and Transcoding Challenges	21
22	Data Management Challenges	22
23	Real-Time Recommendations Challenges	23
24	Live Streaming Scalability Challenges	24
25	Fault Tolerance and Resilience Challenges	25
26	Self Study	26
27	References	27
28	Popular scaling approaches	28
29	SE ZG583, Scalable Services Lecture No. 2	29
30	Partitioning and Sharding	30
31	Introduction	31
32	Types of Partitioning	32
33	Horizontal partitioning (Sharding)	33

34	Vertical partitioning	34
35	Functional partitioning	35
36	Scalability using NoSQL and HDFS	36
37	NoSQL and Scalability	37
38	Use Cases for NoSQL Scalability	38
39	HDFS and Scalability	39
40	Use Cases for HDFS Scalability	40
41	Managing high velocity data streams	41
42	Key Challenges	42
43	Strategies for Managing High- Velocity Video Data	43
44	1. Content Delivery Network	44
45	CDN Case Study: YouTube	45
46	2. Adaptive Bitrate Streaming (ABR)	46
47	3. Efficient Video Compression	47
48	4. Edge Computing	48
49	Case Study: Smart Security Camera	49
50	5. Real-Time Data Processing	50
51	Real Time Analytics	51
52	Case Study: Real Time Fraud Detection	52
53	Web Conferencing	53
54	Steps in Processing Streaming Data for Web Conferencing	54
55	Real-World Example: Zoom	55
56	Managing high velocity Data Streams using Kafka	56
57	BITS Pilani, Pilani Campus	57
58	Kafka for Stream Processing and Analytics	58
59	References	59
60	Popular scaling approaches continued	60
61	SE ZG583, Scalable Services Lecture No. 3	61
62	Agenda	62
63	Managing high volume transactions	63
64	Introduction to replication	64
65	Service Replicas	65
66	Stateful and Stateless Replicas	66
67	Active-Active and Active- Passive Replicas	67
68	Leader-Follower Replica	68

69	Load Balancing	69
70	What are load balancers?	70
71	Why is Load Balancing Important for Scalability?	71
72	Command Query Responsibility Segregation (CQRS)	72
73	Why CQRS for Scalability?	73
74	How CQRS works?	74
75	Some challenges of implementing CQRS	75
76	What is Event sourcing?	76
77	How Event Sourcing and CQRS Work Together	77
78	Event Sourcing and CQRS	78
79	Architecture of CQRS with Event Sourcing	79
80	C. Query Side (Read Model) D. EventHandlers & Projections	80
81	Protocols for communication	81
82	Asynchronous Communication	82
83	Synchronous Vs Asynchronous	83
84	Message Broker	84
85	Benefits of Message Broker	85
86	Drawbacks of Message Broker	86
87	Example	87
88	What is Caching?	88
89	Types of cache	89
90	Distributed Caches	90
91	Advantages of Distributed Cache	91
92	Example: Memcached	92
93	Global Caches	93
94	Google Global Cache (GGC)	94
95	Scalability features in the Cloud	95
96	Auto-scaling	96
97	Types of Auto Scaling	97
98	Comparison	98
99	What is virtualization?	99
100	Types of virtualization	100
101	How Virtualization Improves Scalability?	101
102	What is Serverless Computing?	102
103	Key Characteristics of Serverless Computing	103
104	How Serverless Computing Enhances Scalability?	104

105	Types of Serverless Architectures for Scalability	105
106	Challenges of Serverless Computing for Scalability	106
107	Best Practices for Achieving Scalability	107
108	Case Study: Amazon Prime	108
109	Self Study	109
110	References	110
111	Microservices - Introduction	111
112	SE ZG583, Scalable Services Lecture No. 4	112
113	What is Monolithic Architecture?	113
114	Example Architecture	114
115	Advantages of Monolithic	115
116	Disadvantages of Monolith	116
117	Need for Microservices	117
118	What is Microservices?	118
119	Main characteristics of microservices	119
120	Example Architecture	120
121	SOA	121
122	SOA Vs Microservices	122
123	SOA Vs Microservices	123
124	SOA Vs Microservices	124
125	SOA vs Microservices: Which is best for you?	125
126	Case Study	126
127	FTGO Case Study	127
128	Old Architecture of the FTGO application	128
129	Problems faced in old FTGO architecture	129
130	Monolithic Hell	130
131	Possible solution	131
132	X- axis scaling	132
133	Z-axis scaling	133
134	Y-axis scaling	134
135	Microservices Architecture for FTGO	135
136	Netflix Case Study	136
137	How Netflix worked earlier?	137
138	Challenges in previous architecture	138
139	What they need in new Architecture?	139
140	Updated Netflix architecture	140

141	Netflix Approach	141
142	Latest Netflix Architecture	142
143	Netflix follows continuous learning	143
144	Self Study	144
145	References	145
146	Microservices Contd.	146
147	SE ZG583, Scalable Services Lecture No. 5	147
148	Advantages of Microservices	148
149	Advantages contd.	149
150	Advantages contd.	150
151	Microservices is not a silver bullet	151
152	Amazon Prime Video's Transition	152
153	Process and Organization	153
154	Introduction	154
155	Software development and delivery organization	155
156	Software development and delivery process	156
157	The human side of adopting microservices	157
158	Microservices Design Principles	158
159	Single Responsibility Principle (SRP)	159
160	Abstraction and Information Hiding	160
161	Loose Coupling & High Cohesion	161
162	Resilience & Fault Tolerance	162
163	Observability & Monitoring	163
164	Statelessness	164
165	Important concepts	165
166	What is a Service?	166
167	Size of a Service	167
168	Services: The role of Shared Libraries	168
169	Steps for defining an application's microservice architecture	169
170	BITS Pilani, Pilani Campus	170
171	Step1: Identify system operations	171
172	FTGO domain model	172
173	Defining system operations	173
174	Step 2 and 3: Defining services and service API	174
175	Decomposition based patterns to define services	175
176	Decompose by business capability pattern	176

177	Identifying Business Capabilities	177
178	From business capabilities to services	178
179	Decompose by sub-domain pattern	179
180	From Subdomains to Services	180
181	Decompose by sub-domain pattern	181
182	Decomposition guidelines	182
183	Self Sudy	183
184	References	184
185	Microservices Contd.	185
186	SE ZG583, Scalable Services Lecture No. 6	186
187	Obstacles to decomposing an application into services	187
188	Defining service APIs	188
189	Introduction	189
190	Key Principles of Service APIs	190
191	Assigning system operations to services	191
192	Example	192
193	Determining the APIs required to support collaboration between services	193
194	Example	194
195	Transition from monolith to microservices	195
196	Split the Database First	196
197	Split the Database First	197
198	Split the Code First	198
199	Split the Code First	199
200	Patterns for Monolith to Microservices	200
201	Rebuild From Scratch	201
202	Strangler Pattern	202
203	Strangler Pattern	203
204	Strangler Pattern: Issues	204
205	When not to use Strangler Pattern?	205
206	Communication Protocols	206
207	Aspects of communication	207
208	Synchronous communication	208
209	Representational state transfer (REST)	209
210	Guiding Principles of REST	210
211	gRPC: Introduction	211
212	gRPC	212

213	Why use gRPC in Microservices?	213
214	2. Strongly Typed Contracts	214
215	3. Supports Streaming	215
216	4. Language-Agnostic	216
217	5. Built-in Authentication & Security	217
218	Asynchronous communication	218
219	Asynchronous Communication Example	219
220	Communication Styles	220
221	Push and real-time communication based on HTTP	221
222	Single-receiver message- based communication	222
223	Multiple-receivers message- based communication	223
224	Asynchronous event-driven communication	224
225	Choosing the Right Communication Pattern for Microservices	225
226	REST API (Synchronous, Request-Response)	226
227	gRPC (High-Performance Remote Procedure Calls)	227
228	Kafka (Event-Driven, Asynchronous Messaging)	228
229	RabbitMQ (Message Queuing, Reliable Delivery)	229
230	WebSockets (Persistent, Real-Time Bi-Directional Communication)	230
231	Comparison	231
232	References	232
233	Communication and Transaction management	233
234	SE ZG583, Scalable Services Lecture No. 7	234
235	Introduction	235
236	Direct Communication	236
237	Using an API Gateway	237
238	Using an API Gateway: Request Routing	238
239	Using an API Gateway: Protocol Translation	239
240	API Gateway Architecture	240
241	Benefits of an API Gateway	241
242	Drawbacks of an API Gateway	242
243	What is API design?	243
244	Designing APIs	244
245	How do we design our API program	245
246	Good API design	246
247	External API design issues	247
248	What Is The API Contract?	248

249	When To Create An API Version	249
250	What is API security?	250
251	Why is API security important?	251
252	REST API security	252
253	Most common API security best practices	253
254	Backends for Frontends	254
255	Benefits of BFF pattern	255
256	Database related patterns for Microservices	256
257	Shared Database pattern	257
258	Database View Pattern	258
259	Database Wrapping Service Pattern	259
260	Database-as-a-Service Pattern	260
261	Splitting Apart the Database	261
262	Transaction management in a microservice architecture	262
263	Need for distributed transactions in a MSA	263
264	Challenges with distributed transactions	264
265	Challenges with distributed transactions	265
266	Solution: Saga pattern	266
267	Example	267
268	Challenges in Saga	268
269	Compensating transactions	269
270	Create Order Saga	270
271	Self Study	271
272	References	272
273	Microservices Review	273
274	SE ZG583, Scalable Services Lecture No. 8	274
275	Review	275
276	When to Use Hadoop	276
277	Kafka Vs Rabbit MQ	277
278	Homework: Kafka	278
279	BITS Pilani, Pilani Campus	279
280	When to use CDN	280
281	When to use Load balancer	281
282	Edge computing scenarios	282
283	When to use CQRS pattern	283

284	When serverless architecture is not the right choice	284
285	Building with pipelines	285
286	SE ZG583, Scalable Services Lecture No. 9	286
287	Agenda	287
288	DevOps	288
289	Key Principles of DevOps	289
290	Continuous Integration	290
291	CI Workflow	291
292	Benefits of Continuous Integration	292
293	Continuous Integration for Microservices	293
294	Mapping Continuous Integration to Microservices	294
295	Mapping Continuous Integration to Microservices continued...	295
296	Mapping Continuous Integration to Microservices continued...	296
297	Key Considerations for CI in Microservices	297
298	Continuous Delivery in Microservices	298
299	CD Pipeline for Microservices	299
300	Best Practices for Continuous Delivery in Microservices	300
301	Continuous Deployment in Microservices	301
302	Deployment Strategies in Microservices	302
303	Pipeline	303
304	CI/CD pipeline for Monolith Vs Microservices	304
305	Platform-Specific Artifacts	305
306	Service Configuration	306
307	Repository Strategies for Microservices	307
308	Mono-repo vs. Multi-repo	308
309	Deployment Challenges	309
310	Pipeline for FTGO application	310
311	References	311
312	Designing Reliable Microservices	312
313	SE ZG583, Scalable Services Lecture No. 10	313
314	Agenda	314
315	What is reliability?	315
316	Example	316
317	Sources of failure	317
318	Hardware Failure	318
319	Communication Failure	319

320	Dependency-related failure	320
321	Service Practices	321
322	Cascading failures	322
323	Example: cascading failure	323
324	BITS Pilani, Pilani Campus	324
325	BITS Pilani, Pilani Campus	325
326	Designing reliable communication	326
327	Retries	327
328	Retries	328
329	Retries: few points to remember	329
330	Asynchronous communication	330
331	Circuit breakers	331
332	BITS Pilani, Pilani Campus	332
333	Why Use Circuit Breakers?	333
334	Maximizing service reliability	334
335	Load Balancer	335
336	Rate limits	336
337	Rate limiting in Microservices	337
338	Example Use Cases	338
339	Throttling Pattern	339
340	How to Implement Throttling?	340
341	Queue-Based Load Levelling pattern	341
342	How the Pattern Works?	342
343	Service mesh	343
344	Service Mesh	344
345	How does service mesh work?	345
346	Service Mesh and Microservices	346
347	Service Mesh and Microservices	347
348	Service Mesh and Microservices continued...	348
349	Does service mesh optimize communication?	349
350	Self study	350
351	References	351
352	Securing and Testing scalable services	352
353	SE ZG583, Scalable Services Lecture No. 11	353
354	Agenda	354

355	Securing code and repository	355
356	Measures for securing the code and repository	356
357	What is Authentication?	357
358	Authentication implementation options	358
359	Authentication in API Gateway	359
360	Authentication in API Gateway continued...	360
361	Authentication in API Gateway continued...	361
362	Authorization	362
363	Authorization in API Gateway	363
364	Key aspects and benefits of implementing authorization in an API gateway	364
365	Authorization in Services	365
366	Key aspects and benefits of implementing authorization within services	366
367	Which approach is better?	367
368	What is OAuth?	368
369	Terms used in OAuth	369
370	How OAuth works for Microservices	370
371	Testing in Microservices	371
372	Use DevOps	372
373	Autonomous service teams	373
374	Testing Pyramid and Microservices	374
375	Testing strategies for Microservice architectures	375
376	Unit Tests	376
377	Types of Unit Test	377
378	Component Testing	378
379	Integration Test	379
380	Approaches for integration testing	380
381	Consumer Driven Contact testing	381
382	Contract testing	382
383	Contract testing continued...	383
384	End-to-End Testing	384
385	End to End Testing continued...	385
386	Fault injection testing	386
387	Chaos Engineering At Netflix	387
388	Tools for Unit testing	388
389	Tools for Integration Testing	389
390	Tools for End-to-End Testing	390

391	Self Study	391
392	References	392
393	Deploying Microservices	393
394	SE ZG583, Scalable Services Lecture No. 12	394
395	Agenda	395
396	Challenges in deploying Microservices	396
397	Features of a microservice production environment	397
398	Service Startup	398
399	Service Startup	399
400	Run multiple instances of a service	400
401	Adding a load balancer	401
402	AWS: What is a Classic Load Balancer?	402
403	Models for Deploying Services	403
404	Multiple Service Instances per Host Pattern	404
405	Single Service Instance per Host Pattern	405
406	Service Instance per Container Pattern	406
407	Best Deployment Strategy	407
408	Running a single service on local machine	408
409	Dockerfile	409
410	Dockerfile best practices	410
411	Docker Compose	411
412	Run multiple services on a container	412
413	Run a single service using Docker desktop	413
414	Self Study	414
415	References	415
416	Deploying Microservices Contd.	416
417	SE ZG583, Scalable Services Lecture No. 13	417
418	Agenda	418
419	Introduction of deployment Strategies	419
420	Ramped	420
421	Blue/Green	421
422	Canary	422
423	A/B testing	423
424	Serverless Deployment	424
425	Drawbacks of using Serverless Deployment	425
426	Examples of Serverless Deployment	426

427	Introduction to containers	427
428	Container Image	428
429	Docker Containers	429
430	Key aspects of Docker containers	430
431	What can I use Docker for?	431
432	Docker architecture	432
433	Containerizing a service	433
434	Build and Run Docker Image	434
435	Working with images	435
436	Building your image	436
437	Running containers	437
438	Storing an image	438
439	Deploying to a cluster	439
440	Update the application	440
441	Share the application	441
442	Play with Docker	442
443	Containers as a Service (CaaS)	443
444	CaaS workflow	444
445	Docker and git	445
446	Self Study	446
447	References	447
448	Monitoring	448
449	SE ZG583, Scalable Services Lecture No. 14	449
450	Monitoring	450
451	Observability	451
452	Difference between Observability and Monitoring	452
453	Monitoring Stack	453
454	Monitoring	454
455	Monitoring - Example	455
456	Monitoring	456
457	Golden Signals (Metrics)	457
458	Golden Signals	458
459	Types of Metrics	459
460	Types of Metrics	460
461	Types of Metrics	461
462	AWS CloudWatch Monitoring Tool	462

463	Alerts	463
464	Key aspects of Alerts	464
465	Key aspects of Alerts	465
466	Alerts	466
467	Logs and Traces	467
468	Logs	468
469	Logs	469
470	Log - Centralized	470
471	Log - Distributed	471
472	Log – Distributed – Difficulties in accessing?	472
473	Types of Log	473
474	Sample Log collected from Logstash library Python	474
475	Traces [Distributed Tracing]	475
476	Traces [Distributed Tracing]	476
477	Logs and Traces	477
478	Monitoring in Kubernetes [Self Study – Exercise]	478
479	References	479
480	Thank You!	480
481	Kubernetes	481
482	SE ZG583, Scalable Services Lecture No. 15	482
483	Agenda	483
484	What is Kubernetes?	484
485	What does K8s provides?	485
486	Kubernetes design principles	486
487	Basic Kubernetes Architecture	487
488	Kubernetes Architecture	488
489	Important elements of a Kubernetes deployment	489
490	How Deployment works in Kubernetes	490
491	Kubernetes practice lab sources	491
492	Scalability in Kubernetes	492
493	Introduction	493
494	Horizontal Pod Autoscaler	494
495	Vertical Pod Autoscaler	495
496	Cluster Autoscaler	496
497	Storage in Kubernetes	497
498	Introduction	498

499	Volumes	499
500	Non-Persistent Storage	500
501	Persistent Volumes and Persistent Volume Claims	501
502	Kubernetes Security	502
503	Why Security Matters in Kubernetes	503
504	Core Security Areas in Kubernetes	504
505	Access control in Kubernetes	505
506	Security Tools & Ecosystem	506
507	Other Security measures in Kubernetes	507
508	What is CI/CD pipeline?	508
509	CI/CD with Kubernetes	509
510	Why CI/CD in Kubernetes?	510
511	Core CI/CD Components for Kubernetes	511
512	CI/CD Workflow in Kubernetes	512
513	GitOps Approach to CD	513
514	Testing in CI/CD Pipelines	514
515	Packaging & Deployment	515
516	Security & Observability in CI/CD	516
517	Infrastructure as Code in Kubernetes	517
518	Example of Continuous Deployment to Kubernetes	518
519	Kubernetes dashboard	519
520	Deploying the Dashboard UI	520
521	Self Study	521
522	References	522
523	Kubernetes Cluster	523
524	Agenda	524
525	How to create an Amazon EKS cluster	525
526	Create a Kubernetes cluster on AWS using eksctl	526
527	Pre-requisites	527
528	Pre-requisites	528
529	Pre-requisites	529
530	Pre-requisites	530
531	Pre-requisites	531
532	Create your Amazon EKS cluster and nodes	532
533	Create your Amazon EKS cluster and nodes	533
534	Create your Amazon EKS cluster and nodes	534

By- Akanksha Bharadwaj (Asst. Professor CSIS department)

535	View resources	535
536	View resources	536
537	How to deploy an application on Kubernetes (cloud)	537
538	Deploy application	538
539	Deploy application	539
540	Deploy application	540
541	Deploy application	541
542	Deploy application	542
543	Deploy application	543
544	Understanding Kubernetes Yaml file	544
545	Understanding Kubernetes Yaml file	545
546	Deploying the microservices	546
547	Deploying new version of an existing microservice	547
548	Clean up	548
549	Best Practices to use Kubernetes Cluster	549
550	Amazon EKS security best practices	550
551	Self Study	551
552	References	552
553	Cloud Computing - Notes on the CAP Theorem	553
554	Example System 1 – Direct Access	554
555	Example System 2 – Write - Through Cache	555
556	Example System 3 – Consistent Caches	556
557	How do we quantify CAP?	557
558	Storage Replication	558

By- Akanksha Bharadwaj (Asst. Professor CSIS department)



BITS Pilani
Pilani Campus

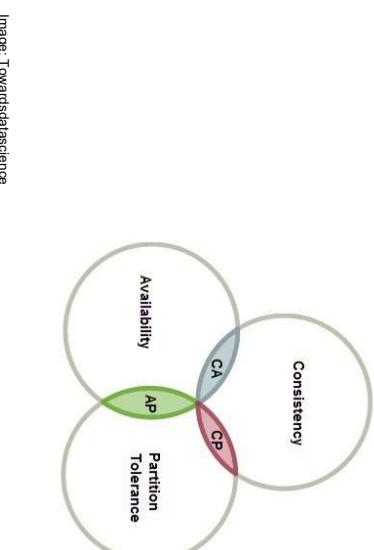
CAP Theorem

- The CAP theorem states that a distributed system can only guarantee two out of these three characteristics: Consistency, Availability, and Partition Tolerance.

SE ZG583, Scalable Services Lecture No. 1



2



BITS Pilani, Pilani Campus

4



BITS Pilani
Pilani Campus

Getting to know Scalability

Prof. Akanksha Bharadwaj
CSIS Department



Introduction to Performance, Consistency and availability

1

BITS Pilani
Pilani Campus

3

Eventual Vs Strong Consistency

- Eventual consistency is a consistency model that enables the data store to be highly available.
- Strong Consistency simply means that all the nodes across the world should contain the same value for an entity at any point in time.



Availability of a system

- Availability refers to a property of software that it is there and ready to carry out its task when you need it to be.
- Here are some of the key resources you can implement to make high availability possible:
 - Use multiple application servers
 - Spread out physically
 - Backup system
 -



CAP Theorem

Consistency
Availability
Partition tolerance

BITS Pilani, Pilani Campus

6

Performance

- The topmost reason for performance concerns is that the tasks we set our systems to perform have become much more complex over a period of time
- Can be measured using response time, throughput etc.



BITS Pilani, Pilani Campus

5

BITS Pilani, Pilani Campus

7

Need for scalable architectures



- Simplicity
- Network latency and security

10

Monolithic Architecture



BITS Pilani, Pilani Campus

12

- Scalability of an architecture refers to the fact that it can scale up to meet increased work loads.
- Types of Scalability
 - Vertical Scalability
 - Horizontal Scalability
- Monolith means composed all in one piece.
- Traditionally, applications were built on a monolithic architecture, a model in which all modules in an application are interconnected in a single, self-contained unit.
- They're typically complex applications that encompass several tightly coupled functions.
- When all functionality in a system had to be deployed together, we consider it a **monolith**.

What is scalability?



BITS Pilani, Pilani Campus

9

Principles of Scalability



- Scale horizontally, not vertically
- Statelessness
- Cache
- Load Balancing
- Modular Architecture
- Asynchronous Processing
- Monitoring and Metrics
- Resilience and Fault Tolerance

All these mainly target three areas **Availability, Performance, and Reliability**

Disadvantages of Monolith



- Scalability
- Slow development
- Long deployment cycle

Guidelines for Building Highly Scalable Systems



- Avoid shared resources as they might become a bottleneck
- Avoid slow services
- Optimize Database Design
- Auto-Scaling
- Prioritize API Design
- Ensure Data Consistency and Partitioning
- Optimize Network Usage
- Minimize Latency
- Plan for Growth

Architecture's scalability requirements



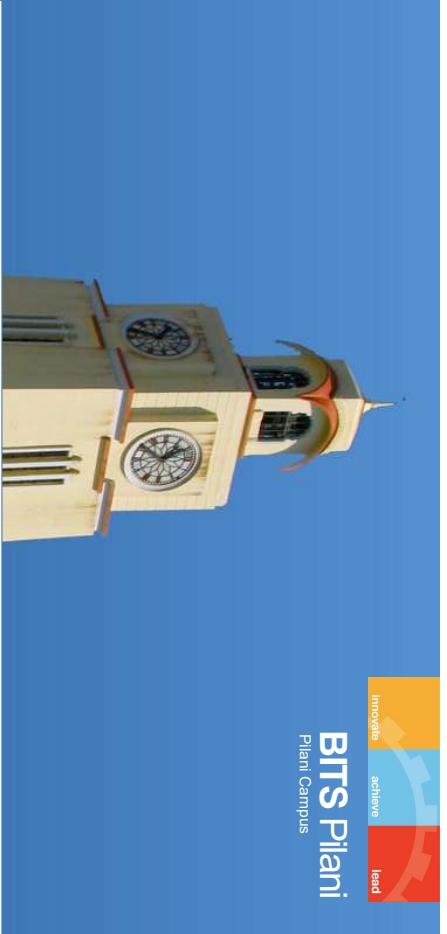
- How important are the scalability requirements?
- Identify the scalability requirements early in the software life cycle so that it allows the architectural framework to become sound enough as the development proceeds.
- System scalability criteria could include the ability to accommodate
 - Increasing number of users,
 - Increasing number of transactions per millisecond,
 - Increase in the amount of data

Scalability in YouTube's Architecture

Video Storage and Distribution

Storage Challenges:

- YouTube stores petabytes of data and must ensure quick access for playback.



Case Study

18

Challenges for Scalability



Here are some common **challenges for scalability** in real-world applications:

- Unanticipated Traffic Spikes
- Database Scalability
- Resource Contention
- Distributed System Complexities
 - Legacy Systems
 - Cost of Scalability
 - Security at Scale
 - Scalability Testing

YouTube case Study



Key Challenges in YouTube's Scalability:

- Massive User Base
- High Data Throughput
- Global Availability
- Real-Time Features.
- Storage
- Traffic Spikes

20



Data Management

Challenges:

- Storing and retrieving metadata for billions of videos efficiently.



Live Streaming Scalability

Challenges:

- Streaming live events like concerts, sports, or news to millions of concurrent viewers.

BITS Pilani, Pilani Campus

22



Video Processing and Transcoding

Challenges:

- Each uploaded video must be transcoded into multiple resolutions and formats to support devices with different network conditions and capabilities.



Real-Time Recommendations

Challenges:

- Generating personalized video recommendations for billions of users in real-time.



BITS Pilani, Pilani Campus

24



BITS Pilani, Pilani Campus

21



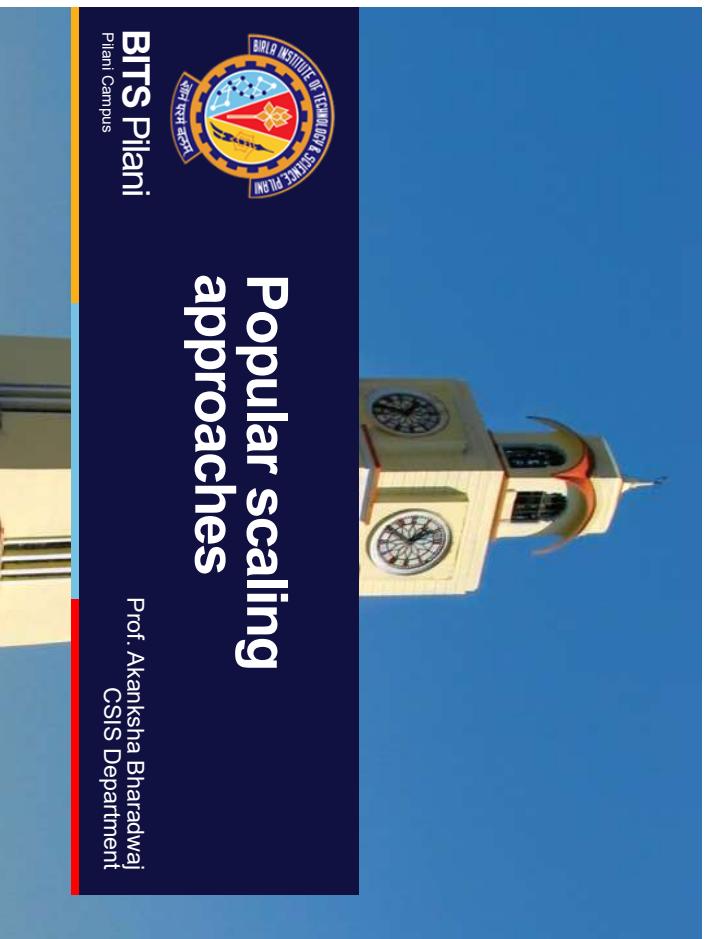
BITS Pilani, Pilani Campus

23

Self Study

Example of scalable architecture:

<https://www.youtube.com/watch?v=VHELcOe1gy0>



Popular scaling approaches

BITS Pilani
Pilani Campus

Prof. Akanksha Bharadwaj
CSIS Department



References

Fault Tolerance and Resilience

Challenges:

- Ensuring uninterrupted service even during server failures or network issues.



26

References

Fault Tolerance and Resilience

Challenges:

- Textbooks and reference books mentioned in the handout
- <https://groups.csail.mit.edu/tds/papers/Gilbert/Brewer2.pdf>
- <http://highscalability.com/youtube-architecture>



BITS Pilani, Pilani Campus

25

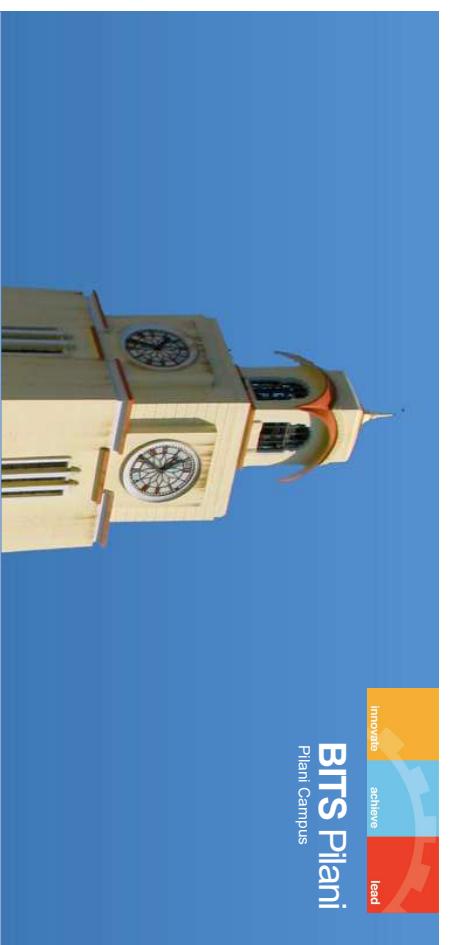
BITS Pilani, Pilani Campus

27

Types of Partitioning

- Horizontal partitioning
- Vertical partitioning
- Functional partitioning

Partitioning and Sharding



30

Introduction

- Partitioning and Sharding are scalable approaches to manage and distribute large volumes of data in a database system.

- In many large-scale solutions, data is divided into *partitions* that can be managed and accessed separately.

Why partition data?

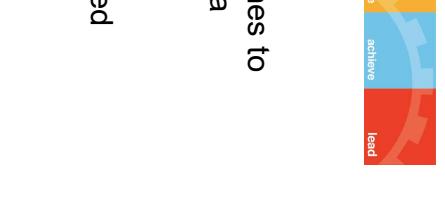
- Improve scalability
- Improve performance
- Improve security
- Provide operational flexibility
- Improve availability

SE ZG583, Scalable Services

Lecture No. 2



29



31

Vertical partitioning



Key	Name	Description	Stock	Price	LastOrdered
ARC1	Arc welder	250Amps	8	119.00	25-Nov-2013
BRK8	Brakcket	250mm	45	5.66	18-Nov-2013
BRK9	Brakcket	400mm	63	6.98	1-Jul-2013
HOS2	Hose	1/2"	27	27.50	10-Feb-2013
WGT4	Widget	Green	16	13.99	3-Feb-2013
WGT6	Widget	Purple	76	13.99	31-Mar-2013

Key	Stock	LastOrdered
ARC1	8	25-Nov-2013
BRK8	46	18-Nov-2013
BRK9	82	1-Jul-2013
HOS2	27	18-Apr-2013
WGT4	16	3-Feb-2013
WGT6	76	31-Mar-2013

BITS Pilani, Pilani Campus

34

Horizontal partitioning (Sharding)



Key	Name	Description	Stock	Date	LastOrdered
ARC1	Arc welder	250Amps	8	119.00	25-Nov-2013
BRK8	Brakcket	250mm	56	5.66	18-Nov-2013
BRK9	Brakcket	400mm	58	6.98	10-Feb-2013
HOS2	Hose	1/2"	27	27.50	3-Feb-2013
WGT4	Widget	Green	16	13.99	3-Feb-2013
WGT6	Widget	Purple	76	13.99	31-Mar-2013



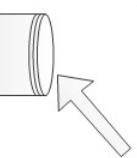
Scalability using NoSQL and HDFS



Functional partitioning



Corporate data domain					
Key	Name	Description	Price
ARC1	Arc welder	250Amps	119.00
BRK8	Brakcket	250mm	5.66
BRK9	Brakcket	400mm	6.98
HOS2	Hose	1/2"	27.50
WGT4	Widget	Green	13.99
WGT6	Widget	Purple	13.99



Use Cases for NoSQL Scalability

- Social media platforms managing large volumes of user interactions.

- E-commerce systems requiring fast, scalable databases for product catalogs and transactions.

- Real-time analytics on streaming data.



NoSQL and Scalability

NoSQL databases (e.g., MongoDB, Cassandra, DynamoDB) are specifically built for high scalability, especially in scenarios requiring distributed systems, real-time data access, and schema flexibility.

Scalability Features in NoSQL:

- Horizontal Scalability (Sharding)
- Dynamic Schema
- Replication
 - High Throughput
 - Decentralized Architecture

BITS Pilani, Pilani Campus
38

HDFS and Scalability

HDFS is a distributed file system designed for storing and processing massive datasets (big data) in a scalable and fault-tolerant manner. It forms the foundation of the Hadoop ecosystem.

Scalability Features in HDFS:

- Horizontal Scalability
- Replication for Fault Tolerance
- Write Once, Read Many (WORM) Design
 - MapReduce Integration
 - Block Storage
 - Cluster Management

BITS Pilani, Pilani Campus
40

Use Cases for HDFS Scalability

- Storing and analyzing massive datasets (e.g., petabytes of log data).

- Batch processing for data warehouses and ETL pipelines.

- Data lakes for structured and unstructured data.



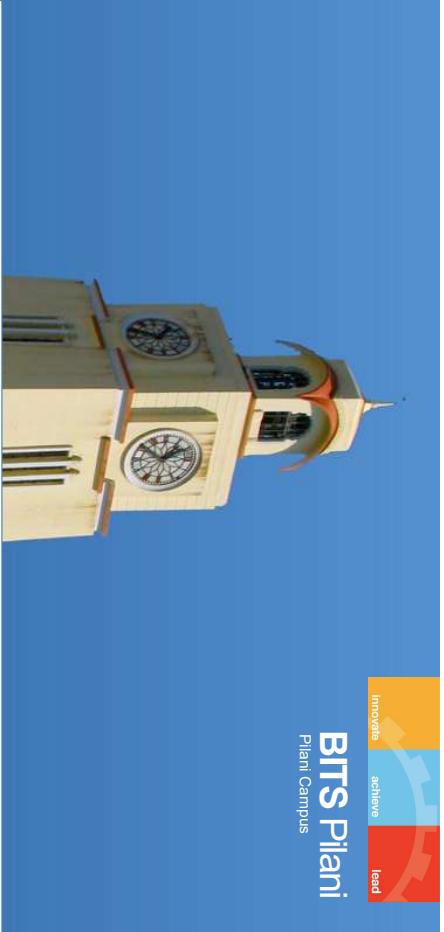
BITS Pilani, Pilani Campus
37

Key Challenges

- High Throughput
- Low Latency
- Scalability
- Fault Tolerance
- Storage Efficiency



Managing high velocity data streams



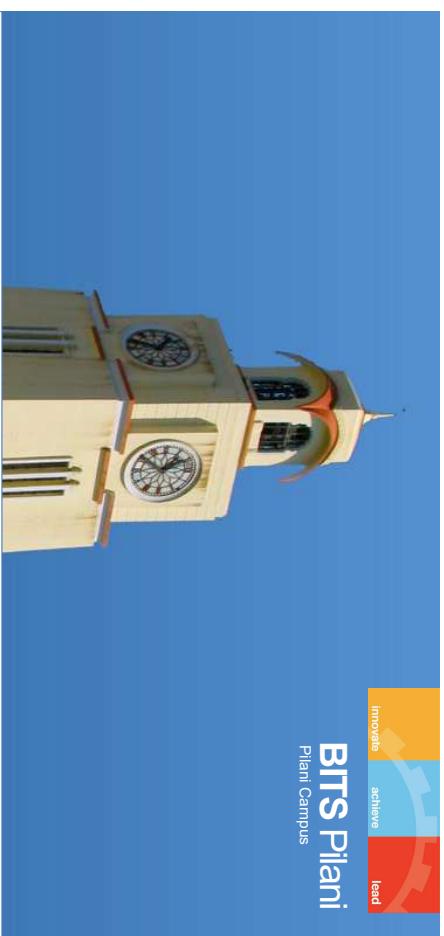
BITS Pilani, Pilani Campus

1. Content Delivery Network

- CDN is nothing more than a bunch of globally distributed computers that are directly connected and move data from one end to another.
- CDN stores multiple copies of video content at edge locations.
- We can use CDNs to cache and distribute video content closer to end-users, reducing latency and bandwidth usage.
- Examples: Akamai, Cloudflare, AWS CloudFront.



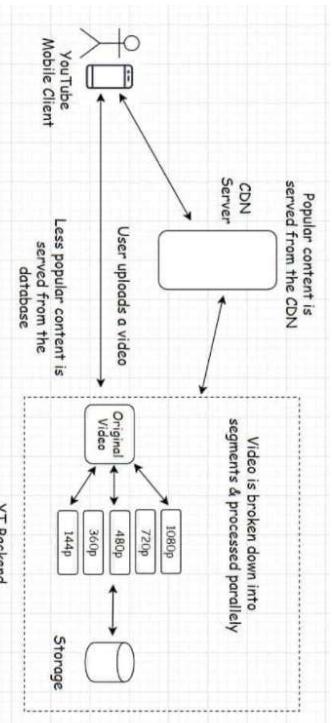
Strategies for Managing High-Velocity Video Data



BITS Pilani, Pilani Campus

2. Adaptive Bitrate Streaming (ABR)

- Deliver video streams in different quality levels (bitrates) and dynamically switch between them based on the user's network speed and device capabilities.
- Example: HLS (HTTP Live Streaming), DASH (Dynamic Adaptive Streaming over HTTP).



YouTube's Video Delivery Architecture

Image: google

4. Edge Computing

- Edge computing is a distributed information technology (IT) architecture in which client data is processed at the periphery of the network, as close to the originating source as possible.

- We can process and cache video content at edge locations to reduce latency and offload computation from the central servers.

- Example: Transcode video streams at the edge for faster delivery.

BITS Pilani, Pilani Campus

46



3. Efficient Video Compression

- Use advanced codecs to reduce video file size while maintaining quality.

- **Example for video Streaming:** To make the videos viewable on different devices Netflix performs transcoding or encoding which involves finding errors and converting the original video into different formats and resolutions.



BITS Pilani, Pilani Campus

48



BITS Pilani, Pilani Campus

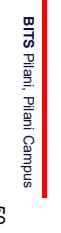
45

Image: google

5. Real-Time Data Processing



- Process telemetry data (e.g., playback statistics, user behavior) in real time to optimize streaming.
- Use stream processing frameworks like Apache Flink or Kafka Streams to analyze and react to user interactions instantly.



50

Case Study: Smart Security Camera

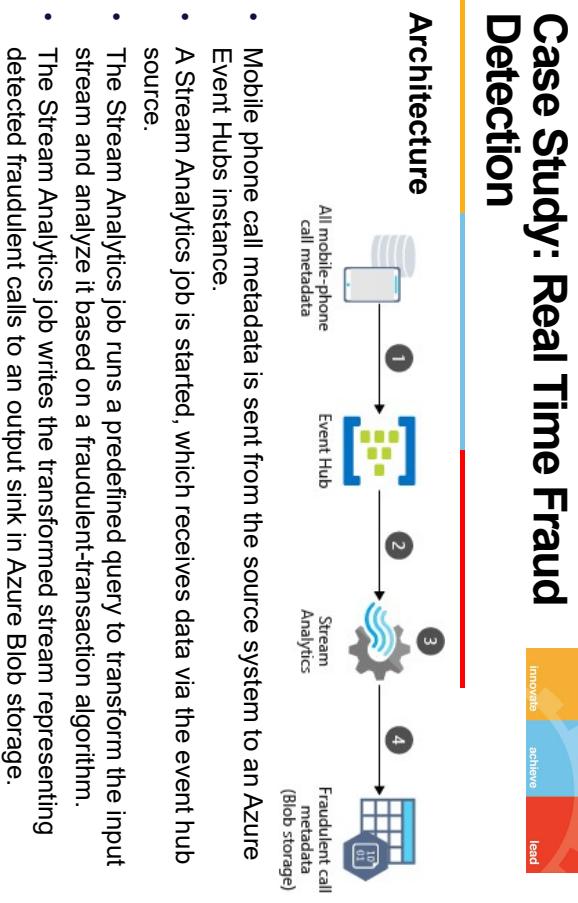
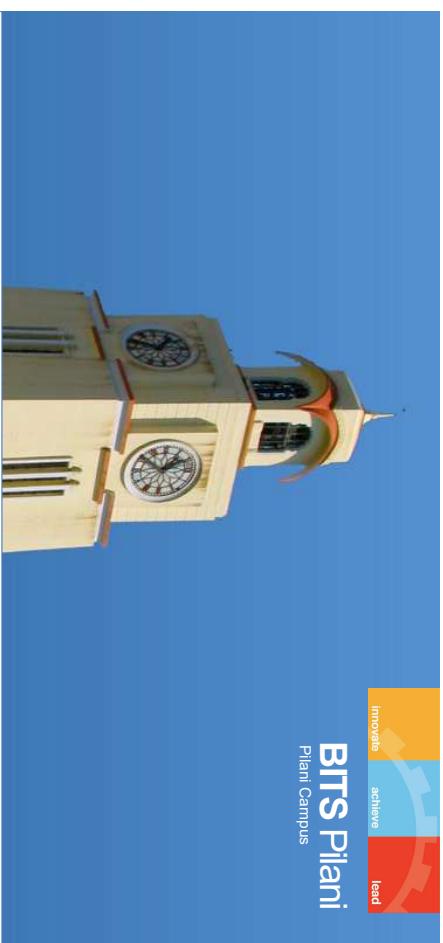


- Device with Edge Capability: A smart security camera equipped with an edge AI chip processes video feeds locally.

- **Edge Processing:** The camera runs motion detection algorithms (e.g., using TensorFlow Lite) directly on the device to identify unusual activity, like someone entering the camera's field of view.

- **Action at the Edge:** If motion is detected: Capture the event (e.g., save a short video clip). Send an alert to the homeowner's mobile device (e.g., "Motion detected at the front door"). If no motion is detected, discard the video feed to save bandwidth.
- **Cloud Involvement:** Only the processed, important data (e.g., the detected motion event) is sent to the cloud for storage or additional analysis.

Real Time Analytics



52

Case Study: Real Time Fraud Detection



Steps in Processing Streaming Data for Web Conferencing

1. Data Capture
2. Data Compression
3. Data Transmission
4. Stream Processing at the Server
5. Adaptive Bitrate Streaming
6. Content Delivery
7. Client-Side Rendering
8. Data Analytics and Feedback



Managing high velocity Data Streams using Kafka

- **Kafka** is a powerful, distributed streaming platform designed to handle massive amounts of data in real-time. It's widely used for building real-time data pipelines and applications.

BITS Pilani, Pilani Campus

54

Real-World Example: Zoom



- Web conferencing involves real-time audio, video, and data sharing, requiring efficient processing of high-velocity data streams to ensure a seamless user experience.
- **Data Capture:** Captures audio, video, and chat data from participants.
- **Compression:** Compresses streams to optimize for bandwidth.
- **Streaming Servers:** Use SFU architecture to minimize latency.
- **Scaling:** Scales dynamically using cloud servers to accommodate user load.



BITS Pilani, Pilani Campus

56

BITS Pilani, Pilani Campus

53

Kafka for Stream Processing and Analytics



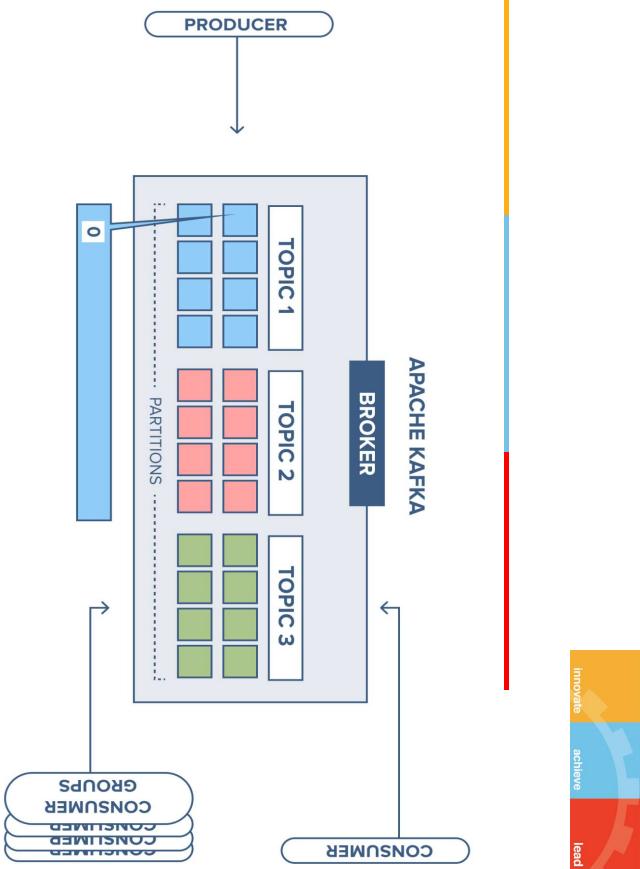
- **Use Case:** Processing and analyzing continuous streams of data for real-time insights, anomaly detection, and predictive analytics.

- **Example:** Financial institutions use Kafka to process market data feeds, detect trading anomalies, and make real-time trading decisions. Retailers analyze customer behavior and preferences in real-time to offer personalized recommendations and promotions.

Ref: <https://medium.com/@pennyy80/>

BITS Pilani, Pilani Campus

58



BITS Pilani, Pilani Campus

57

References

- <https://docs.microsoft.com/en-us/azure/architecture/best-practices/data-partitioning>
- <https://www.mongodb.com/nosql-explained>
- https://hadoop.apache.org/docs/r1.2.1/hdfs_design.html
- <http://highscalability.com/youtube-architecture>
- <https://docs.microsoft.com/en-us/azure/architecture/example-scenario/data/fraud-deletion>
- <https://blog.zoom.us/>
- <https://kafka.apache.org/11/documentation/streams/architecture>
- <https://www.gsma.com/iot/wp-content/uploads/2018/11/IoT-Edge-Opportunities-c.pdf>
- Reading material available on Confluent



BITS Pilani
Pilani Campus

Prof. Akanksha Bharadwaj
Asst. Professor, CSIS Department

Popular scaling approaches continued

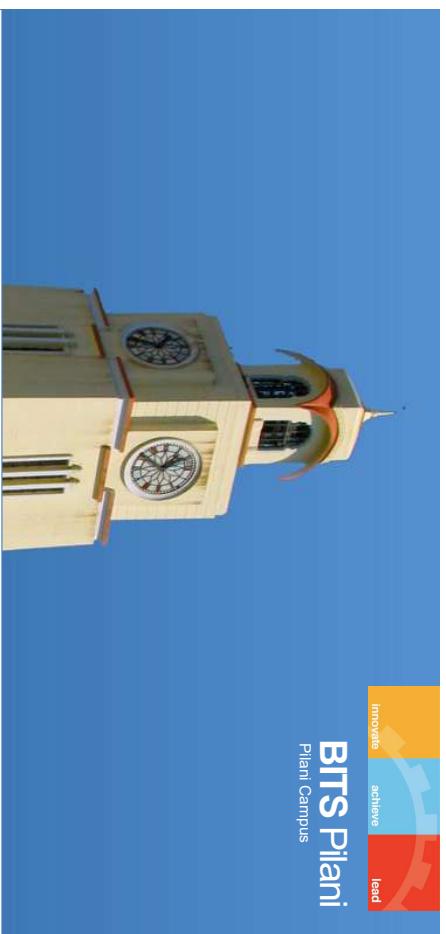


60

Agenda



SE ZG583, Scalable Services Lecture No. 3



Managing high volume transactions



SE ZG583, Scalable Services Lecture No. 3

- Popular scaling approaches
- Managing high volume transactions
 - Service Replicas & load balancing
 - Minimizing event processing: Command Query Responsibility Segregation (CQRS)
 - Asynchronous communication
 - Caching techniques: Distributed cache, global cache
- Scalability features in the Cloud (AWS, Azure, Google)
 - Auto-scaling
 - Horizontal and vertical scaling
 - Use of Load balancers
 - Virtualization
 - Serverless computing
 - Best Practices for Achieving Scalability

Case Study

BITS Pilani, Pilani Campus

62

- Replication is the practice of keeping several copies in different places.
- These replicas help distribute the workload and ensure that failures in one instance do not bring down the entire system.
- What is the need of replication?
 - High Availability
 - Fault Tolerance
 - Load sharing

BITS Pilani, Pilani Campus

64

Introduction to replication



SE ZG583, Scalable Services Lecture No. 3

Managing high volume transactions



SE ZG583, Scalable Services Lecture No. 3

Managing high volume transactions

61

Stateful and Stateless Replicas



- **Stateful Replicas:** Each replica maintains its own persistent state, requiring synchronization across instances.
- **Use Case:** Databases, caching services, distributed file systems.

- **Stateless Replicas:** Each replica operates independently and does not maintain any session-specific data.
- **Use Case:** API gateways, web servers, RESTful microservices.

Service Replicas



- It is a copy of the service logic that runs on one or more nodes of the cluster.
- Types of replicas: These replicas can be categorized based on their deployment model, scaling approach, and management strategy.
 - Stateful
 - Stateless
 - Active-Passive
 - Active-active
 - Leader-follower

Leader-Follower Replica



- A replication model where a **Leader (Primary)** node handles **write operations** and **Follower (Replica)** nodes handle **read operations**. If the Leader fails, a Follower is promoted as the new Leader.
- Use cases: Databases, Caching system, Message Brokers

Active-Active and Active-Passive Replicas



- Active-Active: All replicas handle requests simultaneously in a distributed load-sharing model.
- Use Case: High-traffic applications needing high availability.
- Active-Passive: One primary instance handles traffic, while replicas remain on standby and become active in case of failure.
- Use Case: Database replication, disaster recovery systems.

What are load balancers?



- A load balancer is a software or hardware device that keeps any one server from becoming overloaded.
- A load-balancing algorithm is the logic that a load balancer uses to distribute network traffic between servers

Types Based on Deployment Location:

- Software-based load balancers
- Hardware-based load balancers
- Cloud Load Balancer

Load Balancing



Why is Load Balancing Important for Scalability?

Load balancing ensures:

- Even distribution of requests across multiple instances.
- Reduced latency and improved response time by directing traffic to less loaded servers.
- Fault tolerance and high availability by rerouting traffic in case of server failures.
- Efficient resource utilization and better cost management.
- When a new server is added to the server group, the load balancer automatically starts to send requests to it.
- If a single server goes down, the load balancer redirects traffic to the remaining online servers.
- A load balancer uses to distribute network traffic between servers

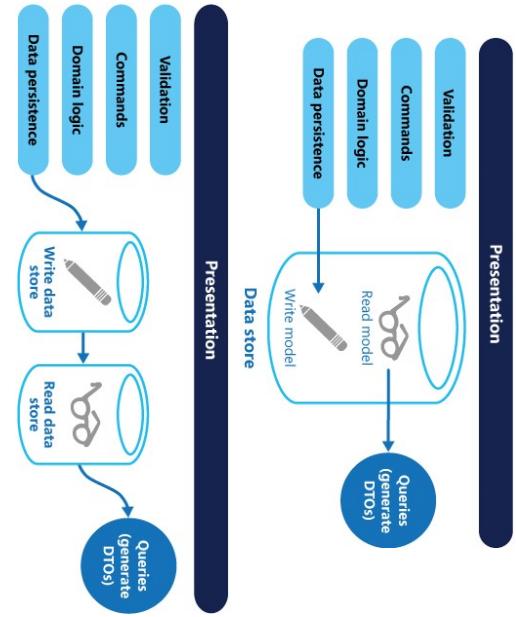
Command Query Responsibility Segregation (CQRS)



- It is a pattern that separates read and update operations for a data store.

- This pattern is particularly useful in **highly scalable and distributed applications**.

How CQRS works?



BITS Pilani, Pilani Campus

74

What is Event sourcing?



- Event Sourcing is an architectural pattern where state changes in a system are stored as a sequence of events instead of modifying the current state directly.

<https://www.confident.io/blog/event-sourcing-cqrs-stream-processing-apache-kafka-wtatis-connection/>

BITS Pilani, Pilani Campus

76

Why CQRS for Scalability?

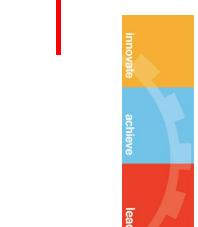


CQRS improves scalability by:

- Separating concerns
- Independent scaling
- Optimizing for reads and writes separately
- Supporting multiple databases
- Simpler queries

Some challenges of implementing CQRS

- Complexity
- Messaging
- Eventual consistency



BITS Pilani, Pilani Campus

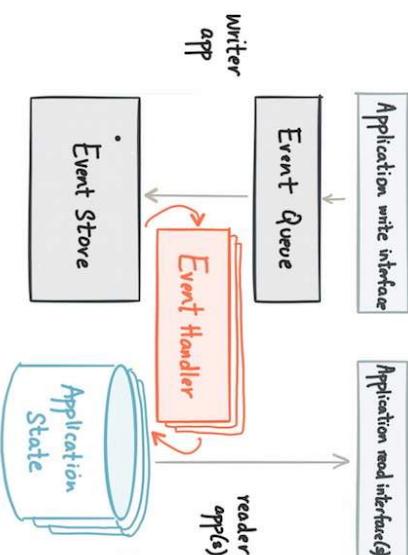
75

Event Sourcing and CQRS



C. Query Side (Read Model)

- Events from the event store are processed to update a read-optimized database.
- Read models can use SQL (denormalized tables), NoSQL (MongoDB, Redis), or search engines (Elasticsearch).



<https://mwaveza.medium.com/cqrs-pattern-with-kafka-streams-part-1-112f381eb998>

BITS Pilani, Pilani Campus

D. Event Handlers & Projections

- As events are stored, event handlers process them to update projections (optimized views for queries).
- Example: A UserRegistered event updates a UserProfile table in a read database

78
BITS Pilani, Pilani Campus

How Event Sourcing and CQRS Work Together



Architecture of CQRS with Event Sourcing



A. Command Side (Write Model)

- Write operations (Commands) → Generate immutable events instead of updating the database directly.
- Events are stored in an event store (e.g., Kafka, EventStoreDB).
- Events are read from the event store and published to event consumers.

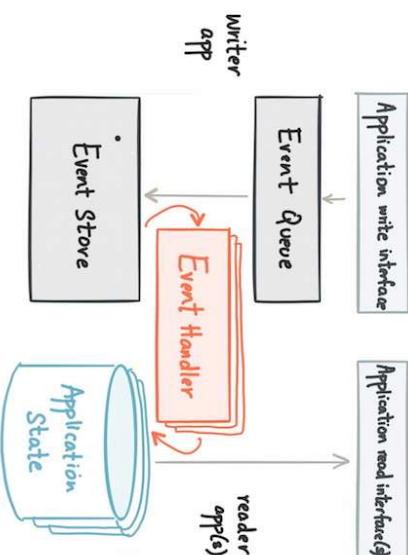
B. Event Store

- Read operations (Queries) → Use a separate, read-optimized database that is updated asynchronously.
- Example storage: Kafka, EventStoreDB, PostgreSQL (JSONB), DynamoDB Streams.

79
BITS Pilani, Pilani Campus

C. Query Side (Read Model)

- Events from the event store are processed to update a read-optimized database.
- Read models can use SQL (denormalized tables), NoSQL (MongoDB, Redis), or search engines (Elasticsearch).



<https://mwaveza.medium.com/cqrs-pattern-with-kafka-streams-part-1-112f381eb998>

BITS Pilani, Pilani Campus

D. Event Handlers & Projections

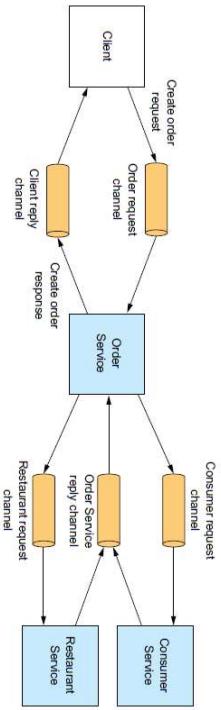
- As events are stored, event handlers process them to update projections (optimized views for queries).
- Example: A UserRegistered event updates a UserProfile table in a read database

80
BITS Pilani, Pilani Campus



Asynchronous Communication

- Services communicating by exchanging messages over messaging channels.



BITS Pilani, Pilani Campus

82

Protocols for communication



In distributed systems, communication between components plays a crucial role in **scalability, performance, and reliability**.

Two main communication patterns are:

- **Synchronous Communication:** The sender waits for a response before proceeding.
- **Asynchronous Communication:** The sender does not wait for an immediate response, improving system efficiency.

Message Broker

- It is a way of implementing asynchronous communication
- A message broker is an intermediary through which all messages flow.

Examples of popular open source message brokers include the following:

- ActiveMQ
- RabbitMQ
- Apache Kafka

BITS Pilani, Pilani Campus

83

Synchronous Vs Asynchronous



Feature	Synchronous (Blocking)	Asynchronous (Non-Blocking)
Use Case	Request-response APIs	Event-driven systems, messaging
Latency	Higher due to waiting	Lower, as operations continue
Scalability	Limited by concurrent requests	High scalability
Resilience	Prone to failures/timeouts	More resilient to failures
Complexity	Easier to implement	More complex (requires event handling)
Examples	REST APIs, RPC calls	Kafka, RabbitMQ, WebSockets

BITS Pilani, Pilani Campus

84

BITS Pilani, Pilani Campus

81

Drawbacks of Message Broker

- Potential performance bottleneck
- Potential single point of failure
- Additional operational complexity



What is Caching?

- Caching is a critical technique in **distributed systems** to reduce latency, improve performance, and optimize resource usage.
- Caching allows you to efficiently reuse previously retrieved or computed data.



Benefits of Message Broker

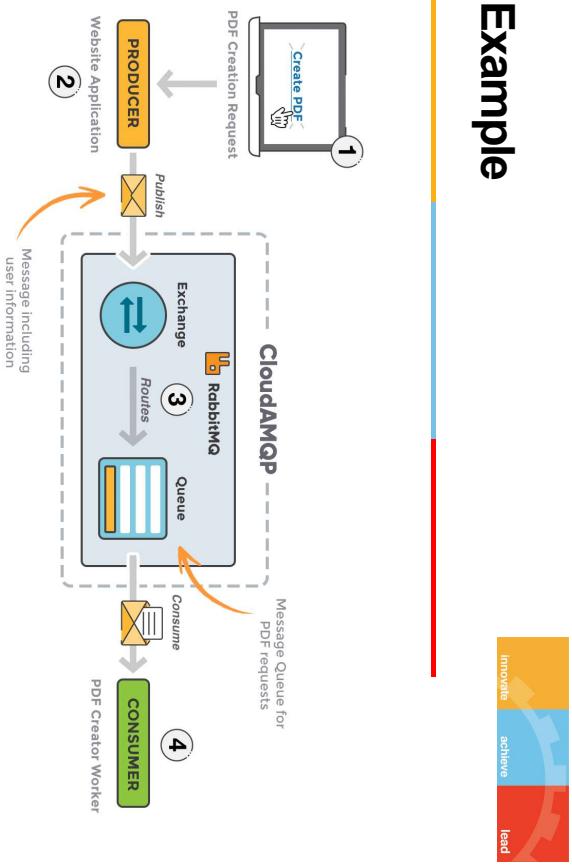
- Loose coupling
- Message buffering
- Explicit interprocess communication
- Resiliency



BITS Pilani, Pilani Campus

86

Example

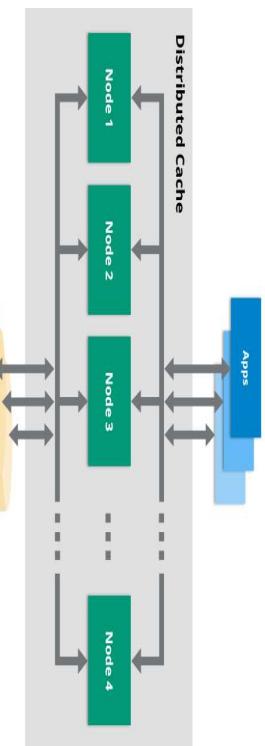


BITS Pilani, Pilani Campus

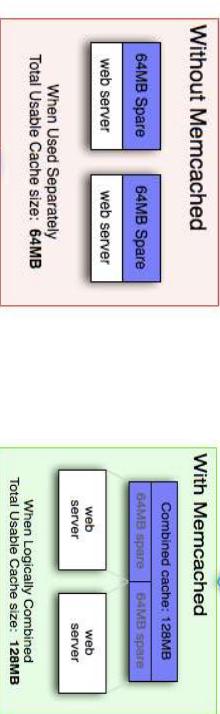
88

Distributed Caches

- A distributed cache may span multiple servers so that it can grow in size and in transactional capacity.
- It is mainly used to store application data residing in database and web session data.



Example: Memcached



Advantages of Distributed Cache

When cached data is distributed, the data:

- Is *coherent* (consistent) across requests to multiple servers.
- Survives server restarts and app deployments.
- Doesn't use local memory.

Different types of caching strategies are used depending on system architecture, data consistency needs, and workload patterns.

- Server side
 - Client side
- Application-Level Cache
- Distributed Cache



Google Global Cache (GGC)



- Google Global Cache (GGC) is a specialized caching system deployed by Google in Internet Service Providers' (ISPs) networks to improve the delivery of Google services such as:
 - YouTube
 - Google Search
 - Google Drive
 - Google Play Store
- GGC is part of Google's Content Delivery Network (CDN), designed to reduce bandwidth costs and improve access speed by caching content closer to users.

Global Caches



- Global Cache refers to a caching mechanism that

operates across multiple geographic regions or data centers, ensuring faster data access and improved system scalability.

Key Characteristics of Global Cache:

- ▽ Distributed Across Multiple Regions
- ▽ Reduces Latency
- ▽ Load Balancing
- ▽ High Availability

Scalability features in the Cloud



Auto-scaling



- Auto Scaling is a technique used in cloud computing and distributed systems to dynamically adjust computing resources based on real-time demand.

Comparison

Vertical Vs Horizontal	Vertical scaling	Horizontal Scaling
Data	Data is executed on a single node	Data is partitioned and executed on multiple nodes
Data Management	Easy to manage – share data reference	Complex task as there is no shared address space
Downtime	Downtime while upgrading the machine	No downtime
Upper limit	Limited by machine specifications	Not limited by machine specifications
Cost	Lower licensing fee	Higher licensing fee

Image: Google

Types of virtualization

- Server virtualization
- Storage virtualization
- Network virtualization
- Data virtualization
- Application virtualization
- Containerization



98

What is virtualization?

- Virtualization is a technology that allows multiple **virtual instances** (e.g., virtual machines, containers) to run on a **single physical machine** by abstracting hardware resources.



BITS Pilani, Pilani Campus

100

- **Vertical Scaling** is an approach to enhance the performance of the server node by adding new instances of the server to the existing servers to distribute the workload equally

Image: Google

BITS Pilani, Pilani Campus

97

BITS Pilani, Pilani Campus

99

What is Serverless Computing?

- Serverless computing is a cloud computing model where cloud providers automatically manage the infrastructure and dynamically allocate resources as needed.
- Stateless
- Managed Load Balancing



BITS Pilani, Pilani Campus
102

How Virtualization Improves Scalability?

- On-Demand Resource Allocation
- Efficient Load Balancing
- Multi-Tenancy Support
- High Availability & Fault Tolerance
- Cost-Effective Scaling



BITS Pilani, Pilani Campus
103

How Serverless Computing Enhances Scalability?

- Auto Scaling Based on Demand
- Instant Scaling and Elasticity
- Cost Efficiency
- Statelessness
- Managed Load Balancing



BITS Pilani, Pilani Campus
104

Key Characteristics of Serverless Computing

- No Server Management
- Event-Driven
- Automatic Scaling
- Pay-as-You-Go



BITS Pilani, Pilani Campus
101

Challenges of Serverless Computing for Scalability

- X Cold Start Latency
- X State Management
- X Vendor Lock-In
- X Debugging and Monitoring



Case Study: Amazon Prime

- Amazon Prime Video Uses AWS to Deliver Solid Streaming Experience to More Than 18 Million Football Fans
 - Amazon Prime needed an architecture that could quickly scale, handle spikes, and have sufficient caching in different layers to manage the demand.
 - Scalable Cloud Infrastructure
 - Content Delivery with Amazon CloudFront
 - High-Performance Video Encoding and Storage
 - Serverless Technology
 - Real-Time Analytics with Amazon Kinesis
 - Global Availability and Reliability
 - Enhanced Security and Compliance

<https://aws.amazon.com/solutions/case-studies/amazon-prime-video/>

BITS Pilani, Pilani Campus
106

Best Practices for Achieving Scalability



- Function-as-a-Service (FaaS): developers write individual functions that are executed in response to specific events
- Backend-as-a-Service (BaaS): use of managed backend services that are provided by cloud vendors
- Serverless Containers: allow the abstraction of infrastructure management while using containerized applications.
- Be Stateless
 - Use Load Balancers
 - Auto Scaling
 - Optimize Database Performance
 - Use Caching Mechanisms
 - Microservices Architecture
 - Event-Driven Architecture
 - Use Serverless Computing
 - Monitor and Analyze Performance

By following these best practices, you can build highly scalable applications that handle increased demand without compromising performance or reliability.

BITS Pilani, Pilani Campus
108

References



Text Books

- Book: *The Art of Scalability: Scalable Web Architecture, Processes, and Organizations for the Modern Enterprise*, Second Edition by Michael T. Fisher, Martin L. Abbott Published by Addison-Wesley Professional, 2015
- <https://docs.microsoft.com/en-us/azure/service-fabric/service-fabric-concepts-replica-lifecycle-balancing.html#text=tutorial%20on%20balancing%20in%20defined%20dataservers%20capable%20of%20fulfilling%20them>
- <https://azurearchitecturepatterns.com/en-us/dotnetarchitecture/microservices/architect-microservice-container-%20and%20scd/>
- <https://aws.amazon.com/>
- <https://www.nginx.com/resources/glossary/load-balancing/>
- <https://www.ibm.com/in-en/coud/learn/virtualization-a-complete-e-guide>
- <https://www.aws.amazon.com/solutions/case-studies/amazon-prime-video/>

SE ZG583, Scalable Services Lecture No. 4



Self Study



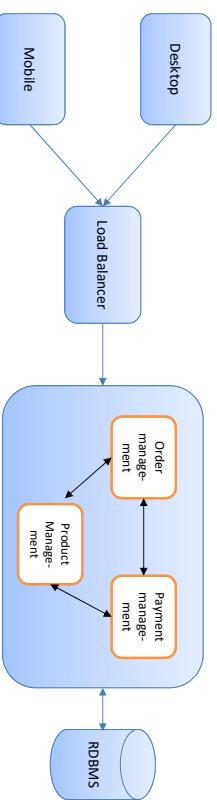
- Amazon Prime case study:
<https://aws.amazon.com/solutions/case-studies/amazon-prime-video/>
- Hotstar case study: <https://aveenai-21.medium.com/cloud-computing-aws-and-disney-hotstar-case-study-f4a3be4669a>
- Article on serverless:
<https://www.simform.com/blog/serverless-architecture-guide/>

Microservices - Introduction

Example Architecture



Online shopping



BITS Pilani, Pilani Campus

114

What is Monolithic Architecture?

- Monolith means composed all in one piece.

- They're typically complex applications that encompass several tightly coupled functions.
- When all functionality in a system had to be deployed together, we consider it a **monolith**.



BITS Pilani, Pilani Campus

113

Disadvantages of Monolith



- Scalability Challenges
- Limited Technology Flexibility
- Development Bottlenecks
- Slow & Risky Updates
- Harder Maintenance in Large Systems



Advantages of Monolithic

- Simplicity
 - Easier to Debug and Troubleshoot
 - Easier Performance Optimization
 - Easy Data Management

BITS Pilani, Pilani Campus

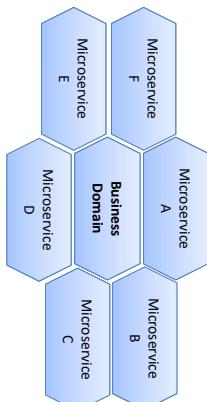
116

BITS Pilani, Pilani Campus

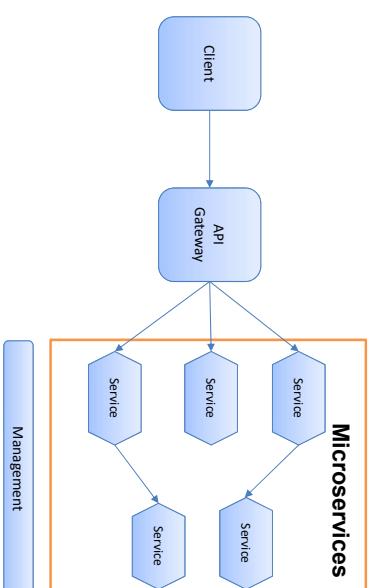
115

What is Microservices?

- Microservices are independently deployable services modeled around a business domain.
- They communicate with each other via networks,
- Each microservice can focus on a single business capability



Example Architecture



Need for Microservices

- Why is there a need to convert a fully functional monolithic application to Microservices ?
- Is the conversion worth the pain and effort?
- Should I be converting all my applications to Microservices?



BITS Pilani, Pilani Campus
118

Main characteristics of microservices

- Independent Deployment
- Decentralized Data Management
- Scalability
- Technology Diversity
- Fault Isolation
- Lightweight Communication



BITS Pilani, Pilani Campus
120

BITS Pilani, Pilani Campus
117

SOA Vs Microservices



- SOA is an enterprise-wide concept.

- Microservices architecture is an application-scoped concept

SOA



- SOA, or service-oriented architecture, defines a way to make software components reusable via service interfaces.

- These interfaces utilize common communication standards in such a way that they can be rapidly incorporated into new applications without having to perform deep integration each time.

SOA Vs Microservices



Data Duplication

- Providing services in SOA applications get hold of and make changes to data directly at its primary source, which reduces the need to maintain complex data synchronization patterns.

- In microservices applications, each microservice ideally has local access to all the data it needs to ensure its independence from other microservices, and indeed from other applications, even if this means some duplication of data in other systems.

BITS Pilani, Pilani Campus
122

SOA Vs Microservices



Communication

- In a microservices architecture, each service is developed independently, with its own communication protocol.
- With SOA, each service must share a common communication mechanism called an enterprise service bus

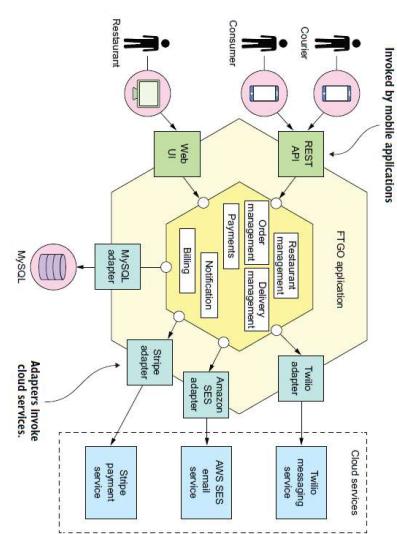
BITS Pilani, Pilani Campus
124

BITS Pilani, Pilani Campus
121

Old Architecture of the FTGO application



Case Study



126

FTGO Case Study

- Since its launch in late 2005, Food to Go, Inc. (FTGO) had grown by leaps and bounds. Today, it's one of the leading online food delivery companies in the United States.
- The business even plans to expand overseas, although those plans are in jeopardy because of delays in implementing the necessary features.
- At its core, the FTGO application is quite simple. Consumers use the FTGO website or mobile application to place food orders at local restaurants.
- FTGO coordinates a network of couriers who deliver the orders. It's also responsible for paying couriers and restaurants. Restaurants use the FTGO website to edit their menus and manage orders.
- The application uses various web services, including Stripe for payments, Twilio for messaging, and Amazon Simple Email Service (SES) for email.



BITS Pilani, Pilani Campus

128

SOA vs Microservices: Which is best for you?

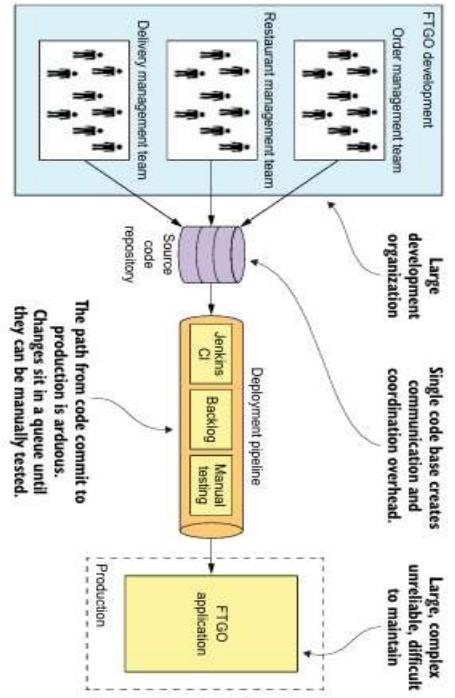
- Both approaches have their advantages, so how can you determine which one will work best for your purposes?
- In general, it depends on how large and diverse your application environment is.



BITS Pilani, Pilani Campus

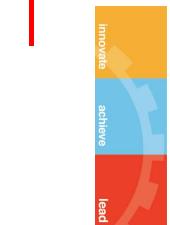
125

Monolithic Hell



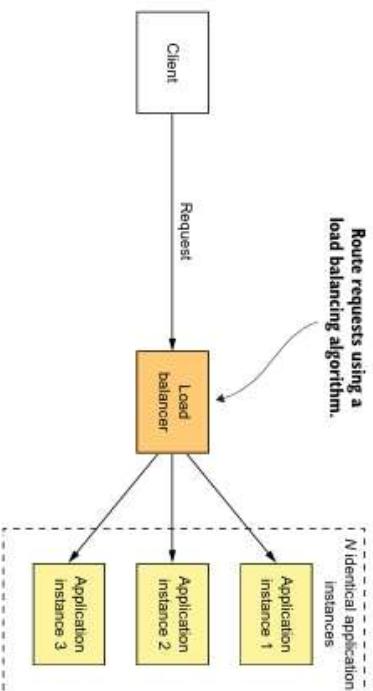
BITs Pilani, Pilani Campus
130

Problems faced in old FTGO architecture



- Complexity
- Slow development and deployment cycle
- Scaling is difficult
- Delivering a reliable monolith is challenging

Possible solution



BITs Pilani, Pilani Campus
132

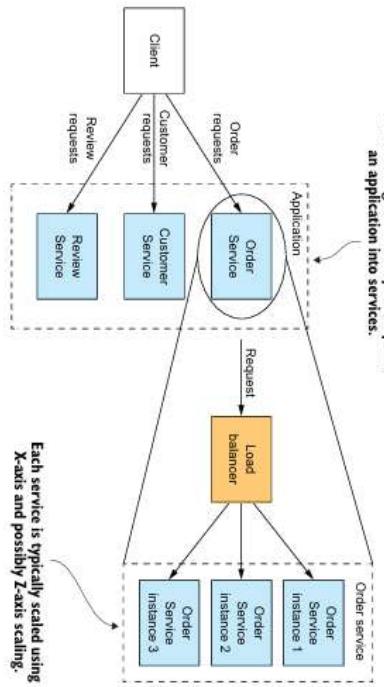
X-axis scaling



BITs Pilani, Pilani Campus
129

Y-axis scaling

Y-axis scaling functionality decomposes an application into services.

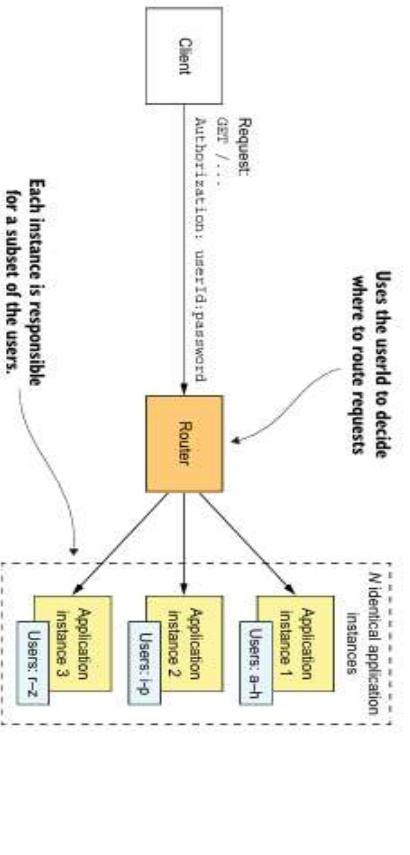


Each service is typically scaled using X-axis and possibly Z-axis scaling.



Z-axis scaling

Uses the user id to decide where to route requests



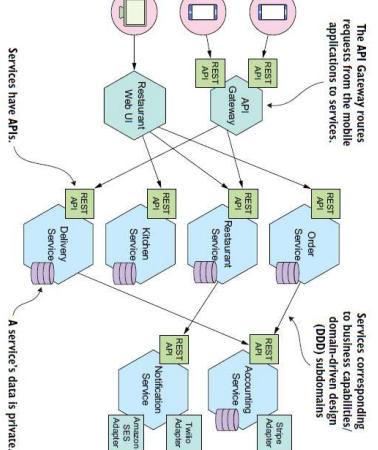
BITS Pilani, Pilani Campus
134

Netflix Case Study

- Netflix launched in 1998. At first they rented DVDs through the US Postal Service. But Netflix saw the future was on-demand streaming video
- Why are we considering this case study?
- In 2007 Netflix introduced their streaming video-on-demand service



Microservices Architecture for FTGO



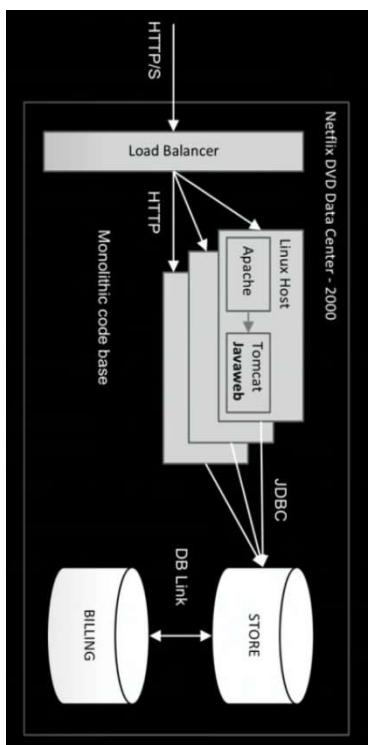
BITS Pilani, Pilani Campus
136

Each instance is responsible for a subset of the users.

BITS Pilani, Pilani Campus
133

Challenges in previous architecture

- Monolithic Code base
- Monolithic Database
- Tightly coupled Architecture

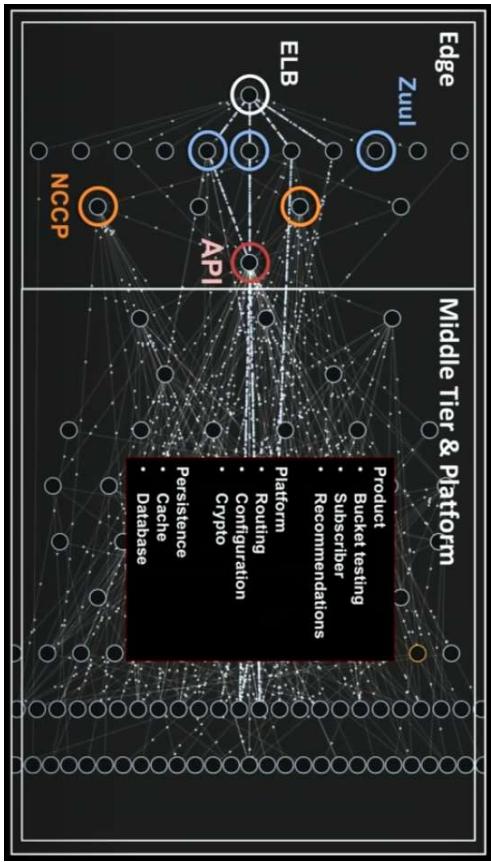


How Netflix worked earlier?

Netflix Architecture earlier

BITS Pilani, Pilani Campus
138

Updated Netflix architecture



What they need in new Architecture?

NETFLIX NEEDS

- Modularity and encapsulation
- Scalability
- Virtualization and Elasticity



BITS Pilani, Pilani Campus
140

BITS Pilani, Pilani Campus
137

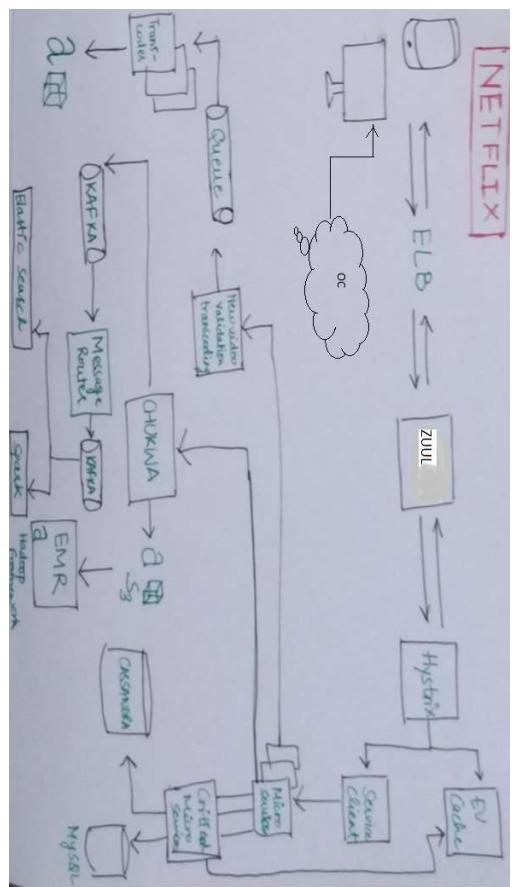
Latest Netflix Architecture



Self Study

- Journey from monolithic to SOA to Microservices:
<https://www.linkedin.com/pulse/evolving-architecture-journey-from-monolithic-soa-avita-katal/>

- Uber case study – Read about Domain Oriented Microservices Architecture. Link: <https://eng.uber.com/github-monolith-microservices/>



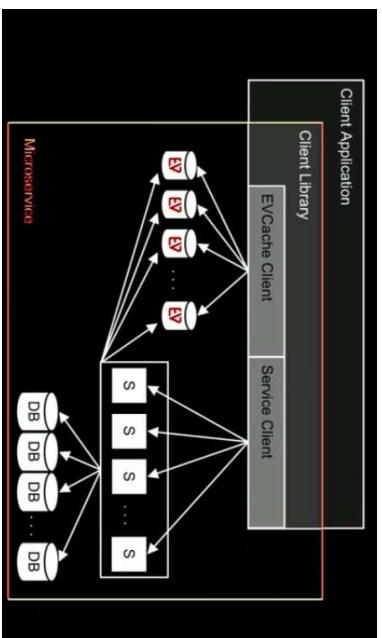
BITS Pilani, Pilani Campus

142

Netflix Approach



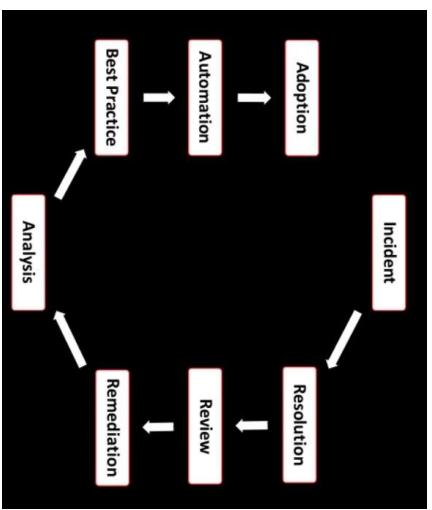
Microservice at Netflix



BITS Pilani, Pilani Campus

141

Netflix follows continuous learning

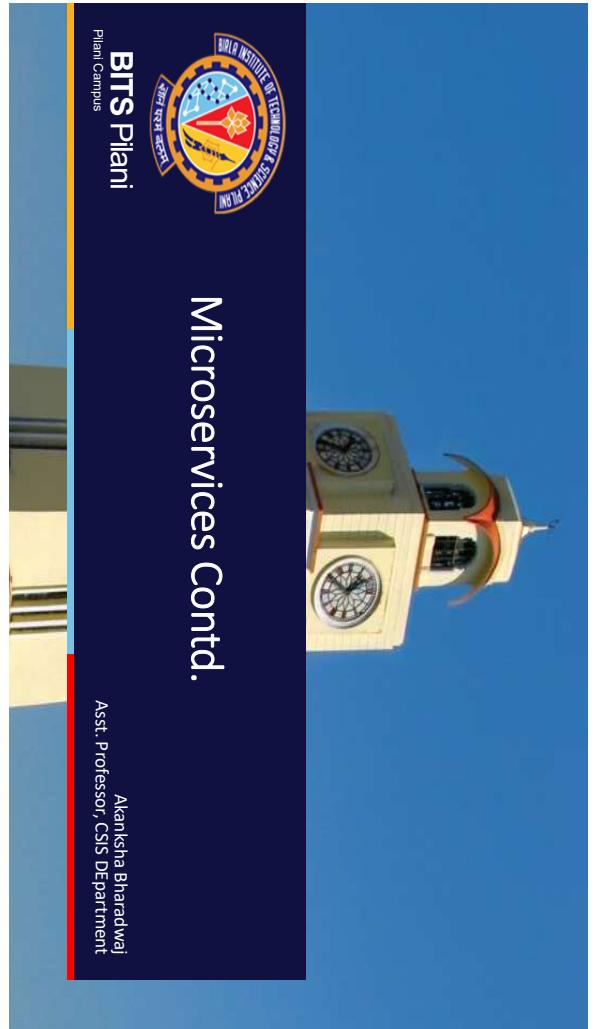


BITS Pilani, Pilani Campus

143

Advantages of Microservices

- Technology Heterogeneity
 - Different microservices can use different programming languages, databases, or frameworks.
 - Teams can choose the best technology for each service.



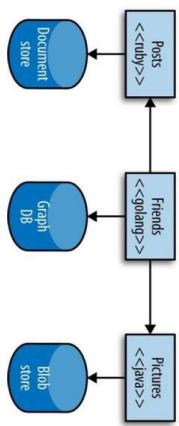
BITS Pilani
Pilani Campus

Microservices Contd.

Asst. Professor, CSIS Department

Innovate achieve lead

146



BITS Pilani, Pilani Campus

148

SE ZG583, Scalable Services Lecture No. 5



BITS Pilani
Pilani Campus

Innovate achieve lead

- Research Paper: Challenges When Moving from Monolith to Microservice Architecture. Miika Kalske, Niko Mäkitalo, and Tommi Mikkonen
- Book: Monolith to Microservices by Sam Newman
- Book: Building Microservices by Sam Newman
- Book: Microservices Vs Service Oriented Architecture by Mark Richards
- Book: Microservices Patterns by Chris Richardson
- Link: <https://www.ibm.com/cloud/blog/soa-vs-microservices>
- Link: <https://www.slideshare.net/adrianc>
- Talks about Netflix by Josh Evans

BITS Pilani, Pilani Campus

145

Advantages contd.

- Better Maintainability & Modularity
- Enhanced Security
- Support for DevOps & Agile Methodologies



Amazon Prime Video's Transition

- Amazon's Prime Video team initially adopted a microservices architecture utilizing serverless components like AWS Lambda and Step Functions to manage their video streaming services. However, they encountered significant challenges

- To address these issues, the team consolidated their services into a monolithic application hosted on Amazon EC2 and ECS.

- It's essential for organizations to carefully assess their specific needs and constraints when deciding between microservices and monolithic architectures.

BITs Pilani, Pilani Campus

150

Advantages contd.

- Resilience
- Faster Development & Deployment
- Scalability



Microservices is not a silver bullet

- Increased Operational Complexity
- Inter-Service Communication Overhead
- Data Consistency & Distributed Transactions
- Security Risks
- Higher Infrastructure & Maintenance Costs
- Steep Learning Curve
- Not Always the Best Fit



BITs Pilani, Pilani Campus

152

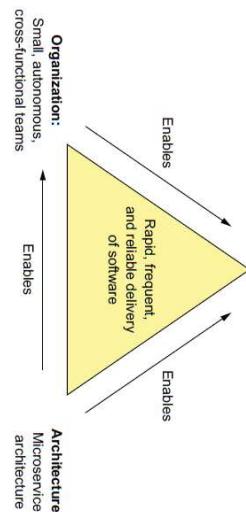
BITs Pilani, Pilani Campus

149

Introduction



Process:
DevOps/continuous delivery/deployment



BITS Pilani, Pilani Campus

154

Software development and delivery process



- If you want to develop an application with the microservice architecture, it's essential that you adopt agile development and deployment practices
- Practice continuous delivery/deployment, which is a part of DevOps.
- A key characteristic of continuous delivery is that software is always releasable

BITS Pilani, Pilani Campus

156

Software development and delivery organization



- Success means that the engineering team will grow
- The solution is to refactor a large single team into a team of teams.
- The velocity of the team of teams is significantly higher than that of a single large team.
- Moreover, it's very clear who to contact when a service isn't meeting its SLA.



Process and Organization



153

155

Microservices Design Principles



Abstraction and Information Hiding

- A service should only be consumed through a standardized API and should not expose its internal implementation details to its consumers

158

The human side of adopting microservices



Single Responsibility Principle (SRP)



- Ultimately, it changes the working environment of people and thus impact them emotionally
- Three stage transition model
- Ending, Losing, and Letting Go
 - The Neutral Zone
 - The New Beginning

160

Microservices Design Principles

Resilience & Fault Tolerance



- Each service is necessarily fault tolerant so that failures on the side of its collaborating services will have minimal impact.

Statelessness



- Services should be stateless and avoid storing session data.
- Use external caching mechanisms (e.g., Redis, Memcached) for performance optimization.

Loose Coupling & High Cohesion



Observability & Monitoring

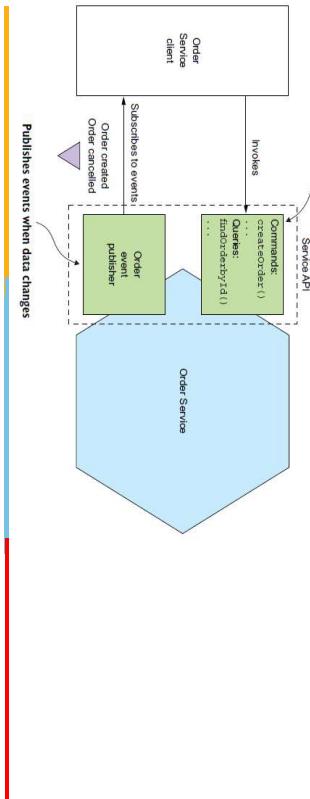


- Services should be loosely coupled to minimize dependencies.
- High cohesion ensures that each service is self-contained and modular.

- Implement logging, monitoring, and tracing (e.g., Prometheus, ELK Stack).
- Enable distributed tracing to track request flows across microservices.

What is a Service?

- It is a standalone, independently deployable software component that implements some useful functionality



BITs Pilani, Pilani Campus

166

Services: The role of Shared Libraries

- On the surface, it looks like a good way to reduce code duplication in your services.
- But you need to ensure that you don't accidentally introduce coupling between your services.
- You should strive to use libraries for functionality that's unlikely to change.



Size of a Service

- Each service should be small enough to be independently developed, deployed, and scaled, yet large enough to encapsulate a meaningful business capability.

- ✓ **Too Small** → **Increased Complexity:** Leads to too many inter-service calls, higher network latency, and dependency issues.
- ✓ **Too Large** → **Monolithic Behavior:** Reduces scalability, reusability, and independent deployments.
- ✓ **Right Size** → **Business-Oriented, Autonomous, Scalable**

Important concepts

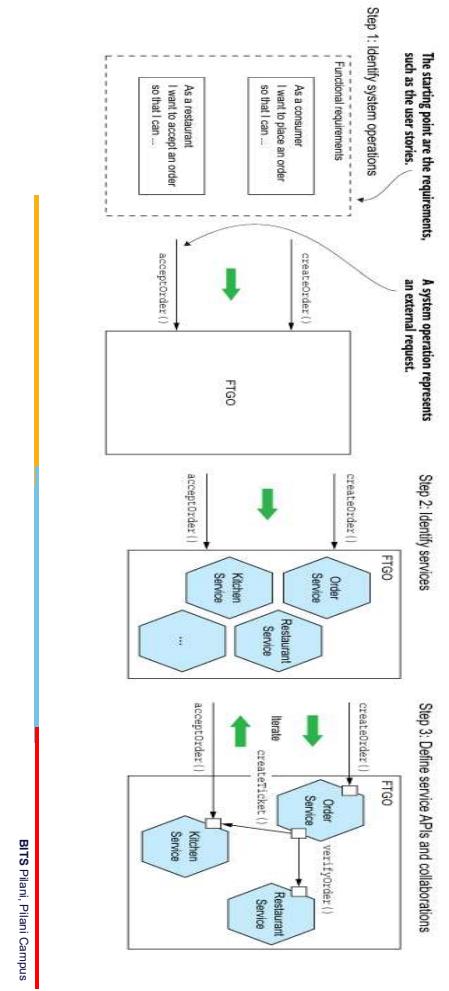


BITs Pilani, Pilani Campus

168



FTGO domain model



BITS Pilani, Pilani Campus

170

BITS Pilani, Pilani Campus

172

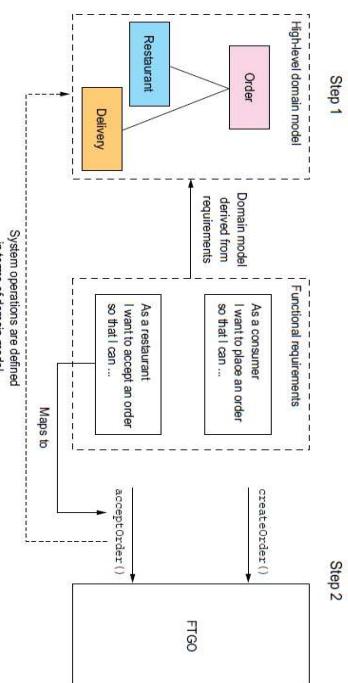
Steps for defining an application's microservice architecture

Step 1: Identify system operations

Step 2: Identify services

Step 3: Define service APIs and collaborations

Step1: Identify system operations



BITS Pilani, Pilani Campus

169

Step 2 and 3: Defining services and service API



- One strategy is defining services by applying the Decompose by business capability pattern and another is defining services by applying the Decompose by sub-domain pattern
- The next step is to define each service's API: its operations and events.
- A service API operation exists for one of two reasons: some operations correspond to system operations. They are invoked by external clients and perhaps by other services.
- The other operations exist to support collaboration between services
- The starting point for defining the service APIs is to map each system operation to a service.

BITs Pilani, Pilani Campus

174

Defining system operations



There are two types of system operations:

- Commands—System operations that create, update, and delete data
- Queries—System operations that read (query) data

Actor	Story	Command	Description
Consumer	Create Order	createOrder()	Creates an order.
Restaurant	Accept Order	acceptOrder()	Indicates that the restaurant has accepted the order and is committed to preparing it by the indicated time

Operation
restaurant.createOrder(consumerId, paymentMethod, deliveryAddress, deliveryTime, orderId, ...)

- Preconditions
- The consumer exists and can place orders.
 - The line items correspond to the restaurant's menu items.
 - The delivery address and time can be serviced by the restaurant.

- Post-conditions
- An order was created in the PENDING_ACCEPTANCE state.

Decompose by business capability pattern



- Business capability is something that a business does in order to generate value.
- Example:** The capabilities of an online store include Order management, Inventory management, Shipping, and so on.

BITs Pilani, Pilani Campus

176

BITs Pilani

Pilani Campus

Decomposition based patterns to define services

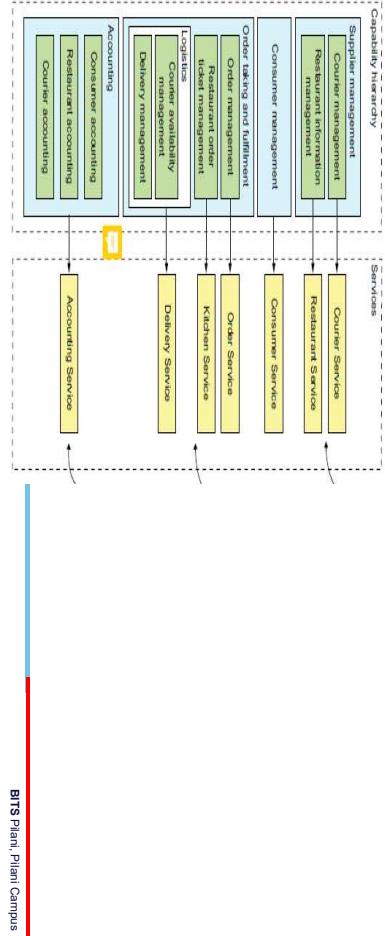


173

175

From business capabilities to services

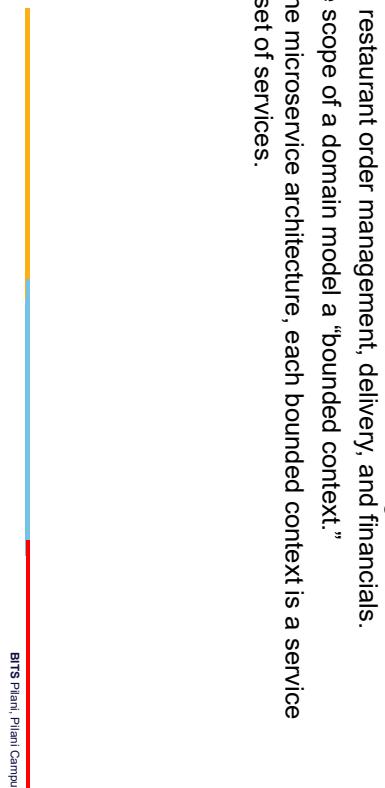
- Once you've identified the business capabilities, you then define a service for each capability or group of related capabilities



178

From Subdomains to Services

- DDD defines a separate domain model for each subdomain
- The examples of subdomains in FTGO include order taking, order management, restaurant order management, delivery, and financials
- DDD calls the scope of a domain model a "bounded context."
- When using the microservice architecture, each bounded context is a service or possibly a set of services.



180

Identifying Business Capabilities

Business capabilities for FTGO include the following:

- Supplier management
- Consumer management
- Order taking and fulfilment
- Accounting

Decompose by sub-domain pattern

Business capabilities for FTGO include the following:

- DDD is an approach for building complex software applications centered on the development of an object-oriented, domain model.
- DDD has two concepts that are incredibly useful when applying the microservice architecture: **subdomains** and **bounded contexts**.

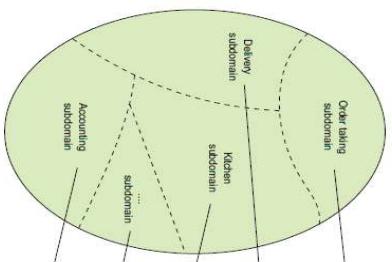
Decompose by sub-domain pattern

Business capabilities for FTGO include the following:

- Supplier management
- Consumer management
- Order taking and fulfilment
- Accounting

Decomposition guidelines

- Single Responsibility Principle
- Common Closure Principle



Decompose by sub-domain pattern

BITs Plani, Plani Campus

182



Self Study

Article: <https://kittrum.com/blog/is-microservice-architecture-still-a-trend/>
Article: <https://devops.com/microservices-amazon-monolithic-richibw/>

BITs Plani, Plani Campus

184

References

- Book: Microservices Patterns by Chris Richardson
- Book: Building Microservices by Sam Newman
- Book: Monolith to Microservices by Sam Newman
- Link: <https://docs.microsoft.com/en-us/azure/architecture/microservices/design/patterns>
- Amazon Prime Case Study: <https://theodore.ie/the-monolith-returns-why-companies-are-moving-away-from-microservices-and-serverless-architecture/>



BITs Plani, Plani Campus

181

SE ZG583, Scalable Services

Lecture No. 6



Microservices Contd.

Obstacles to decomposing an application into services

- Network latency
- Reduced availability due to synchronous communication
- Maintaining data consistency across services
- Obtaining a consistent view of the data
- God services preventing decomposition

Defining service APIs



Key Principles of Service APIs



Example



- **Loose Coupling** – Minimize dependencies between services.
- **High Cohesion** – Each API should be specific to a microservice's functionality.
- **Standardized Communication** – Use REST, gRPC, GraphQL, or event-driven messaging.
- **Backward Compatibility** – Ensure versioning to prevent breaking changes.
- **Security** – Implement authentication, authorization, and encryption.

BITS Pilani, Pilani Campus
190

Service	Operations
Consumer Service	createConsumer()
Order Service	createOrder()
Restaurant Service	findAvailableRestaurants()
Kitchen Service	■ acceptOrder() ■ noteOrderReadyForPickup()
Delivery Service	■ noteUpdatedLocation() ■ noteDeliveryPickedup() ■ noteDeliveryDelivered()

BITS Pilani, Pilani Campus
192

Introduction



A service API operation exists for one of two reasons:

- **To Expose Business Capabilities** – The operation provides functionality that aligns with a business use case, such as processing orders, managing users, or retrieving product details.
- **To Support System Integration** – The operation facilitates interaction between services, allowing them to communicate, share data, and maintain consistency across the system.

Assigning system operations to services



- Many system operations neatly map to a service, but sometimes the mapping is less obvious.

We can either

- Assign an operation to a service that needs the information provided by the operation

- Assign an operation to the service that has the information necessary to handle it

BITS Pilani, Pilani Campus
189

Example



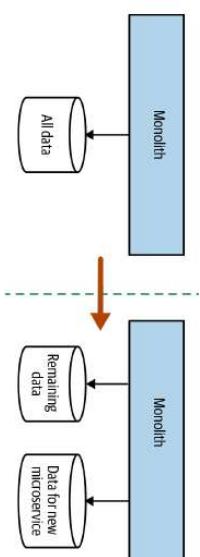
Split the Database First



Service	Operations	Collaborators
Consumer Service	verifyConsumerDetails()	—
Order Service	createOrder()	■ Consumer Service ■ verifyConsumerDetails() ■ Restaurant Service ■ verifyOrderDetails() ■ Kitchen Service ■ createTicket() ■ Accounting Service ■ authorizeCard()
Restaurant Service	findAvailableRestaurants()	—
Kitchen Service	■ createTicket() ■ noteOrderReadyForPickup()	—
Delivery Service	■ scheduledDelivery() ■ noteUpdatedLocation() ■ noteDeliverypickup() ■ noteDeliveryDelivered()	■ Delivery Service ■ scheduleDelivery() ■ acceptOrder() —
Accounting Service	■ authorizeCard()	—

BITS Pilani, Pilani Campus

194



BITS Pilani, Pilani Campus

195

Determining the APIs required to support collaboration between services



- Some system operations are handled entirely by a single service

- Some system operations span across multiple services.

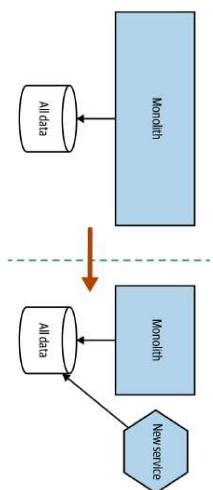
Transition from monolith to microservices



Split the Code First



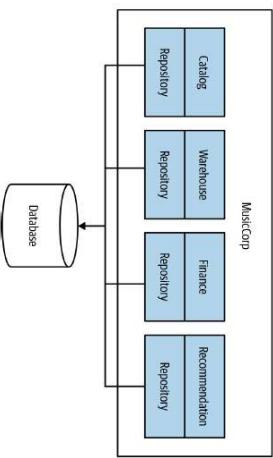
BITS Pilani
Pilani Campus



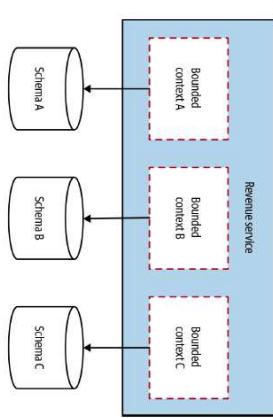
BITS Pilani, Pilani Campus
198

Split the Database First

Pattern: Repository per bounded context

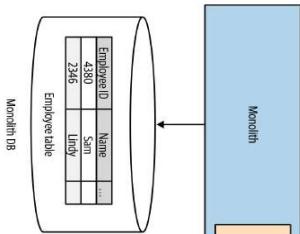


Pattern: Database per bounded context

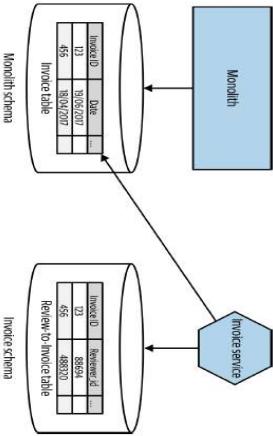


Patterns for Monolith to Microservices

Pattern: Monolith as data access layer



Pattern: Multi-schema storage



200

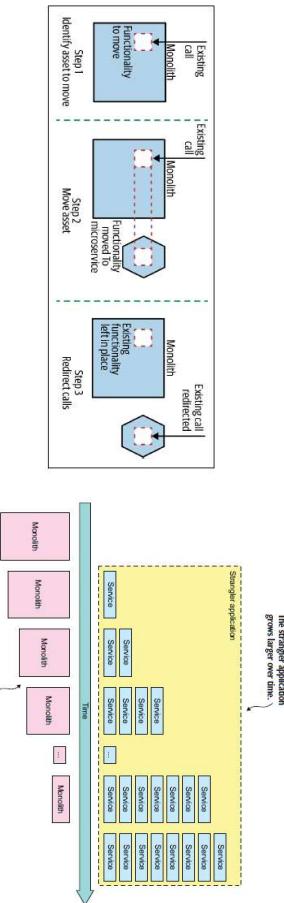


BITS Pilani, Pilani Campus
197

BITS Pilani, Pilani Campus
199

Strangler Pattern

- The Strangler Pattern is a popular design pattern to incrementally transform your monolithic application into microservices by replacing a particular functionality with a new service.
- Any new feature to be added is done as part of the new service

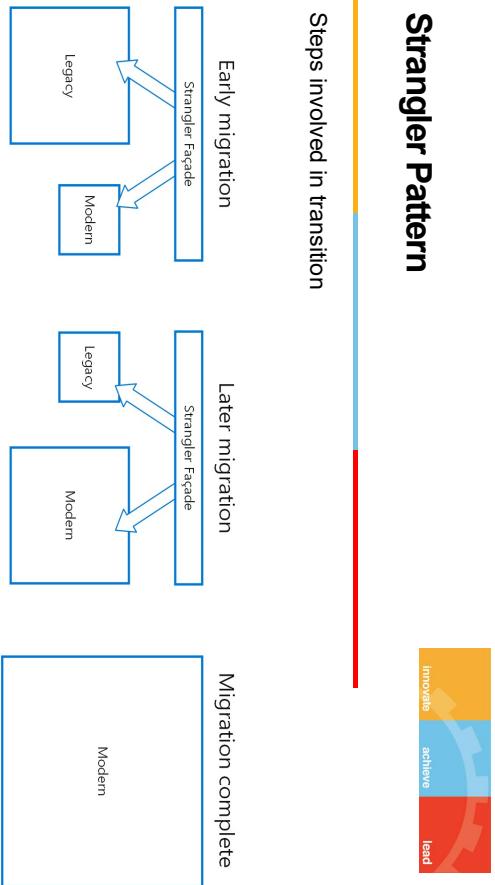


Rebuild From Scratch

- One of the biggest challenges for us was having a good understanding of the legacy system.
- Can not use the system until complete
- Longer duration required

Strangler Pattern

Steps involved in transition



- What component to start with?
- How to handle services and data stores that are potentially used by both new and legacy systems?
- Migration

Strangler Pattern: Issues

The strangler application grows bigger over time.

BTS Plan! Plan Campus
202

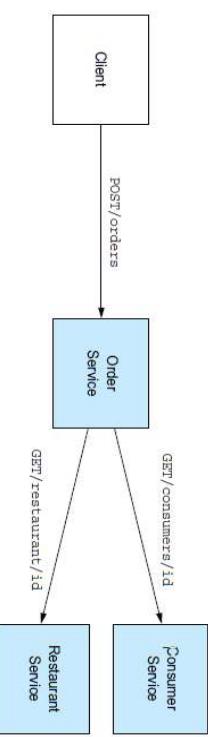
The monolith shrinks over time.

204

Synchronous communication

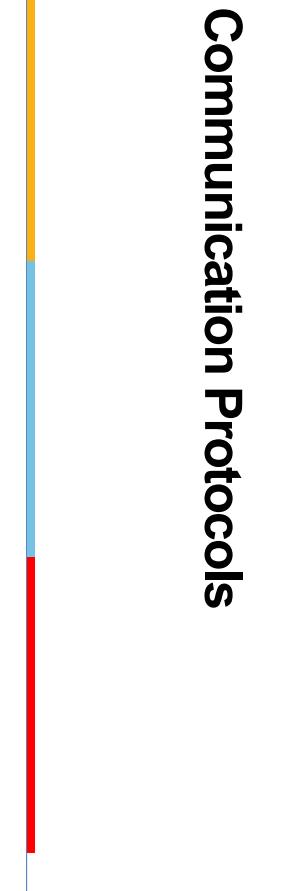
- REST is an extremely popular IPC mechanism under this category

Example: FTGO CreateOrder request



208

Communication Protocols



206

When not to use Strangler Pattern?

- When requests to the back-end system cannot be intercepted.
- For smaller systems where the complexity of wholesale replacement is low.

Aspects of communication

Communication Type

- Synchronous protocol:** The client sends a request and waits for a response from the service.
- Asynchronous protocol:** The client code or message sender usually doesn't wait for a response.



Number of Receivers

- Single receiver
- Multiple receivers



Guiding Principles of REST



- Resource-Oriented Design
- Stateless
- Cacheable
- Uniform interface
- API Versioning
- Service Discovery & Load Balancing

BITS Pilani, Pilani Campus
210

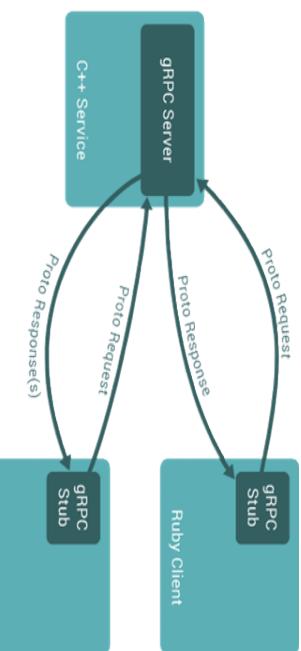
Representational state transfer (REST)

- It is a software architectural style that was created to guide the design and development of the architecture for the World Wide Web.

BITS Pilani, Pilani Campus
212

gRPC: Introduction

- gRPC (Google Remote Procedure Call) is a high-performance, language-agnostic RPC (Remote Procedure Call) framework that is well-suited for microservices-based architectures



gRPC



- gRPC clients and servers can run and talk to each other in a variety of environments - from servers inside Google to your own desktop - and can be written in any of gRPC's supported languages.

- So, for example, you can easily create a gRPC server in Java with clients in Go, Python, or Ruby.

- In addition, the latest Google APIs will have gRPC versions of their interfaces, letting you easily build Google functionality into your applications

BITS Pilani, Pilani Campus
212

BITS Pilani, Pilani Campus
209



2. Strongly Typed Contracts

- Enforces a strict schema using .proto files, ensuring consistency between services.
- Auto-generates client and server code in multiple languages, reducing manual coding errors.



4. Language-Agnostic

- gRPC supports multiple programming languages (**Java, Python, Go, C++, etc.**), making it ideal for polyglot microservices environments.

BITS Pilani, Pilani Campus
214



Why use gRPC in Microservices?

1. High Performance & Efficiency

- Uses Protocol Buffers (ProtocolBuf), a compact, binary format that is faster and smaller than JSON over HTTP.
- Uses HTTP/2 for multiplexed connections, reducing latency and improving efficiency.

3. Supports Streaming

Unlike REST, gRPC supports real-time communication using:

- Unary RPC (Single request-response, like REST)
- Server Streaming (Server sends multiple responses for a single request)
- Client Streaming (Client sends multiple requests and gets a single response)
- Bi-directional Streaming (Both client and server send multiple messages in real-time)

BITS Pilani, Pilani Campus
216



Why use gRPC in Microservices?

1. High Performance & Efficiency

- Uses Protocol Buffers (ProtocolBuf), a compact, binary format that is faster and smaller than JSON over HTTP.
- Uses HTTP/2 for multiplexed connections, reducing latency and improving efficiency.

3. Supports Streaming

Unlike REST, gRPC supports real-time communication using:

- Unary RPC (Single request-response, like REST)
- Server Streaming (Server sends multiple responses for a single request)
- Client Streaming (Client sends multiple requests and gets a single response)
- Bi-directional Streaming (Both client and server send multiple messages in real-time)

BITS Pilani, Pilani Campus
213

Asynchronous communication

- It means communication which happens 'out of sync' or in other words; not in real-time.
- For example, an email to a colleague would be classed as asynchronous communication

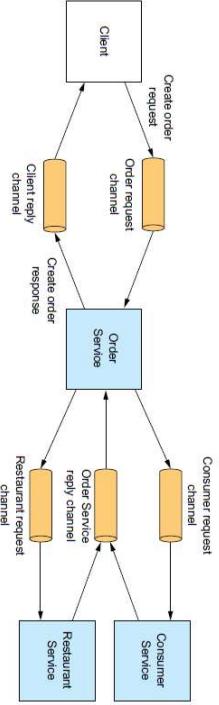


Communication Styles



Asynchronous Communication Example

- Services communicating by exchanging messages over messaging channels.



5. Built-in Authentication & Security

- Supports **TLS encryption** for secure communication.
- Works with authentication mechanisms like **OAuth 2.0** and **JWT**.

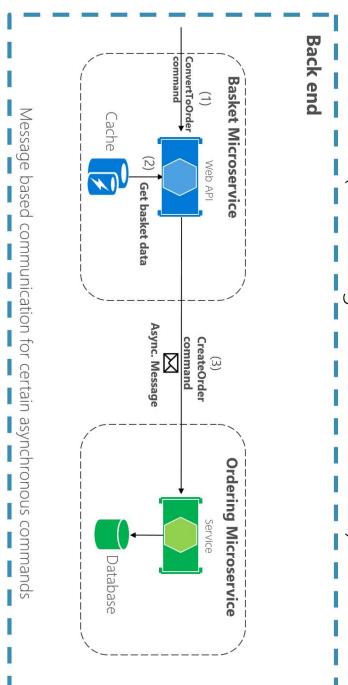


Single-receiver message-based communication



Single receiver message-based communication

(i.e. Message-based Commands)

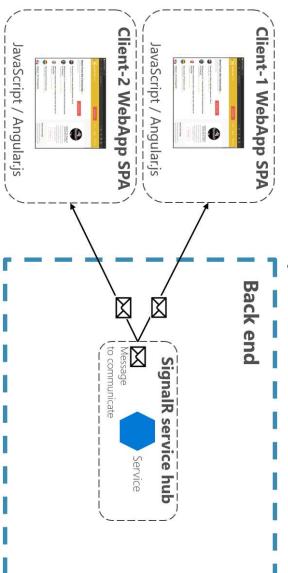


BITS Pilani, Pilani Campus
222

Push and real-time communication based on HTTP

Push and real-time communication based on HTTP

One-to-many communication



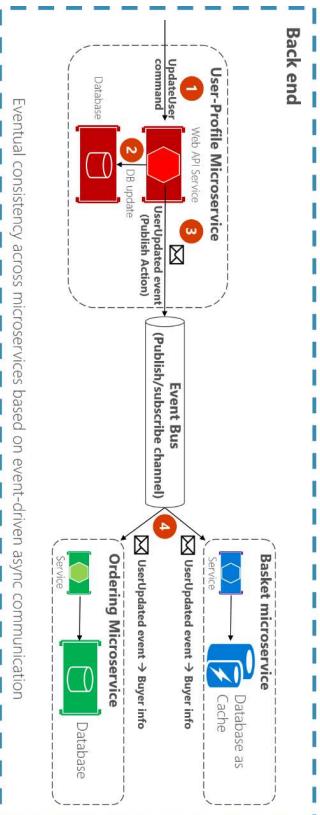
BITS Pilani, Pilani Campus
222

Asynchronous event-driven communication



Asynchronous event-driven communication

Multiple receivers



BITS Pilani, Pilani Campus
224

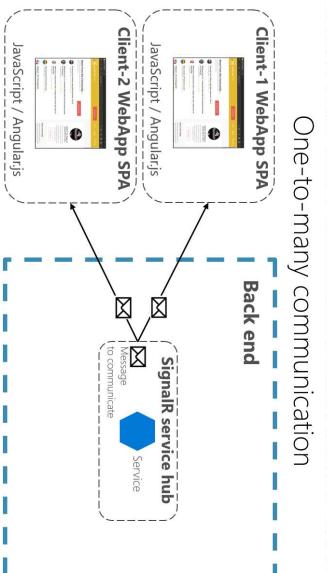
Multiple-receivers message-based communication

Multiple-receivers message-based communication

Use a publish/subscribe mechanism so that your communication from the sender will be available to all the subscriber microservices or to external applications.



Back end



Back end

BITS Pilani, Pilani Campus
221

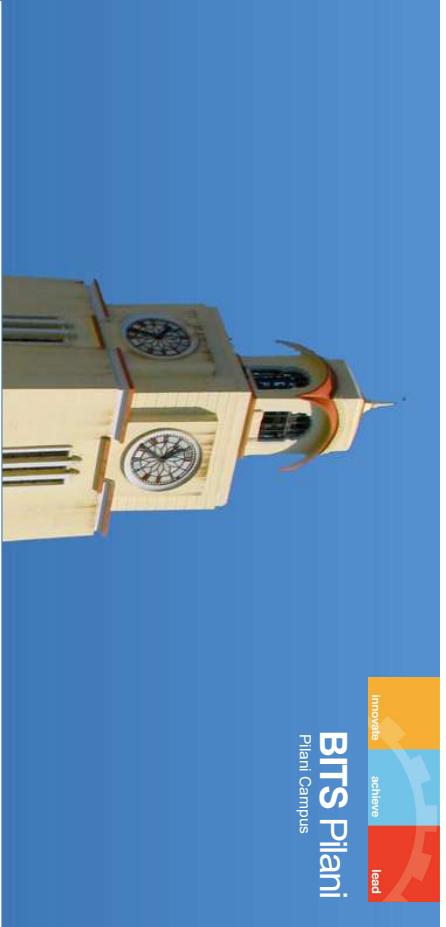
223

REST API (Synchronous, Request-Response)



Example Scenario: User Authentication & Profile Management

- A user logs into an **e-commerce platform** using email and password. The system validates the credentials and returns a response.
- Similarly, fetching user details (name, address, order history) should be a simple request-response operation.



Kafka (Event-Driven, Asynchronous Messaging)



Example Scenario: Order Processing in an E-Commerce System

- When a user places an order, multiple microservices need to react to this event:
 - Order Service** → Saves order details.
 - Inventory Service** → Updates stock levels.
 - Payment Service** → Charges the user.
 - Notification Service** → Sends confirmation emails/SMS.
- Kafka acts as an **event bus** that ensures all these services process the event without direct coupling.



gRPC (High-Performance Remote Procedure Calls)



Example Scenario: Ride-Sharing App's Driver Location Tracking

- A passenger requests a ride, and the system continuously updates the **driver's real-time location** every few seconds. Since performance and efficiency are crucial, gRPC's low latency and streaming features make it ideal.

Choosing the Right Communication Pattern for Microservices

WebSockets (Persistent, Real-Time Bi-Directional Communication)



Example Scenario: Stock Market Price Updates

- A stock trading platform needs to show **real-time stock price changes** to users without them having to refresh the page.

- Book: Microservices Patterns by Chris Richardson
- Book: Monolith to Microservices by Sam Newman
- Link: <https://microservices.io/patterns>
- Link: <https://docs.microsoft.com/en-us/dotnet/architecture/microservices/architect-microservice-containet-applications/communication-in-microservice-architecture>
- https://www.ics.uci.edu/~fielding/pubs/dissertation/test_arch_style.htm
- <https://gRPC.io/docs/what-is/gRPC/introduction/>

BITS Pilani, Pilani Campus

230

RabbitMQ (Message Queuing, Reliable Delivery)



Example Scenario: Asynchronous Task Processing in a Banking System

- A banking system needs to process transactions in a **fraud detection microservice**. Instead of waiting synchronously, the **Transaction Service** places a message in a queue, and the **Fraud Detection Service** processes it asynchronously.

BITS Pilani, Pilani Campus

231

Comparison



Feature	REST	gRPC	Kafka	RabbitMQ	WebSockets
Type	Request-Response	Remote Procedure Call	Event Streaming	Message Queue	Persistent Connection
Best For	Web APIs, CRUD	Microservices, Low Latency	High-Volume Events	Reliable Async Tasks	Real-Time Updates
Data Format	JSON/XML	Protobuf	Binary	Any	JSON/Binary
Communication	Synchronous	Synchronous/Streaming	Asynchronous	Asynchronous	Full Duplex
Browser Support	Yes	No	No	Yes	
Throughput	Medium	High	Very High	High	Medium
Latency	High	Low	Medium	Low	Very Low
Persistence	No	No	Yes	Yes	No

References



BITS Pilani, Pilani Campus

229

Direct Communication

A client could make requests to each of the microservices directly

But there are challenges and limitations with this option

- Inefficient to call directly.
- When the client directly is calling the microservices they might use protocols that are not web-friendly.



SE ZG583, Scalable Services

Lecture No. 7

234

Introduction

Let's imagine that you are developing a native mobile client for a shopping application.

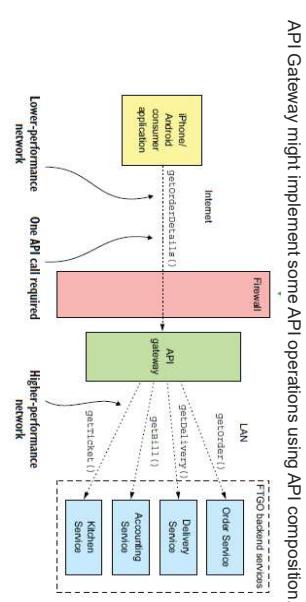
The product details page displays a lot of information

- Number of items in the shopping cart
- Order history
- Basic Product Information
 - Customer reviews
 - Low inventory warning
 - Shipping options
 - Various suggested items

How the mobile client accesses these services?

Using an API Gateway: Request Routing

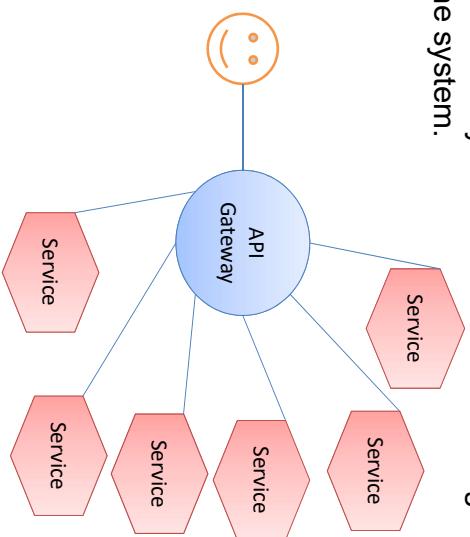
An API gateway implements some API operations by routing requests to the corresponding service.



BITS Pilani, Pilani Campus
238

Using an API Gateway

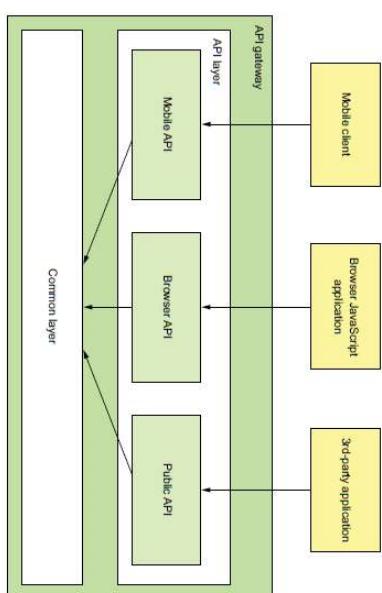
An API Gateway is a server that is the single entry point into the system.



BITS Pilani, Pilani Campus
237

API Gateway Architecture

An API gateway implements some API operations by routing requests to the corresponding service.



BITS Pilani, Pilani Campus
240

Using an API Gateway: Protocol Translation

It might provide a RESTful API to external clients

API Gateway provides each client with client-specific API

Android client and IOS client separate

Implementing edge functions

- Authentication
- Authorization
- Rate limiting
- Caching
- Metrics collection
- Request logging

BITS Pilani, Pilani Campus
239

Drawbacks of an API Gateway

- Another component that must be developed, deployed, and managed.
- API gateway can become a development bottleneck

Designing APIs

- A service's API is a contract between the service and its clients.
- The challenge is that a service API isn't defined using a simple programming language construct.

- The nature of the API definition depends on which IPC mechanism you're using.

- APIs invariably change over time as new features are added, existing features are changed, and old features are removed

Benefits of an API Gateway

- It encapsulates internal structure of the application.
- It also simplifies the client code

What is API design?

API design refers to the process of developing application programming interfaces (APIs) that expose data and application functionality for use by developers and users.



Good API design

- Simplicity
- Flexibility
- Consistency
- Scalability



What Is The API Contract?

The API contract is the agreement between the API producer and consumer.

An API contract is also a way to keep track of the changes made to an API.

BITS Pilani, Pilani Campus
246

247

How do we design our API program

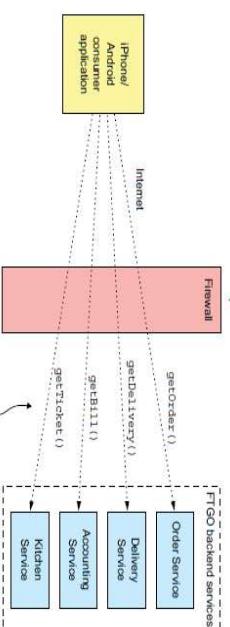
Teams must ask themselves:

- What technology is used to build the APIs?
- How are the APIs designed?
- How are the APIs maintained?
- How are the APIs promoted inside the organization or marketed to the outside world?
- What resources are available?
- Who should be on the team?



External API design issues

- One approach to API design is for clients to invoke the services directly.
- But this approach is rarely used in a microservice architecture because of a lot of drawbacks



BITS Pilani, Pilani Campus
248

Figure 8.2 | A client can retrieve the order details from the monolithic F1GO application with a single request. But the client must make multiple requests to retrieve the same information in a microservice architecture.

BITS Pilani, Pilani Campus
245

246

What is API security?



- API security is the protection of the integrity of APIs, both the ones you own and the ones you use.

- REST APIs use HTTP and support Transport Layer Security (TLS) encryption.
- TLS is a standard that keeps an internet connection private and checks that the data sent between two systems (a server and a server, or a server and a client) is encrypted and unmodified.

BITS Pilani, Pilani Campus
250

When To Create An API Version



API versioning should occur any time you:

- Change fields or routing after an API is released.
- Change payload structures, such as changing a variable from an integer to float, for instance.
- Remove endpoints to fix design or poor implementation of HTTP.

BITS Pilani, Pilani Campus
252

Why is API security important?



- Broken, exposed, or hacked APIs are behind major data breaches.
- If your API connects to a third party application, understand how that app is funneling information back to the internet.

BITS Pilani, Pilani Campus
251

REST API security

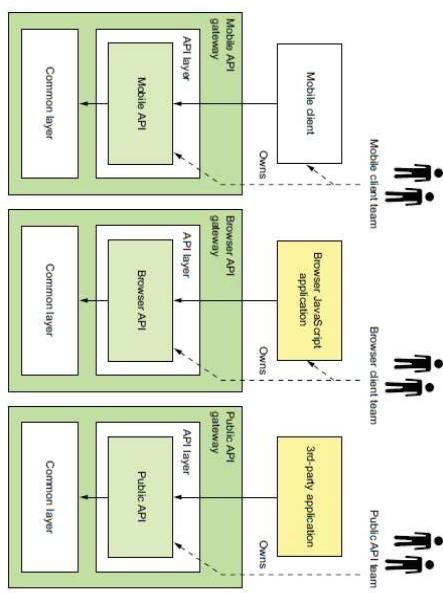


BITS Pilani, Pilani Campus
249

Backends for Frontends



BITS Pilani
Pilani Campus



BITS Pilani, Pilani Campus
254

Most common API security best practices



BITS Pilani
Pilani Campus

- Use tokens
- Use encryption and signatures
- Identify vulnerabilities
- Use quotas and throttling
- Use an API gateway

Benefits of BFF pattern

Database related patterns for Microservices



BITS Pilani
Pilani Campus

- The API modules are isolated from one another, which improves reliability.

- It also improves observability, because different API modules are different processes

- Another benefit of the BFF pattern is that each API is independently scalable.
- The BFF pattern also reduces startup time because each API gateway is a smaller, simpler application.

256

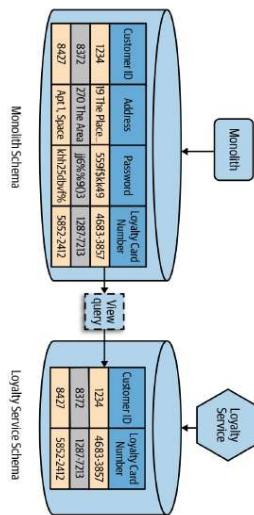
BITS Pilani, Pilani Campus
253

Database View Pattern

- For a single source of data for multiple services, a view can be used to mitigate the concerns regarding coupling

Limitations:

- Views are read-only.

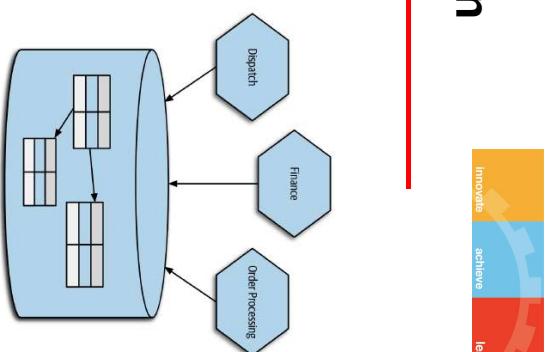


BITS Pilani, Pilani Campus
258

Shared Database pattern

Problems:

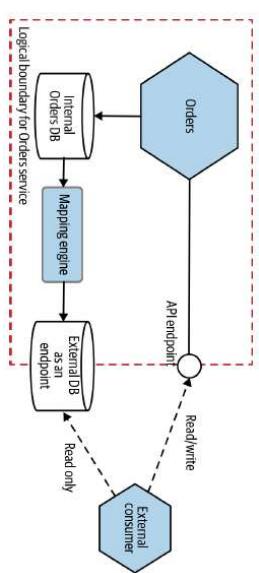
- Shared Vs hidden data
- Control on data



BITS Pilani, Pilani Campus
258

Database-as-a-Service Pattern

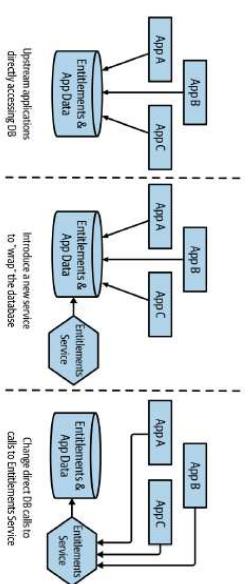
Can be used when clients just need a database to query.



BITS Pilani, Pilani Campus
260

Database Wrapping Service Pattern

- Here we hide the database behind a service that acts as a thin wrapper



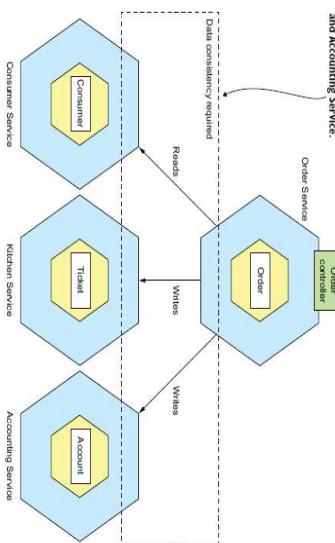
BITS Pilani, Pilani Campus
260



Challenges with distributed transactions

- The traditional approach to maintaining data consistency across multiple services, databases, or message brokers is to use distributed transactions

The `createOrder()` operation reads from Consumer Service, Kitchen Service, and Accounting Service.



262
BITS Pilani, Pilani Campus

Need for distributed transactions in a MSA

- Imagine that you're the FTGO developer responsible for implementing the `create- Order()` system operation

- this operation must verify that the consumer can place an order, verify the order details, authorize the consumer's credit card, and create an Order in the database

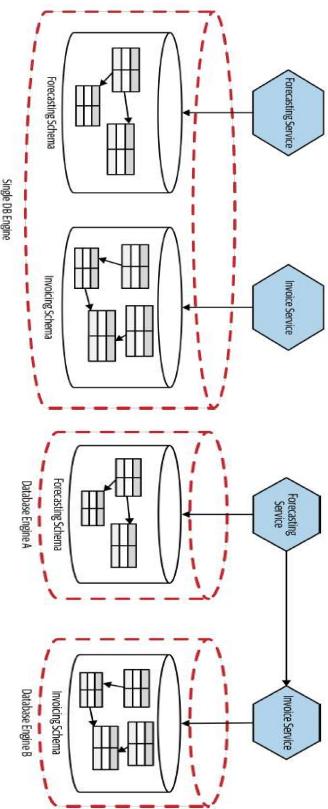
Transaction management in a microservice architecture



262

Splitting Apart the Database

Physical Versus Logical Database Separation



Solution: Saga pattern



- Sagas are mechanisms to maintain data consistency in a microservice architecture without having to use distributed transactions.

Challenges in Saga

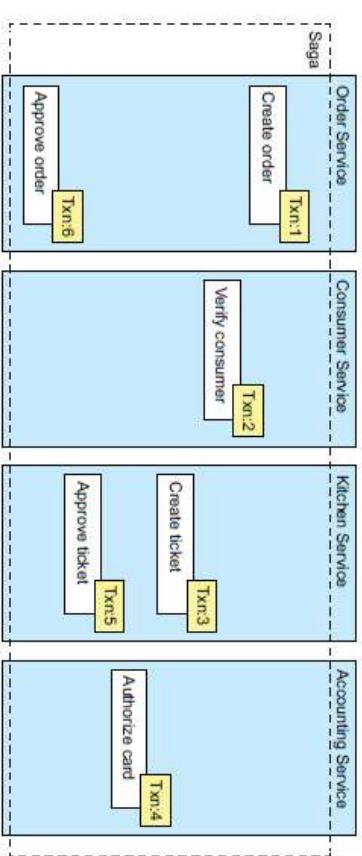


- lack of isolation between sagas
- rolling back changes when an error occurs



BITS Pilani, Pilani Campus
266

Example



BITS Pilani, Pilani Campus
268

BITS Pilani, Pilani Campus
265

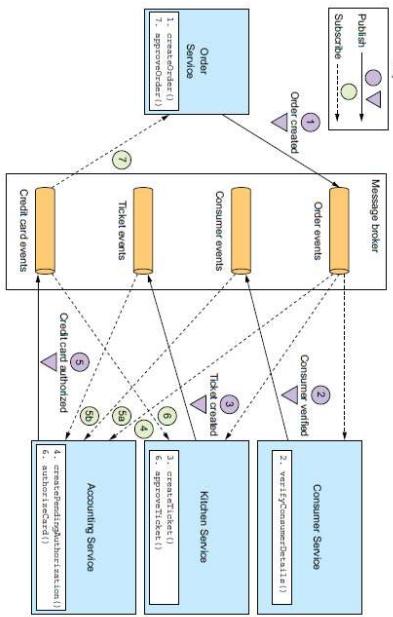
267

Create Order Saga



References

1. Book: Microservices Patterns by Chris Richardson
2. Link: <https://www.nginx.com/blog/building-microservices-using-an-api-gateway/>
3. <https://www.redhat.com/en/topics/api/what-is-api-design>
4. <https://www.redhat.com/en/topics/security/api-security>
5. <https://marutitech.com/api-gateway-in-microservices-architecture/>
6. <https://www.getambassador.io/blog/optimizing-microservices-architecture>
7. <https://www.postman.com/api-platform/api-design/>



BITS Pilani, Pilani Campus
270

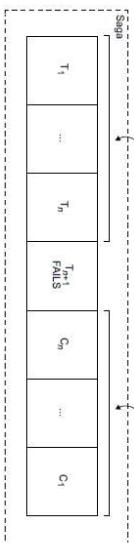
Compensating transactions



- Unfortunately, sagas can't be automatically rolled back, because each step commits its changes to the local database.

The changes made by $T_1 \dots T_n$ have been committed.

The compensating transactions undo the changes made by $T_1 \dots T_n$.



Step	Service	Transaction	Compensating transaction
1	order service	createOrder()	rejectOrder()
2	consumer service	verifyConsumerDetails()	-
3	Kitchen service	createTicket()	rejectTicket()
4	Accounting service	authorizeCreditCard()	-
5	Kitchen service	approveTicket()	-
6	order service	approveOrder()	-

BITS Pilani, Pilani Campus
269

Self Study



- API in detail
- Sagas in detail

BITS Pilani, Pilani Campus
272



BITS Pilani
Pilani Campus

- For Processing Really BIG Data
- For Storing a Diverse Set of Data
- For Parallel Data Processing

When to Use Hadoop



SE ZG583, Scalable Services Lecture No. 8

274



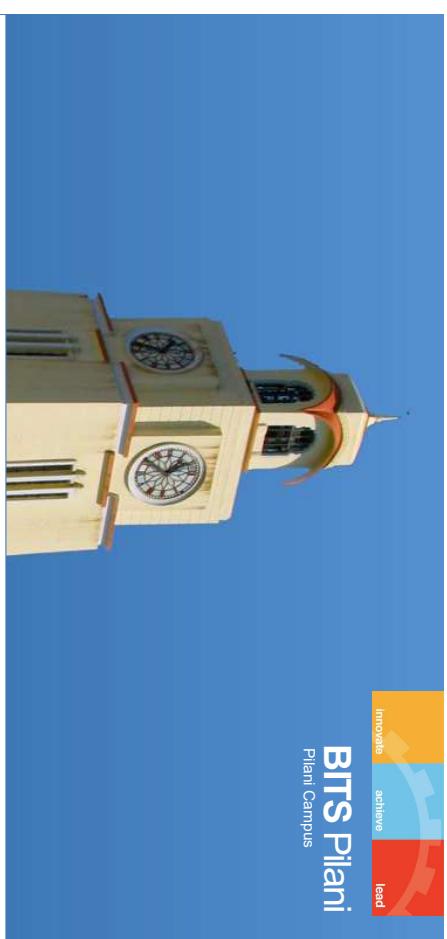
BITS Pilani
Pilani Campus

BITS Pilani, Pilani Campus

276



BITS Pilani
Pilani Campus



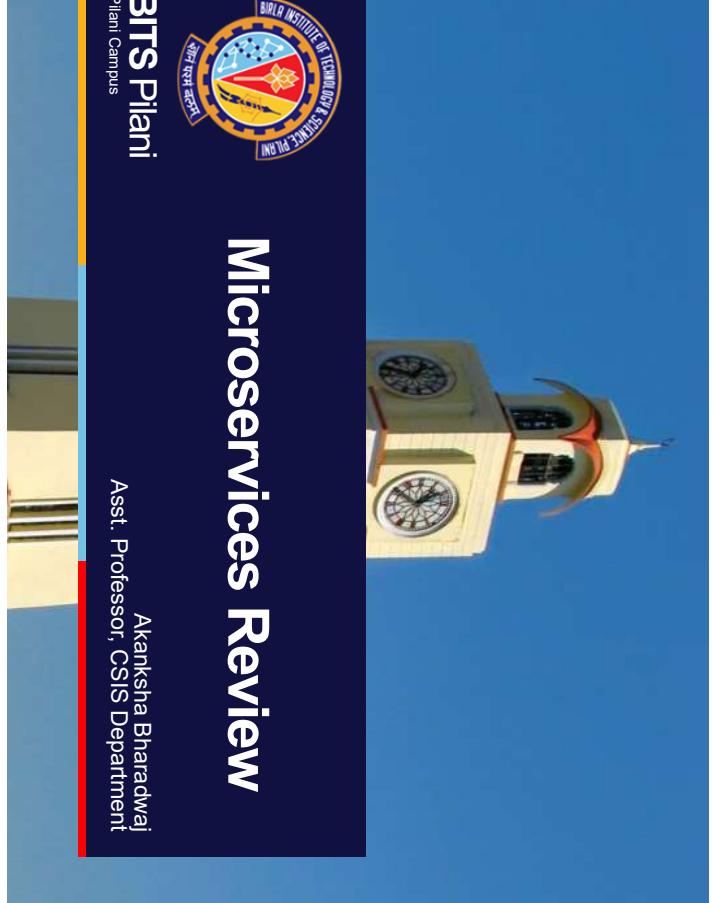
Review



BITS Pilani
Pilani Campus



BITS Pilani
Pilani Campus



Microservices Review



BITS Pilani
Pilani Campus



BITS Pilani
Pilani Campus

BITS Pilani
Pilani Campus

Microservices Review
Asst. Professor, CSIS Department

273

Homework: Kafka

Compare and contrast Kafka's event streaming architecture with traditional message queuing systems. What advantages does Kafka provide in terms of scalability and fault tolerance?



When to use CDN

- If your platform serves a global audience, and you aim to provide a seamless viewing experience, then a CDN is a must-have.
- CDN for large high-load websites.



Kafka Vs Rabbit MQ

- RabbitMQ and Apache Kafka allow producers to send messages to consumers.
- Producers and consumers interact differently in RabbitMQ and Kafka.
 - In RabbitMQ, the producer sends and monitors if the message reaches the intended consumer.
 - On the other hand, Kafka producers publish messages to the queue regardless of whether consumers have retrieved them.
- Kafka is suitable for applications that need to reanalyze the received data. You can process streaming data multiple times within the retention period or collect log files for analysis. Log aggregation with RabbitMQ is more challenging, as messages are deleted once consumed.
- Rabbit MQ suits applications that must adhere to specific sequences and delivery guarantees when exchanging and analyzing data.



BITS Pilani, Pilani Campus

278



BITS Pilani, Pilani Campus

280

Aspect	Kafka Event Streaming	Traditional Message Queuing
Data Model	Log-centric, events stored in immutable logs.	Point-to-point or publish-subscribe messaging model.
Storage	Persists messages for a configurable retention period.	Messages are often transient and may be deleted after consumption.
Message Distribution	Publis-subscribe model with partitioned topics.	Follows either point-to-point or publish-subscribe patterns.
Scalability	Scales horizontally by adding more brokers and partitions.	Scales often involves vertical scaling by increasing resources.
Fault Tolerance	Achieves fault tolerance through replication of partitions.	Relies on mechanisms like clustering for fault tolerance.
Acknowledgments and Guarantees	Configurable delivery guarantees (at most once, at least once, exactly once).	Various acknowledgment and delivery guarantee mechanisms.
Use Cases	High-throughput, durability, real-time data pipelines, streaming analytics.	Point-to-point communication, traditional messaging scenarios.
Integration Ecosystem	Rich ecosystem with connectors for various data sources and sinks.	Often integrated with middleware and enterprise messaging systems.
Horizontal Scalability	Enabled through partitioning and adding more brokers.	Scaling may involve vertical scaling by increasing resources.

BITS Pilani, Pilani Campus

277

BITS Pilani, Pilani Campus

279

Edge computing scenarios



- Enterprise edge
- Operations edge
- Provider edge

BITS Pilani, Pilani Campus
282

When to use Load balancer

- Load Balancing for Scale
- High availability



BITS Pilani, Pilani Campus
283

When serverless architecture is not the right choice



- Entirely Serverless application is not suitable for real-time applications that use WebSockets because FaaS functions have limited lifetime
- After some time of being idle, function will require to go through a cold start which can take up to a few seconds.
- Different FaaS providers may differ in some particularities of using their services which will make the switch to another provider troublesome.

When to use CQRS pattern

- Collaborative domains where many users access the same data in parallel.
- You can use the CQRS pattern to separate updates and queries if they have different requirements for throughput, latency, or consistency.



BITS Pilani, Pilani Campus
284

BITS Pilani, Pilani Campus
281



DevOps

- DevOps is a **software development methodology** that integrates **development (Dev)** and **IT operations (Ops)** to improve collaboration, automate processes, and accelerate software delivery.

SE ZG583, Scalable Services

Lecture No. 9



286



BITS Pilani, Pilani Campus

288

Agenda

Building with pipelines

- Continuous integration
- Tooling
- Repository patterns – Multi-repo, mono-repo



Building with pipelines



BITS Pilani
Pilani Campus

Prof. Akanksha Bharadwaj
Asst. Professor, CSIS Department

285

BITS Pilani, Pilani Campus

287

Continuous Integration



- Continuous integration is a DevOps software development practice where developers regularly merge their code changes into a central repository, after which automated builds and tests are run.

- As part of this process, we often create artifact(s) that are used for further validation, such as deploying a running service to run tests against it.

- To enable the artifacts to be reused, we place them in a repository of some sort, either provided by the CI tool itself or on a separate system.

BITS Pilani, Pilani Campus

290

Key Principles of DevOps



- Collaboration & Communication
- Automation
- Continuous Integration (CI)
- Continuous Delivery (CD)
- Monitoring & Feedback
- Security (DevSecOps)

BITS Pilani, Pilani Campus

289

Benefits of Continuous Integration



- Early Bug Detection
- Faster Release Cycles
- Improved Code Quality
- Reduced Integration Risks
- Efficient Collaboration

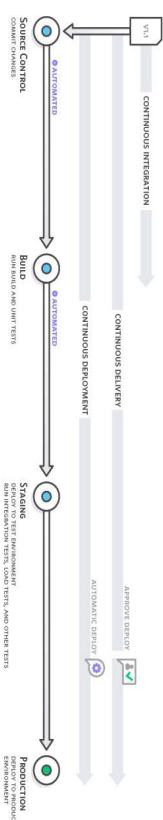
BITS Pilani, Pilani Campus

292

CI Workflow



- Code Commit
- Automated Build
- Automated Testing
- Feedback and Notifications
- Integration into Main Branch



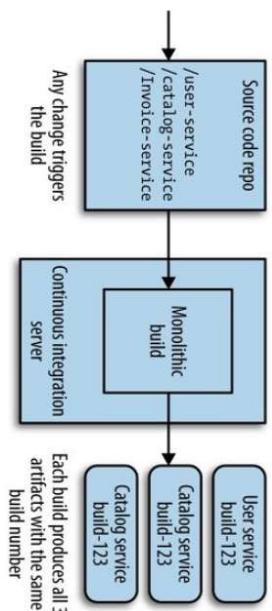
BITS Pilani, Pilani Campus

291

Mapping Continuous Integration to Microservices



- The simplest option is we could lump everything in together. We will have a single, giant repository storing all our code, and have one single build.

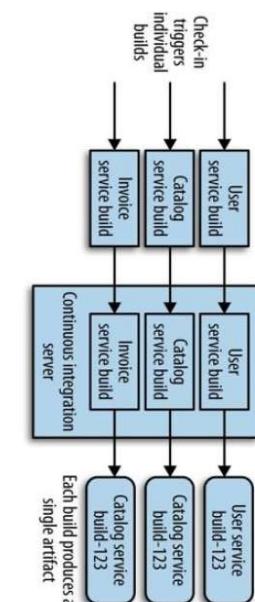


BITS Pilani, Pilani Campus
294

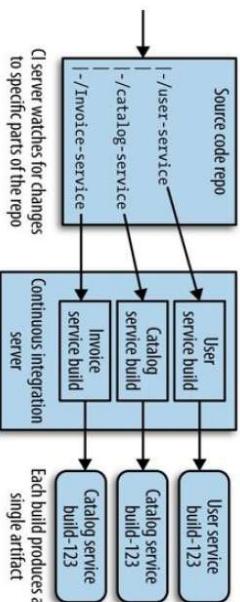
Continuous Integration for Microservices



- Continuous Integration (CI) in microservices ensures that each service is built, tested, and integrated frequently while maintaining **independence and scalability**.



BITS Pilani, Pilani Campus
295



BITS Pilani, Pilani Campus
293

Continuous Delivery in Microservices



- Continuous Delivery (CD) is the practice of **automating the deployment process** so that software can be released at any time with minimal manual intervention.
- In a **microservices** architecture, each service is independently built, tested, and deployed, enabling faster and more flexible software delivery.

BITS Pilani, Pilani Campus
298

Key Considerations for CI in Microservices



- Independent Pipelines
- API Contracts & Testing
- Parallel Testing
- Dependency Management

BITS Pilani, Pilani Campus
300

CD Pipeline for Microservices

A typical **Continuous Delivery pipeline** for microservices consists of:

- Continuous Integration (CI)
- Artifact Storage
- Deployment to Staging
- Automated Deployment Strategies
- Monitoring & Feedback Loop

Best Practices for Continuous Delivery in Microservices



- Independent Deployments
- Versioning & Rollbacks
- Deployment Automation
- Monitoring & Logging

BITS Pilani, Pilani Campus
300

BITS Pilani, Pilani Campus
297

Deployment Strategies in Microservices

- Rolling Update: Gradually replaces old instances with new ones.
- Blue-Green: Deploys to a parallel environment, then switches traffic.
- Canary Release: Releases to a small % of users before full rollout.
- Feature Flags: Enables/disables features dynamically without redeploying.

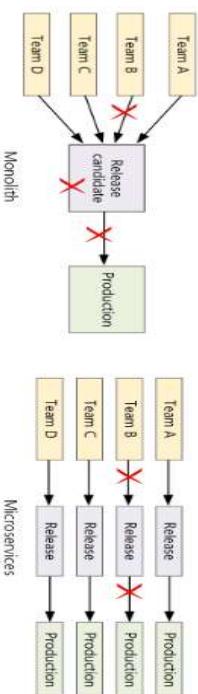


CI/CD pipeline for Monolith Vs Microservices

- Monolithic
 - Single build pipeline
 - Bug fixes impacts release

Microservice

- One pipeline per service
 - Issue in one service doesn't impact the other service



302

304



Continuous Deployment in Microservices

- Continuous Deployment takes automation a step further by automatically deploying every code change that passes automated tests directly to the production environment without manual intervention.



Pipeline

- Continuous delivery is the approach whereby we get constant feedback on the production readiness of each and every check-in, and furthermore treat each and every check-in as a release candidate



Figure 6-4. A standard release process modeled as a build pipeline

- In a microservices world, where we want to ensure we can release our services independently of each other, it follows that as with CI, we'll want one pipeline per service.

BITS Pilani, Pilani Campus

301

BITS Pilani, Pilani Campus

303

Service Configuration



- If we have some configuration for our service that does change from one environment to another, how should we handle this as part of our deployment process?

- One option is to build one artifact per environment, with configuration inside the artifact itself.
- A better approach is to create one single artifact, and manage configuration separately.

BITS Pilani, Pilani Campus
306

Platform-Specific Artifacts



- From the point of view of a microservice, depending on your technology stack, the artifact may not be enough by itself.

- While a Java JAR file can be made to be executable and run an embedded HTTP process, for things like Ruby and Python applications, you'll expect to use a process manager running inside Apache or Nginx.
- So we may need some way of installing and configuring other software that we need in order to deploy and launch our artifacts.
- This is where **automated configuration management** tools like Puppet and Chef can help

BITS Pilani, Pilani Campus
308

Mono-repo vs. Multi-repo



	Monorepo	Multiple repos
Advantages	Code sharing Easier to standardize code and tooling Easier to refactor code Discoverability - single view of the code	Clear ownership per team Potentially fewer merge conflicts Helps to enforce decoupling of microservices
Challenges	Changes to shared code can affect multiple microservices Greater potential for merge conflicts Tooling must scale to a large code base Access control More complex deployment process	Harder to share code Harder to enforce coding standards Dependency management Diffuse code base, poor discoverability Lack of shared infrastructure

Repository Strategies for Microservices

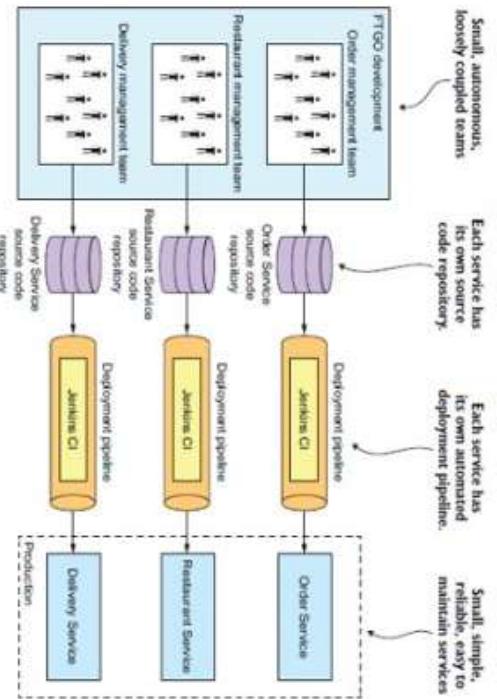


- Managing code repositories in a **microservices architecture** is crucial for efficient development, version control, and deployment.
- Types of Repository Structures for Microservices:
 - Monorepo (Single Repository for All Microservices): A single Git repository containing code for all microservices.
 - Multirepo (One Repository per Microservice): Each microservice has its own separate Git repository.

BITS Pilani, Pilani Campus
305

BITS Pilani, Pilani Campus
307

Pipeline for FTGO application



310
BITS Pilani, Pilani Campus

Deployment Challenges

- Many small independent code bases.
- Multiple languages and frameworks
- Release management
- Service updates
- Integration and load testing



References

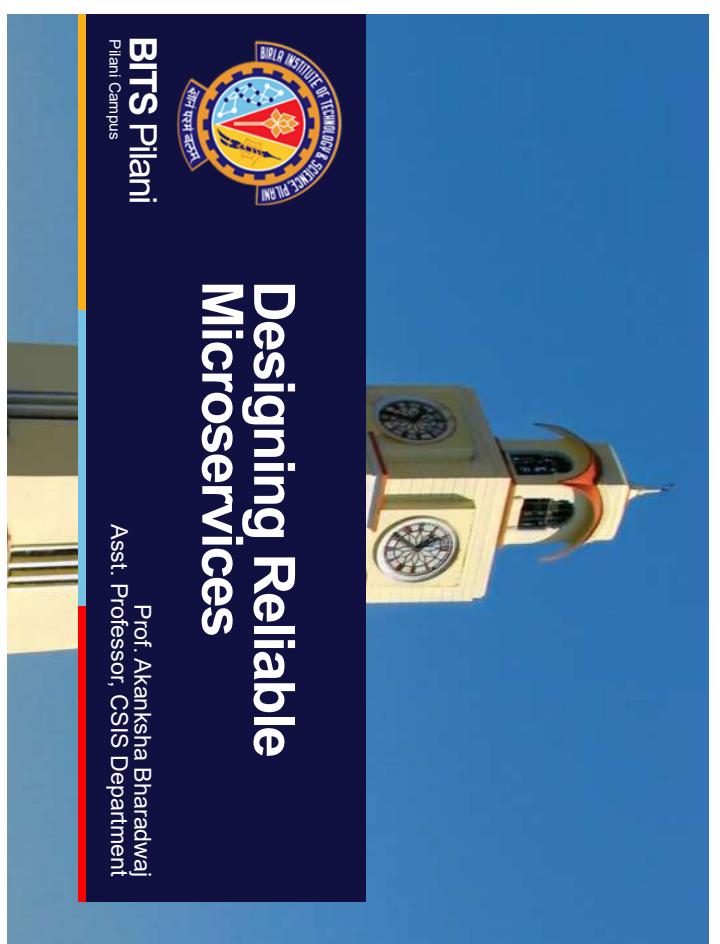
- Books as per the handout
- <https://aws.amazon.com/devops/continuous-integration/>
- <https://docs.microsoft.com/en-us/azure/architecture/microservices/ci-cd>
- <https://medium.com/@dmosyan/ci-cd-for-microservices-1b5582f3e1fd>



Designing Reliable Microservices



Asst. Professor, CSIS Department
Prof. Akanksha Bharadwaj



312
BITS Pilani, Pilani Campus

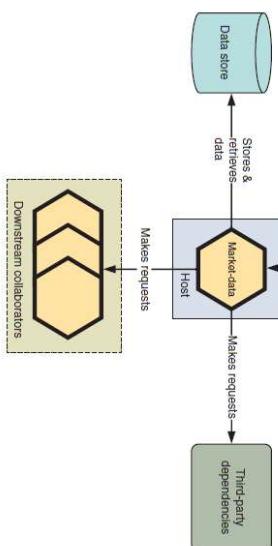
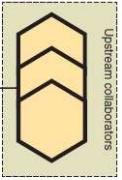
Agenda

Designing reliable Microservices

- Sources of failure, cascading failures
- Designing reliable communication: Retries, async. Comm., circuit breakers
- Maximizing service reliability: Load balancing, Rate limiting (Queues, Throttling)
- Service mesh



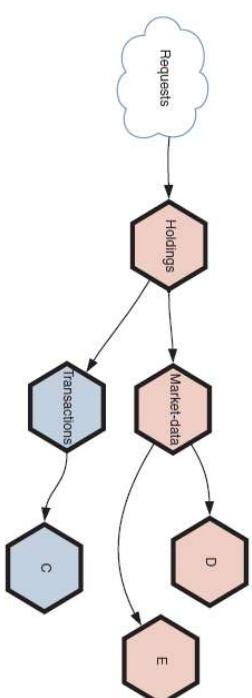
Example



BITS Pilani, Pilani Campus
314

What is reliability?

- A service's availability is a measure of how reliable you can expect it to be.



SE ZG583, Scalable Services Lecture No. 10

Hardware Failure



Source of failure	Frequency	Description
Host	Often	Individual hosts (physical or virtual) may fail.
Data center	Rare	Data centers or components within them may fail.
Host configuration	Occasionally	Hosts may be misconfigured—for example, through errors in provisioning tools.
Physical network	Rare	Physical networking (within or between data centers) may fail.
Operating system and resource isolation	Occasionally	The OS or the isolation system—for example, Docker—may fail to operate correctly.

Sources of failure



Every point of interaction between your service and another component indicates a possible area of failure. Failures could occur in four major areas:

- *Hardware*
- *Communication*
- *Dependencies*
- *Internal*

Dependency-related failure



Source of failure	Description
Timeouts	Requests to services may time out, resulting in erroneous behavior.
Decommissioned or nonbackwards-compatible functionality	Design doesn't take service dependencies into account, unexpectedly changing or removing functionality.
Internal component failures	Problems with databases or caches prevent services from working correctly.
External dependencies	Services may have dependencies outside of the application that don't work correctly or as expected—for example, third-party APIs.

Communication Failure



Source of failure	Description
Network	Network connectivity may not be possible.
Firewall	Configuration management can set security rules inappropriately.
DNS errors	Hostnames may not be correctly propagated or resolved across an application.
Messaging	Messaging systems—for example, RPC—can fail.
Inadequate health checks	Health checks may not adequately represent instance state, causing requests to be routed to broken instances.

Cascading failures

- In a microservice application, overload can cause a domino effect: failure in one service increases failure in upstream services, and in turn their upstream services. At worst, the result is widespread unavailability



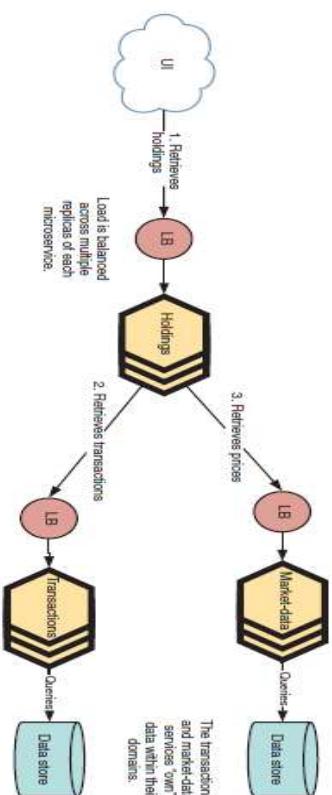
Service Practices

- Inadequate or limited engineering practices when developing and deploying services may lead to failure in production
- Because each service contributes to the effectiveness of the whole system, one poor quality service can have a detrimental effect on the availability of swathes of functionality.

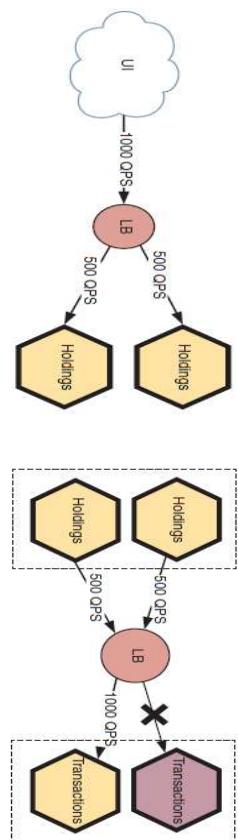


BITS Pilani, Pilani Campus
322

Example: cascading failure



BITS Pilani, Pilani Campus
324



BITS Pilani, Pilani Campus
321

Designing reliable communication

Several techniques for ensuring that services behave appropriately

- Retries
- Fallbacks, caching, and graceful degradation
- Timeouts and deadlines
- Circuit breakers
- Communication brokers



Retries

- How can you use retries to improve your resiliency in the face of intermittent failures without contributing to wider system failure if persistent failures occur?

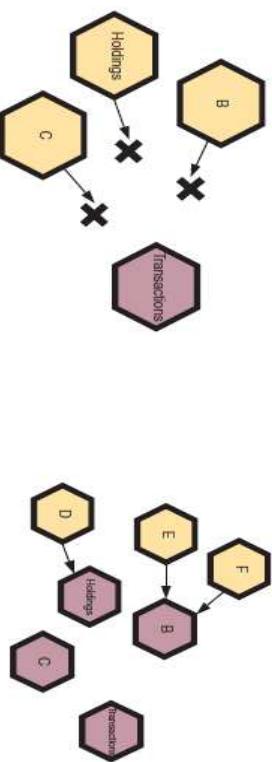
BITS Pilani, Pilani Campus
326



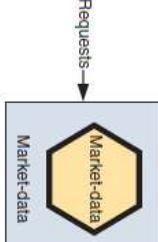
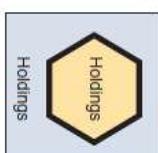
Retries

Upstream dependencies are unable to service requests that rely on transactions.

Increased failure in upstream dependencies leads to retries, repeating the cycle of failure.



- Imagine that a call from the holdings service to retrieve prices fails, returning an error.
- From the perspective of the calling service, it's not clear yet whether this failure is isolated
- Repeating that call is likely to succeed or systemic i.e. the next call has a high likelihood of failing.



BITS Pilani, Pilani Campus
328



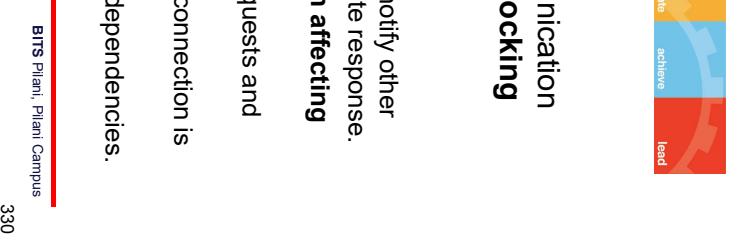
BITS Pilani, Pilani Campus
325

Asynchronous communication

- It's recommended to use asynchronous communication when microservices need to interact **without blocking execution**.

Use Cases:

- **Event-Driven Architectures:** When a service needs to notify other services about an event but does not require an immediate response.
- When you want to **prevent a failure in one service from affecting others**.
- When multiple microservices handle a high volume of requests and blocking calls can degrade performance.
- When an operation takes too long and holding an HTTP connection is impractical.
- When multiple services must consume an event without dependencies.



330

Retries: few points to remember

- Always limit the total number of retries.
- Use exponential back-off with jitter to smoothly distribute retry requests and avoid compounding load.
- Consider which error conditions should trigger a retry and, therefore, which retries are unlikely to, or will never, succeed.



331

Circuit breakers

- A circuit breaker is a pattern for pausing requests to a failing service to prevent cascading failures.

Approach

- When making a request to a service, you can track the number of times that request succeeds or fails
- In case of normal operation, we consider the circuit to be closed.
 - If the number of failures seen or the rate of failures within a certain time window passes a threshold, then the circuit is opened



332



Rate limits

- Technique used to control the number of requests a client can make to a service within a specific time period.



Maximizing service reliability

334

Why Use Circuit Breakers?

- Prevent System Overload:** Stops excessive requests to a failing service, reducing unnecessary resource consumption.

- Improve System Stability:** Ensures that failures in one microservice don't bring down the entire system.

- Enable Fast Failure & Recovery:** Redirects traffic or returns fallback responses instead of waiting for a failed service.

Load Balancer

The load balancer plays two important roles:

- Health Check:** Identifying which underlying instances are healthy and able to serve requests
- Routing requests** to different underlying instances of the service

We can classify health checks based on two criteria:

- Liveness:** check that the application has started and is running correctly.
- Readiness:** indicates whether a service is ready to serve traffic



336

Maximizing service reliability

334

Why Use Circuit Breakers?

- Prevent System Overload:** Stops excessive requests to a failing service, reducing unnecessary resource consumption.

- Improve System Stability:** Ensures that failures in one microservice don't bring down the entire system.

- Enable Fast Failure & Recovery:** Redirects traffic or returns fallback responses instead of waiting for a failed service.

Load Balancer

The load balancer plays two important roles:

- Health Check:** Identifying which underlying instances are healthy and able to serve requests
- Routing requests** to different underlying instances of the service

We can classify health checks based on two criteria:

- Liveness:** check that the application has started and is running correctly.
- Readiness:** indicates whether a service is ready to serve traffic



336

Maximizing service reliability

334

Why Use Circuit Breakers?

- Prevent System Overload:** Stops excessive requests to a failing service, reducing unnecessary resource consumption.

- Improve System Stability:** Ensures that failures in one microservice don't bring down the entire system.

- Enable Fast Failure & Recovery:** Redirects traffic or returns fallback responses instead of waiting for a failed service.

Load Balancer

The load balancer plays two important roles:

- Health Check:** Identifying which underlying instances are healthy and able to serve requests
- Routing requests** to different underlying instances of the service

We can classify health checks based on two criteria:

- Liveness:** check that the application has started and is running correctly.
- Readiness:** indicates whether a service is ready to serve traffic



336

Example Use Cases



✓ Protecting Authentication Services

Limit failed login attempts (e.g., **5 attempts per minute** to prevent brute force attacks).

✗ Ensuring Fair API Usage

Free-tier users can make **100 API calls per hour**, while premium users get **1000 API calls per hour**.

✗ Preventing Traffic Spikes

A content delivery API limits requests to **200 per second** to avoid overwhelming backend servers.

- ◆ **API Gateway-Based Throttling:** Use **AWS API Gateway**, Kong, NGINX, or Azure API Management to throttle requests per client or user.
- ◆ **Application-Level Throttling:** Implement in code using Redis + Express Rate Limit (Node.js) or Spring Boot Bucket4j (Java).
- ◆ **Message Queue-Based Throttling:** Use **Kafka**, RabbitMQ, or **SQS** to control the rate at which requests are processed asynchronously.
- ◆ **Load Balancer-Based Throttling:** Use **AWS ALB**, Nginx, or HAProxy to distribute and throttle traffic before reaching backend services.

BITS Pilani, Pilani Campus
338

Rate limiting in Microservices



- API Gateway-based Rate Limiting
- Middleware-based Rate Limiting
- Cloud Provider Solutions

BITS Pilani, Pilani Campus
340

Throttling Pattern



- **Throttling** is a technique used to **control the rate at which requests are processed** to prevent system overload and ensure fair resource distribution.

Types of Throttling

- User-Based Throttling
- Server-Based Throttling
- Network-Based Throttling
- IP-Based Throttling

339

How to Implement Throttling?



◆ API Gateway-Based Throttling: Use AWS API

Gateway, Kong, NGINX, or Azure API Management to throttle requests per client or user.

◆ Application-Level Throttling: Implement in code using

Redis + Express Rate Limit (Node.js) or Spring Boot Bucket4j (Java).

◆ Message Queue-Based Throttling: Use Kafka, RabbitMQ, or SQS to control the rate at which requests

are processed asynchronously.

◆ Load Balancer-Based Throttling: Use AWS ALB,

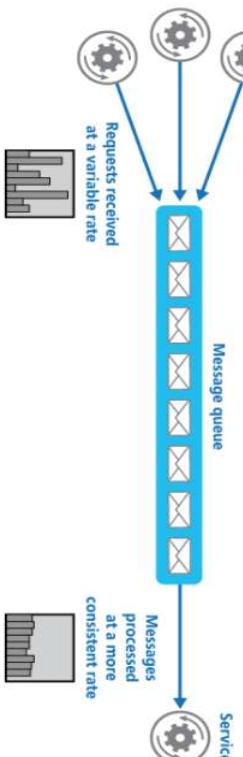
Nginx, or HAProxy to distribute and throttle traffic before reaching backend services.

BITS Pilani, Pilani Campus
340

BITS Pilani, Pilani Campus
337

How the Pattern Works?

Solution:



BITS Pilani, Pilani Campus

342

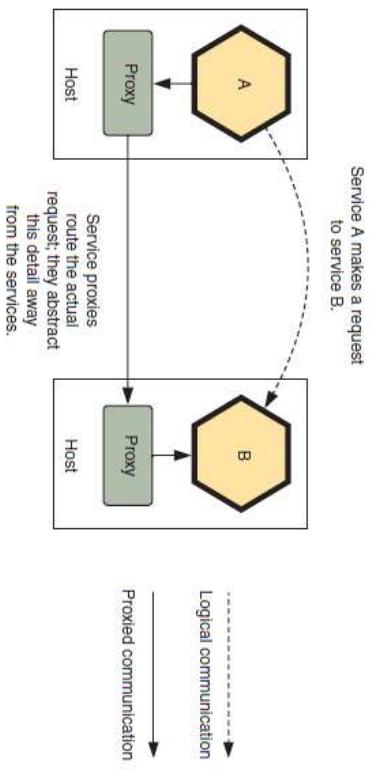
Queue-Based Load Levelling pattern

- The **Queue-Based Load Levelling Pattern** is a **scalability and resilience pattern** used in microservices to handle **high traffic spikes** by placing requests in a queue.
- This prevents services from being overwhelmed and allows them to process tasks **asynchronously at their own pace**.

Service Mesh

- A service mesh is a dedicated infrastructure layer designed to handle service-to-service communication within a network of microservices.

- It manages and controls how different parts of an application or service interact with each other.



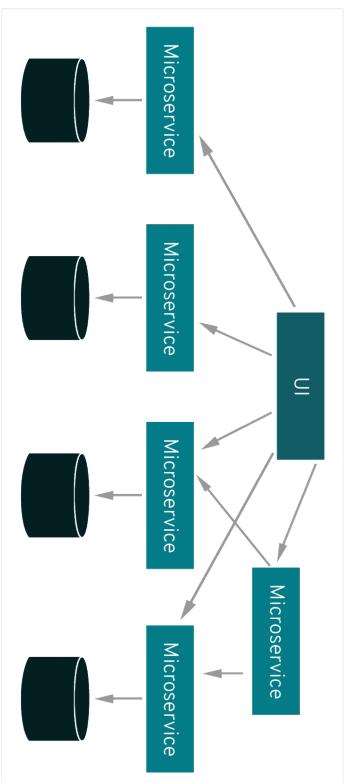
BITS Pilani, Pilani Campus

343

BITS Pilani, Pilani Campus

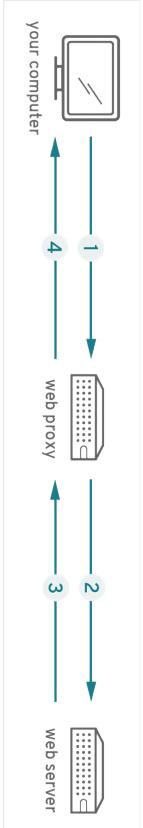
341

Service Mesh and Microservices



BTS Plan! Plan Campus
346

How does service mesh work?

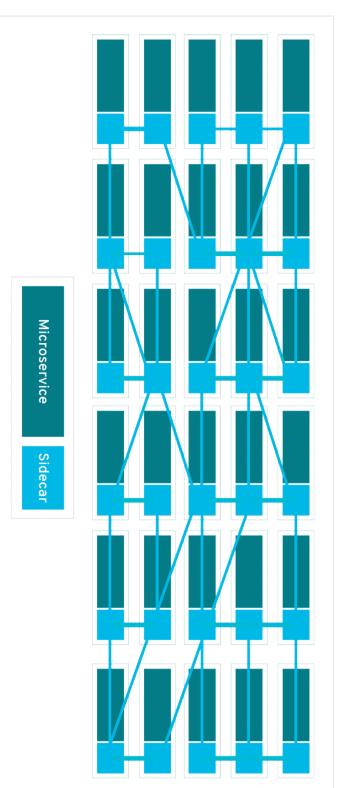


346

- In a service mesh, requests are routed between microservices through proxies in their own infrastructure layer.
- For this reason, individual proxies that make up a service mesh are sometimes called “sidecars,” since they run alongside each service, rather than within them.
- Taken together, these “sidecar” proxies—decoupled from each service—form a mesh network.
- Without a service mesh, each microservice needs to be coded with logic to govern service-to-service communication, which means developers are less focused on business goals.

BTS Plan! Plan Campus
348

Service Mesh and Microservices



348

- A service mesh is built into an app as an array of network proxies.
- As your request for this page went out, it was first received by your company's web proxy
- After passing the proxy's security measure, it was sent to the server that hosts this page
- Next, this page was returned to the proxy and again checked against its security measures
- And then it was finally sent from the proxy to you.

345

BTS Plan! Plan Campus

347

Self study

- <https://cloud.google.com/architecture/rate-limiting-strategies-techniques>
- <https://aws.amazon.com/blogs/architecture/chaos-testing-with-aws-fault-injection-simulator-and-aws-codepipeline/>
- <https://learn.microsoft.com/en-us/azure/architecture/patterns/sidecar>



Securing and Testing scalable services

BITS Pilani
Pilani Campus

Akanksha Bharadwaj
Asst. Professor, CSE Department



Does service mesh optimize communication?



350

References



352

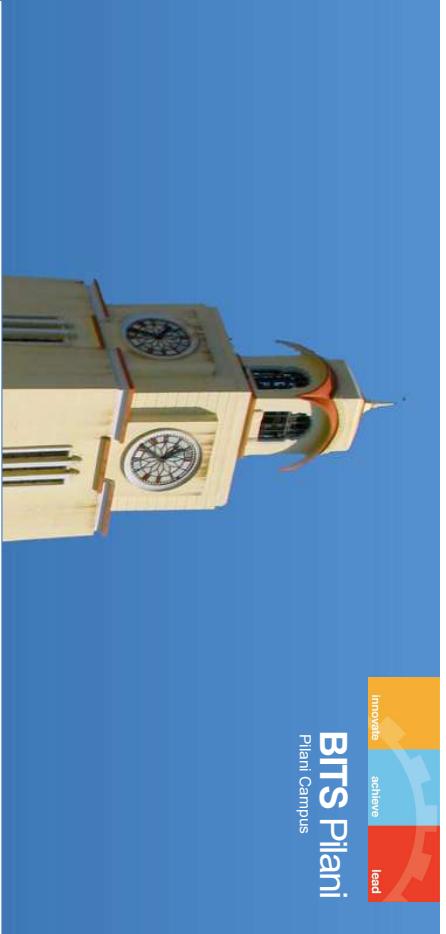
- Within a complex microservices architecture, it can become nearly impossible to locate where problems have occurred without a service mesh.
- Over time, data made visible by the service mesh can be applied to the rules for interservice communication, resulting in more efficient and reliable service requests.

Agenda



Securing and Testing scalable services

- Securing code and repositories
- Using Authentication and Authorization
- Testing



BITS Pilani, Pilani Campus
354

Measures for securing the code and repository



- Choose a repository you trust
- Consider the exposure of your repository
- Protect access credentials
- Access to the repository should be revoked when no longer required, or in the event of compromise

- Encryption
- Code Signing
- Ensure your code is backed up
- Auditing and Logging

Securing code and repository



- Your code is only as secure as the systems used to create it.
- Version control, peer review and built-in auditing are some of the advantages which come with using a code repository.
- If proper attention is paid to security measures, the benefits of using a repository far outweigh the risks.

SE ZG583, Scalable Services

Lecture No. 11

Authentication implementation options

- One option is for the individual services to authenticate the user.

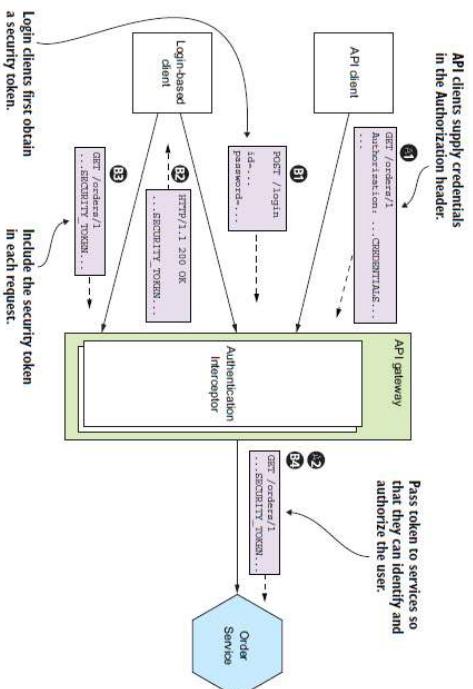
- Other approach is for the API gateway to authenticate a request before forwarding it to the services.



What is Authentication?

- Verifying the identity of the application or human that's attempting to access the application.

Authentication in API Gateway



BITS Pilani, Pilani Campus
358



The sequence of events for API clients is as follows:

- A client makes a request containing credentials.
- The API gateway authenticates the credentials.
- It creates a security token, and passes that to the service or services.



BITS Pilani, Pilani Campus
360



BITS Pilani, Pilani Campus
357

BITS Pilani, Pilani Campus
359

Authorization



Key aspects and benefits of implementing authorization in an API gateway

- Verifying that the principal is allowed to perform the requested operation on the specified data.
- Applications often use a combination of rolebased security and access control lists (ACLs).

Explore: What are the possible drawbacks of this approach?

BITS Pilani, Pilani Campus
362



Authorization in API Gateway

The sequence of events for login-based clients is as follows:

- A client makes a login request containing credentials.
- The API gateway returns a security token.
- The client includes the security token in requests that invoke operations.
- The API gateway validates the security token and forwards it to the service or services.

BITS Pilani, Pilani Campus
364



- Centralized Access Control
- Scalability and Performance
- Flexibility
- Logging and Auditing
- Protection Against Attacks

BITS Pilani, Pilani Campus
361



Key aspects and benefits of implementing authorization within services

- Fine-Grained Control
- Decentralization
- Service Autonomy
- Scalability
- Flexibility
- Security Isolation



What is OAuth?

- OAuth is an open-standard authorization protocol or framework
- The simplest example of OAuth is when you go to log onto a website and it offers one or more opportunities to log on using another website's/service's logon.

- **Explore:** What are the possible drawbacks of this approach?

BITS Pilani, Pilani Campus
366

Authorization in Services



- The other place to implement authorization is in the services.
- A service can implement role-based authorization for URLs and for service methods.
- It can also implement ACLs to manage access to aggregates.

Which approach is better?

The preferred approach depends on the specific requirements and characteristics of the system:

- **Centralized Control:** If uniformity and centralized management of access control policies are paramount, implementing authorization in an API gateway might be preferred.
- **Fine-Grained Control:** For systems with diverse access control requirements or where services have unique authorization needs, implementing authorization within services may be more suitable.
- **Scalability and Performance:** If scalability and performance are critical concerns, especially in high-traffic environments, offloading authorization to an API gateway might be preferable.
- **Security Isolation:** If security isolation is a priority and potential single points of failure are a concern, embedding authorization within services could be favored.
- In many cases, a hybrid approach may be appropriate, combining elements of both centralized and decentralized authorization mechanisms to achieve a balance between control, scalability, performance, and security.

BITS Pilani, Pilani Campus
365

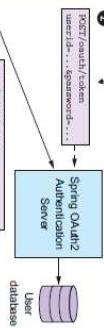


BITS Pilani, Pilani Campus
368

How OAuth works for Microservices



Password grant request



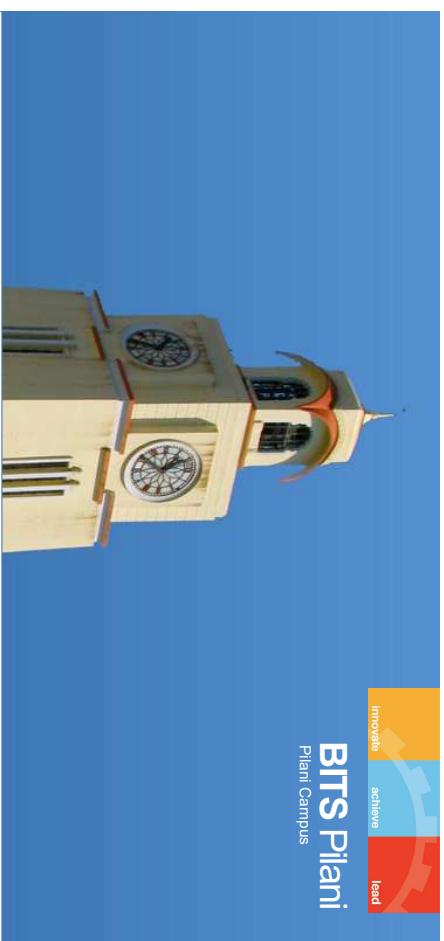
Contains the user ID and their roles

370

Terms used in OAuth

- Authorization Server
- Access Token
- Refresh Token
- Resource Server
- Client

Testing in Microservices

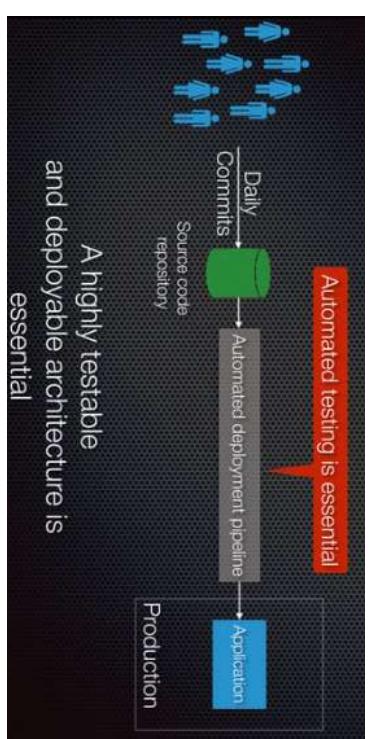


BITS Pilani
Pilani Campus

Use DevOps



Automated testing is essential

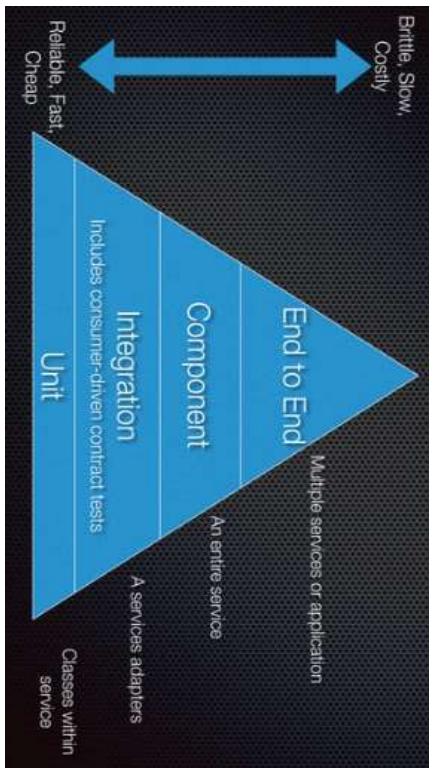


372

<https://microservices.io/testing/>

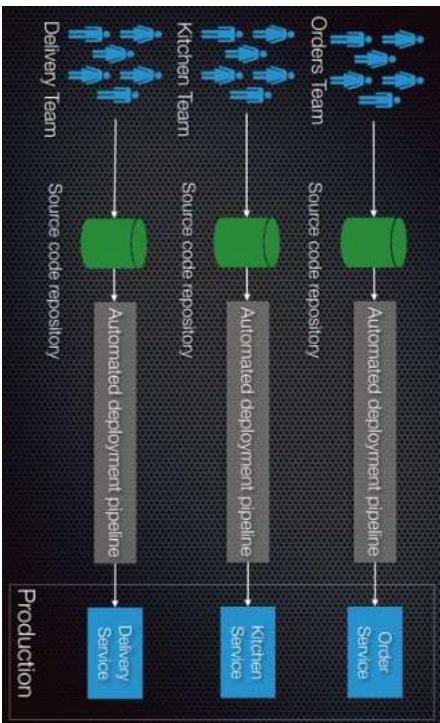
BITS Pilani, Pilani Campus
369

Testing Pyramid and Microservices



BITS Pilani, Pilani Campus
374

Autonomous service teams



BITS Pilani, Pilani Campus
373

Unit Tests



- The size of the unit under test is not defined anywhere. In some firms, consider writing them at a class level or at a group of related classes.

- Consider keeping the testing units as small as possible. It becomes easier to express the behavior if the test is small since the branching complexity of the unit is lower.

Testing strategies for Microservice architectures



- Unit Testing
- Component Testing
- Integration Testing
- Contract Testing
- End-to-End testing
- Fault Injection Testing
- Monitoring and Observability

BITS Pilani, Pilani Campus
376

Component Testing

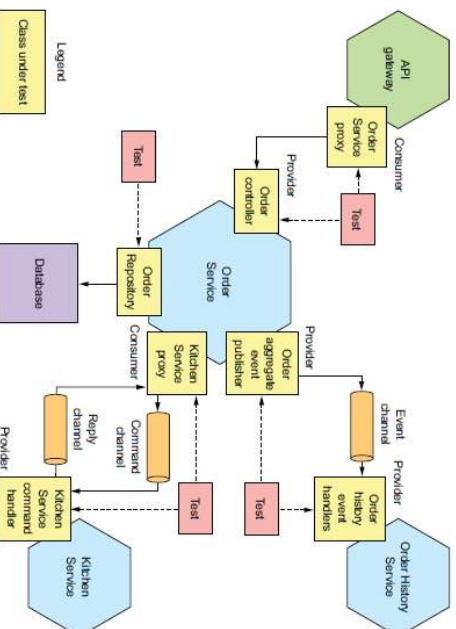


- A component or microservice is a well-defined coherent and independently replaceable part of a larger system. Once we execute unit tests of all the functions within microservices, it's time to test microservice itself in isolation.
- To test a single microservice in isolation, we need to mock the other microservices. Isolating the microservices in this way using test doubles avoids any complex behavior they exert on execution.

BITS Pilani, Pilani Campus
378

Types of Unit Test

- *Solitary unit* — looks at the interactions and collaborations between an object and its dependencies, which are replaced by the test doubles.
- *Sociable unit test*— focuses on testing the behavior of the module by observing changes in their state.



Approaches for integration testing



- Service Integration Testing
- Testing Service Contracts
- Mocking External Dependencies
- End-to-End Testing

BITS Pilani, Pilani Campus
380

Integration Test

BITS Pilani, Pilani Campus
377

Contract testing

- For simplifying integration tests that verify interactions between application services is to use contracts

Interaction style	Consumer	Provider	Contract
REST-based, request/response	API Gateway	order Service	HTTP request and response
Publish/subscribe	Order History service	Order Service	Domain event
Asynchronous request/response	Order Service	Kitchen Service	Command message and reply message

BITS Pilani, Pilani Campus

382

Consumer Driven Contact testing

Verify that a service and its clients can communicate while testing them in isolation



BITS Pilani, Pilani Campus

384

End-to-End Testing

- We usually treat the system as a black box while performing end-to-end tests.

- As Microservice includes more moving parts for the same behavior, it provides coverage of the gaps between the services.



- Consumer-side contract tests:** These are tests for the consumer's adapter. They use the contracts to configure stubs that simulate the provider, enabling you to write integration tests for a consumer that don't require a running provider.
- Provider-side contract tests:** These are tests for the provider's adapter. They use the contracts to test the adapters using mocks for the adapters's dependencies.

BITS Pilani, Pilani Campus

381

BITS Pilani, Pilani Campus

383

Fault injection testing

- Fault injection testing, also known as chaos testing or failure testing, is a technique used to validate the resilience and fault tolerance of systems by deliberately introducing faults or failures into the system and observing how it responds.



BITS Pilani, Pilani Campus
386

Chaos Engineering At Netflix

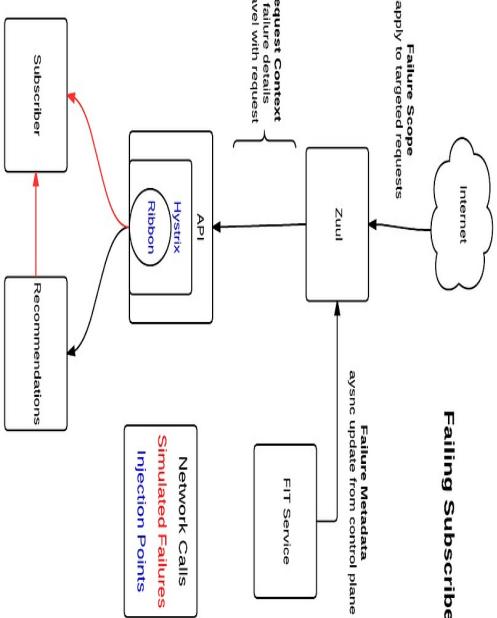
- JUnit (for Java), Pytest (for Python), JUnit (for .NET): These are popular unit testing frameworks for writing and executing unit tests for individual components, functions, or classes within microservices.
- Mockito, Mockito.NET: These mocking frameworks help simulate dependencies and external systems during unit testing by creating mock objects or stubs.



BITS Pilani, Pilani Campus
388

There are various best practices to adhere to while writing efficient E2E tests for microservices:

- Define end-to-end test scenarios that cover critical user journeys or business processes across multiple Microservices.
- To write automated tests that are simple to execute and manage, testing team should make use of testing frameworks and tools.
- Developers should use a continuous integration and delivery (CI/CD) strategy, in which the tests are incorporated into the development cycle and executed automatically anytime the application is changed or released.



BITS Pilani, Pilani Campus
385

- Fault injection testing, also known as chaos testing or failure testing, is a technique used to validate the resilience and fault tolerance of systems by deliberately introducing faults or failures into the system and observing how it responds.



BITS Pilani, Pilani Campus
386

Tools for End-to-End Testing



References

- **Selenium:** Selenium is a widely used tool for automating web browser interactions and testing web applications. It allows you to create automated end-to-end tests that simulate user interactions across multiple microservices.
- **Cypress:** Cypress is a modern testing tool for web applications that provides an elegant API for writing end-to-end tests. It offers features such as real-time testing, automatic waiting, and debugging capabilities.

BITS Pilani, Pilani Campus
390

Tools for Integration Testing



- **Postman:** Postman is a popular API testing tool that allows you to create and execute tests for RESTful APIs. It supports automated testing, scripting, and validation of API responses.
- **SoapUI:** SoapUI is another API testing tool that specializes in testing SOAP and RESTful web services. It provides features for functional testing, security testing, and performance testing of APIs.
- **WireMock:** WireMock is a lightweight HTTP server for stubbing and mocking HTTP responses. It is commonly used for simulating dependencies and external services during integration testing.

BITS Pilani, Pilani Campus
389

Self Study



- <https://docs.github.com/en/code-security/getting-started/securing-your-repository>
- OAuth in detail from textbook
 - Testing in detail from textbook
- <https://www.netvys.com/insights/microservices-load-testing>
- <https://github.io/chaosmonkey/>

BITS Pilani, Pilani Campus
392

Challenges in deploying Microservices



Installing microservices can be complex due to their distributed nature. Here are the key challenges:

- **Dependency Management:** Microservices often rely on multiple libraries, frameworks, and external services.
- **Configuration Complexity:** Each microservice may require its own configuration for databases, APIs, or environments (dev, staging, production).
- **Networking and Communication:** Microservices communicate over networks, introducing issues like latency, timeouts, and service discovery.
- **Containerization and Orchestration:** Most microservices use containers (e.g., Docker) and orchestration tools (e.g., Kubernetes).
- **Monitoring and Logging:** With many services running independently, tracking performance, errors, and logs becomes challenging.

Agenda

-
- Deploying Microservices
 - Service startup
 - Running multiple instances
 - Adding load balancer
 - Service to host models
 - Single Service Instance to Host
 - Multiple static Service Per Host
 - Multiple scheduled services per host



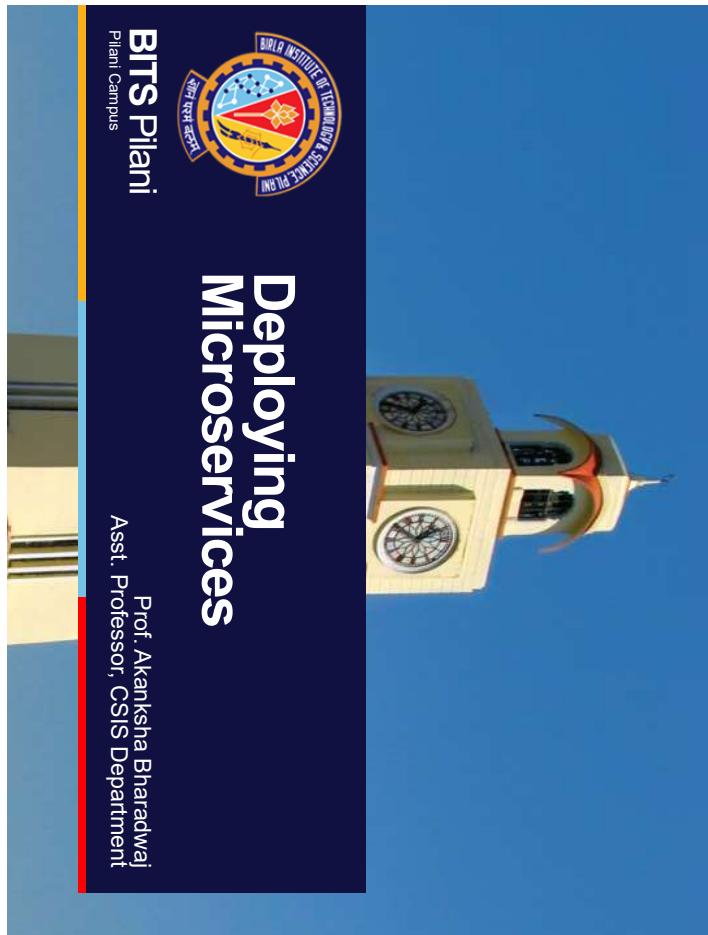
Deploying Microservices

BITS Pilani
Pilani Campus

Prof. Akanksha Bharadwaj
Asst. Professor, CSIS Department

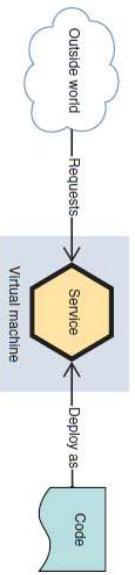


SE ZG583, Scalable Services Lecture No. 12



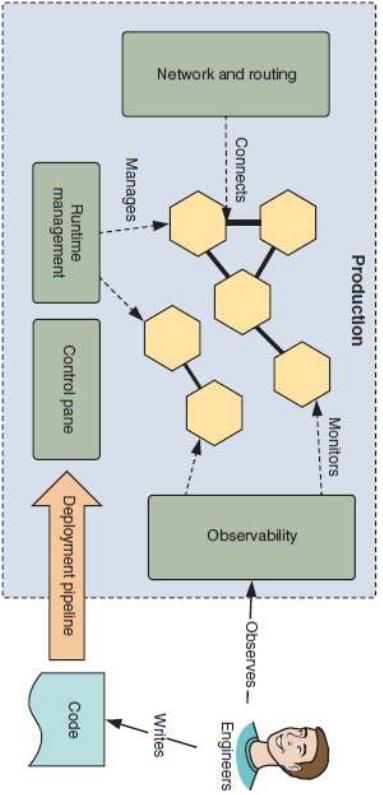
Service Startup

- You need to take your code, get it running on a virtual machine, and make it accessible from the outside world



BITS Pilani, Pilani Campus
398

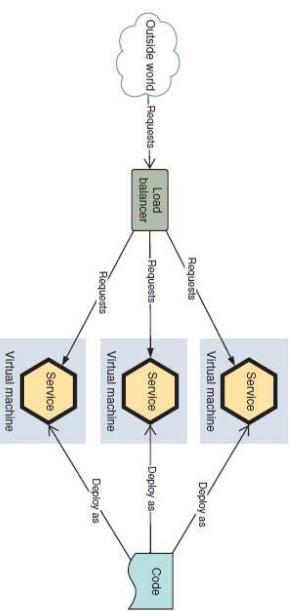
Features of a microservice production environment



BITS Pilani, Pilani Campus
397

Run multiple instances of a service

- Think about Scale horizontally



BITS Pilani, Pilani Campus
400

Service Startup

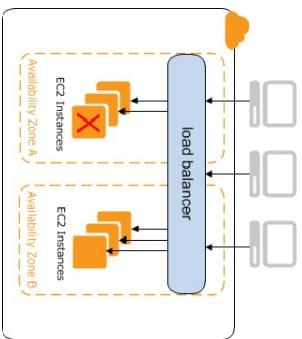
- Installs binary dependencies required to run the application
- Downloads your service code from Github or any other repository
- Installs that code's dependencies
- Configure a supervisor if required

BITS Pilani, Pilani Campus
399

AWS: What is a Classic Load Balancer?



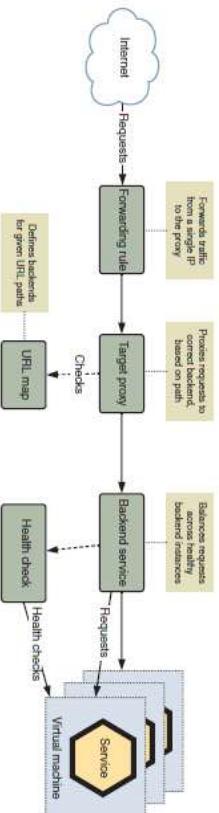
- Elastic Load Balancing automatically distributes your incoming traffic across multiple targets, such as EC2 instances, containers, and IP addresses, in one or more Availability Zones.



BITS Pilani, Pilani Campus
402

Adding a load balancer

- The load balancer uses routing rules, proxies, and maps to forward requests from the outside world to a set of healthy service instances.



Multiple Service Instances per Host Pattern



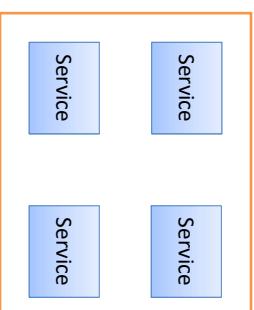
Deploying multiple services to a single host in production is synonymous with deploying multiple services to a local dev workstation or laptop.

Benefits:

- Attractive from a host management point of view
- Cost

Challenges:

- Monitoring
- Isolation of instances
- Resource consumption



BITS Pilani, Pilani Campus
404

Models for Deploying Services



Service Instance per Container Pattern

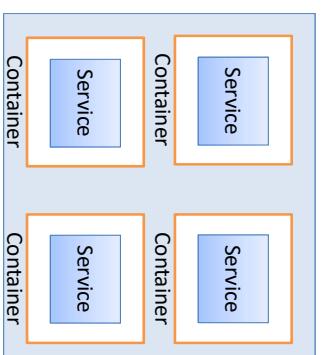
Package the service as a container image and deploy each service instance as a container

Benefits

- Easy to scale up and down a service.
- Isolation
- Limit the resource utilization by a service instance
- Faster deployment

Drawbacks

- Infrastructure for deploying containers is not rich



BITS Pilani, Pilani Campus
406

Single Service Instance per Host Pattern

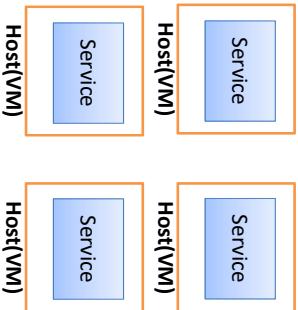
We can avoid the side effects of multiple services living on a single host, by making monitoring and remediation much simpler

Benefits

- Service instance Isolation
- Easy to manage, monitor and deploy
- No conflicting resource requirements

Drawbacks

- Less efficient resource utilization
- cost



BITS Pilani, Pilani Campus
405

Running a single service on local machine

- Open the service that you have created
- Use an appropriate command in the command prompt to run the service
- Use local browser/postman application to check the functioning

BITS Pilani, Pilani Campus
408

Best Deployment Strategy

3 main practical and popular approaches to deploy Microservices:

- Deploy each Microservice instance on a separate Virtual Machine

- Deploy each Microservice instance as a (Docker) container on Kubernetes
- Deploy each Microservice as a Serverless Function



407

Dockerfile best practices



- Pick the right base image
- Keep the number of steps in the Dockerfile to minimum
- Control the structure of your dockerfile

BITS Pilani, Pilani Campus
410

Dockerfile



- Docker can build images automatically by reading the instructions from a Dockerfile

Basic Docker commands:

- FROM
- RUN
- ENV
- COPY
- EXPOSE
- ENTRYPOINT
- CMD
- VOLUME
- WORKDIR
- LABEL
- ADD
- ARG

Run multiple services on a container



If you need to run more than one service within a container, you can accomplish this in a few different ways.

- Use wrapper script
- Use a process manager like supervisord

BITS Pilani, Pilani Campus
412

Docker Compose



- Compose is a tool for defining and running multi-container Docker applications.

Basic Structure of Docker compose file

- version
- services
- build
- command
- ports
- volumes
- links
- image
- environment
- restart
- depends_on

BITS Pilani, Pilani Campus
409

Self Study

- <https://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-getting-started.html>
- <https://docs.docker.com/>



Deploying Microservices Contd.

Prof. Akanksha Bharadwaj
Asst. Professor, CSIS Department

416

Run a single service using Docker desktop



Deploying Microservices Contd.

Prof. Akanksha Bharadwaj
Asst. Professor, CSIS Department

416

References

- Chapter8, Microservices in Action
- <https://docs.aws.amazon.com/elasticloadbalancing>



Agenda

Deploying microservices

- Deploying services without downtime: Canaries, Blue-Green, & rolling deploys
- Deploying microservices using Serverless deployment

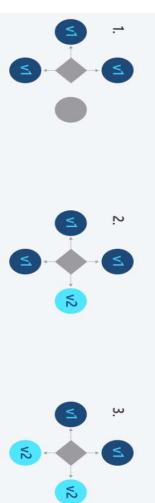
Deployment with Containers

- Introduction to containers
- Containerizing a service
- Deploying to a cluster



Ramped

- This is the default roll-out method in Kubernetes.
- This deployment slowly replaces pods one at a time to avoid downtime.



BITS Pilani, Pilani Campus
418

Introduction of deployment Strategies

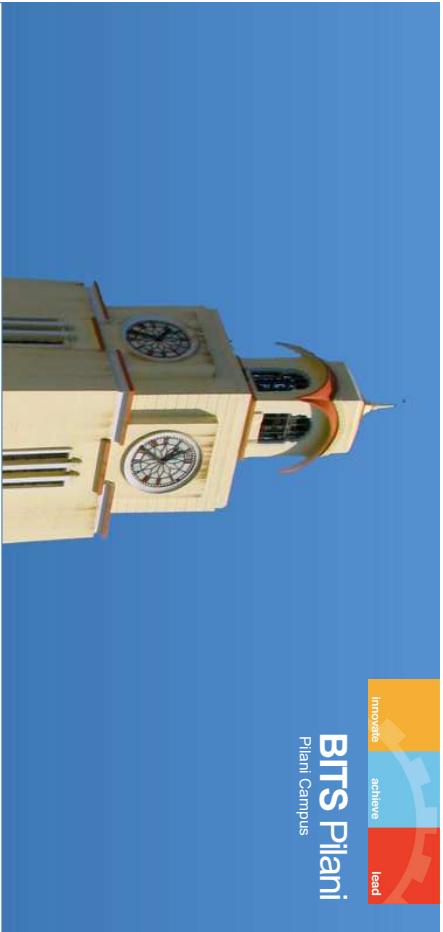
Here are four common deployment strategies that organizations use in production

- Ramped
- Blue/Green
- Canary
- A/B testing



SE ZG583, Scalable Services

Lecture No. 13



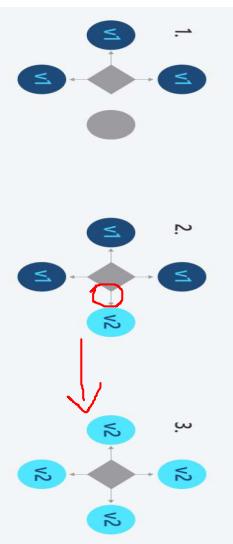
BITS Pilani
Pilani Campus

417

BITS Pilani, Pilani Campus
420

Canary

- Allow the customers to test your Kubernetes deployment by releasing the new version to a small group of them.



Serverless Deployment

- It is a deployment infrastructure that hides any concept of servers physical or virtual hosts, or containers.
- To deploy your service using this approach, you package the code (e.g. as a ZIP file), upload it to the deployment infrastructure and describe the desired performance characteristics.

BITS Pilani, Pilani Campus
422

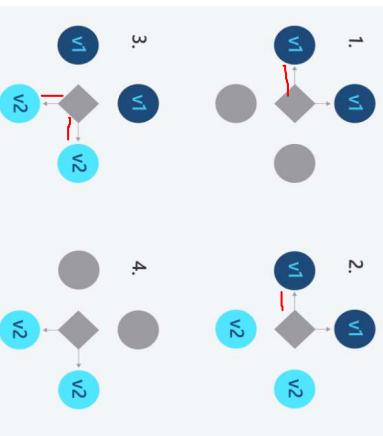
A/B testing

- Much like the canary Kubernetes deployment strategy, an A/B testing strategy targets a specific group of customers.



Blue/Green

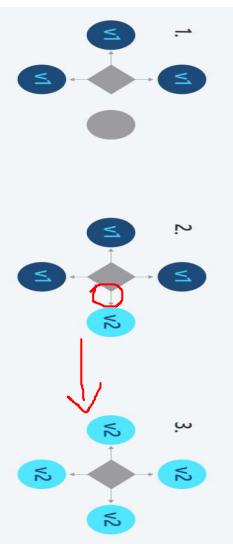
- In this you release a new version of your application or workflow while your current version is still running.



BITS Pilani, Pilani Campus
421

Canary

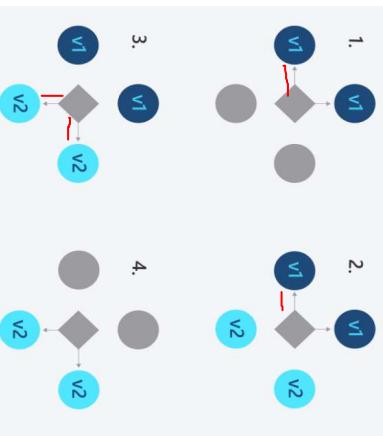
- Allow the customers to test your Kubernetes deployment by releasing the new version to a small group of them.



BITS Pilani, Pilani Campus
422

Blue/Green

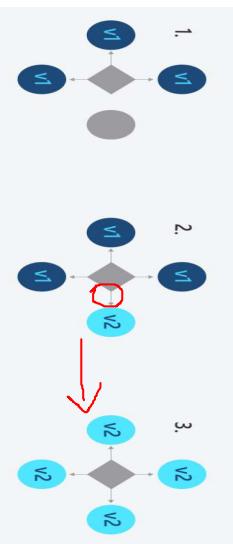
- In this you release a new version of your application or workflow while your current version is still running.



BITS Pilani, Pilani Campus
421

Canary

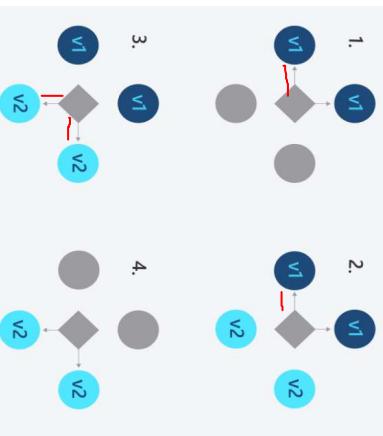
- Allow the customers to test your Kubernetes deployment by releasing the new version to a small group of them.



BITS Pilani, Pilani Campus
422

Blue/Green

- In this you release a new version of your application or workflow while your current version is still running.



BITS Pilani, Pilani Campus
421

Examples of Serverless Deployment

- AWS Lambda
- Google Cloud Functions
- Azure Functions



Container Image

- A container image is a lightweight, standalone, executable package that contains everything needed to run a piece of software, including the application code, runtime, libraries, dependencies, and configuration files.



Drawbacks of using Serverless Deployment

- It's not the best scenario for executing long-running applications.



BITS Pilani, Pilani Campus
426

Introduction to containers

- A container is a lightweight, portable, and self-contained unit of software that packages together an application and its dependencies, libraries, and configuration files, ensuring consistency and reliability across different computing environments.



BITS Pilani, Pilani Campus
428

- Cold start

- Containers provide a standardized runtime environment for applications, enabling them to run reliably and consistently on any platform that supports containerization.

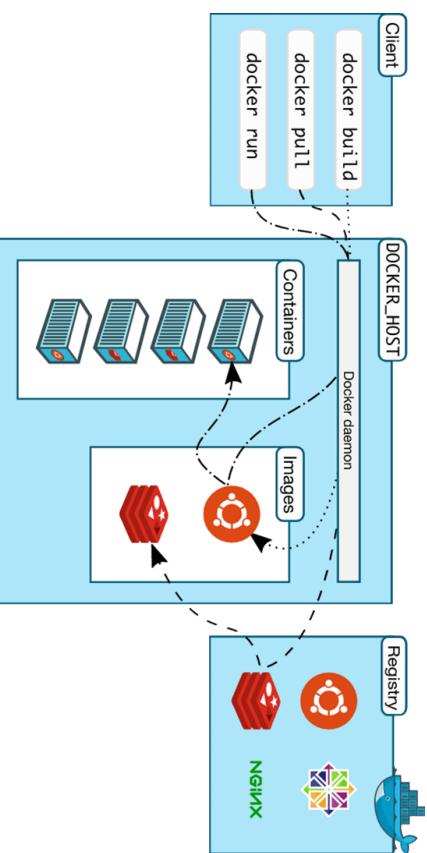
BITS Pilani, Pilani Campus
425

Key aspects of Docker containers

- Image-Based
- Isolation
- Portability
- Efficiency
- Orchestration
- Microservices
- DevOps



Docker architecture



What can I use Docker for?

- Fast and consistent delivery of your applications

- Microservices Architecture

- Scalability



Docker Containers

- Docker containers are lightweight, portable, and self-contained units of software that encapsulate an application along with its dependencies, libraries, and configuration files.

- Docker is a containerization platform that allows developers to create, deploy, and manage containers easily, providing a consistent runtime environment across different computing environments.
- Docker provides the ability to package and run an application in a loosely isolated environment called a container.

BITS Pilani, Pilani Campus
430

BITS Pilani, Pilani Campus
429

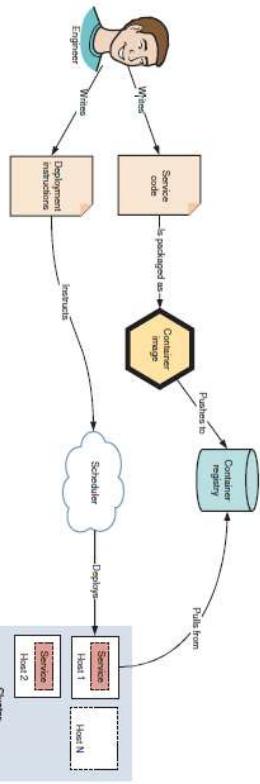
BITS Pilani, Pilani Campus
432

Build and Run Docker Image



Steps required:

1. Install Docker
2. Build your Docker image
3. Run your Docker image



BITS Pilani, Pilani Campus
434

Containerizing a service

- Build an image for a service
- Run multiple instances or containers of your image
- Push your image to a shared repository, or registry

Building your image



- To build image, first you need to create a Dockerfile

• Example: docker build -t market-data:first-build



Working with images

- The image will include the file system that your application needs to run code and dependencies and other metadata, such as the command that starts your application.

- Example: docker pull python:3.6

BITS Pilani, Pilani Campus
436

BITS Pilani, Pilani Campus
433

BITS Pilani, Pilani Campus
435

Storing an image

- To push to this repository, you need to tag your image with an appropriate name.
- Docker image names follow the format `<registry>/<repository>:<tag>`.
- After deleting old container, start a new container using the updated code.



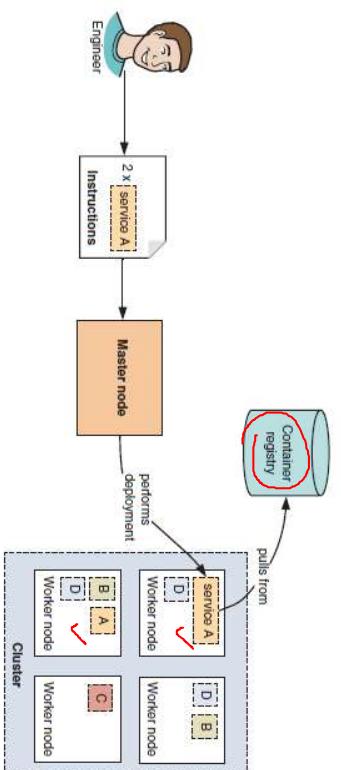
Update the application

Steps

- Update the source code
- Build our updated version of the image, using the same command we used before.
`docker build -t getting-started:v2`.
- After deleting old container, start a new container using the updated code.
`docker run -dp 3000:3000 getting-started:v2`

BITS Pilani, Pilani Campus
438

Deploying to a cluster



BITS Pilani, Pilani Campus
440

High-level scheduler architecture and deployment process

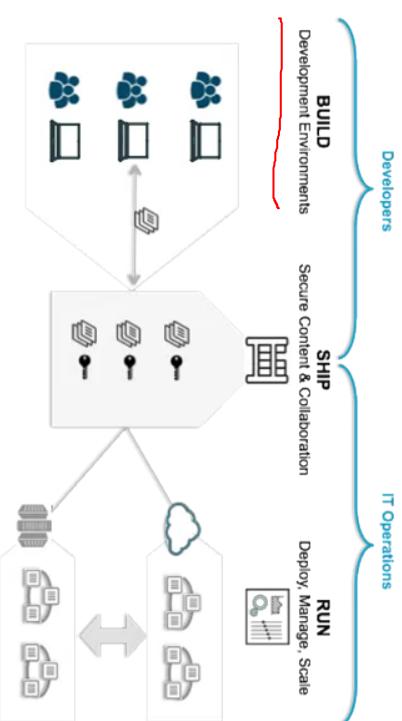
BITS Pilani, Pilani Campus
437

Play with Docker



1. Open your browser to Play with Docker:
<https://labs.play-with-docker.com/>
2. <https://training.play-with-docker.com/>

CaaS workflow



Containers as a Service (CaaS)



Create a repo

1. Sign in to Docker Hub.
2. Click the Create Repository button.
3. Give the desired repo name and make sure the Visibility is Public.
4. Click the Create button!

Share the application



BITS Pilani, Pilani Campus
442

Containers as a Service (CaaS)



BITS Pilani, Pilani Campus
444

- Containers as a Service is a model where IT organizations and developers can work together to build, ship and run their applications anywhere.
- CaaS enables an IT secured and managed application environment consisting of content and infrastructure, from which developers are able to build and deploy applications in a self service manner.

BITS Pilani, Pilani Campus
441

Self Study

- <https://aws.amazon.com/blogs/apn/deploying-code-faster-with-serverless-framework-and-aws-service-catalog/>

- How to use GitHub with Docker: <https://docs.docker.com/ci-cd/github-actions/>

- Minikube: <https://minikube.sigs.k8s.io/docs/start/>



Monitoring

BITS Pilani
Pilani Campus



Prof. Akanksha Bharadwaj
Asst. Prof. CSIS



References

- Ques. How to use GitHub with Docker

- Chapter 8,9 Microservices in Action
- Link: <https://azure.microsoft.com/en-in/overview/kubernetes-deployment-strategy/>
- <https://www.docker.com/blog/containers-as-a-service-caas/>
- <https://docs.docker.com/>



Docker and git



446

- Self study:
Link: <https://docs.docker.com/ci-cd/github-actions/>

- Link: <https://docs.docker.com/ci-cd/github-actions/>



Docker and git



- Ques. How to use GitHub with Docker

- Chapter 8,9 Microservices in Action
- Link: <https://azure.microsoft.com/en-in/overview/kubernetes-deployment-strategy/>
- <https://www.docker.com/blog/containers-as-a-service-caas/>
- <https://docs.docker.com/>

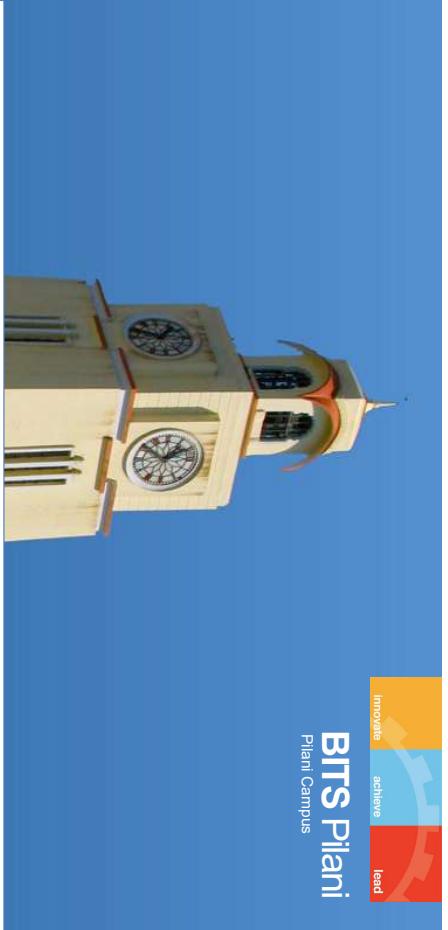
447

Monitoring

- Monitoring involves the collection, analysis, and visualization of various metrics and data points from a system.

- It focuses on tracking predefined metrics and observing the system's behavior to ensure its health, performance, and availability.

- Monitoring is typically based on pre-established monitoring tools or frameworks that capture specific aspects of the system, such as CPU usage, memory utilization, response times, error rates, and network traffic.



450

Difference between Observability and Monitoring

In a web application

- Monitoring might involve tracking the number of requests per second, response times, HTTP error rates, database connection pool utilization, and other relevant metrics. These metrics are typically collected at regular intervals and displayed on dashboards or sent as alerts to notify system administrators or developers about potential issues or anomalies in the system's performance.
- Observability could involve collecting and analyzing detailed logs of requests and responses, tracing individual requests across different microservices, and capturing custom events or signals that indicate specific conditions or scenarios within the application. This richer set of data enables engineers to explore and investigate the system's behavior in greater depth, even in complex distributed architectures.

Observability



452

- The ability to measure a system's current state based on the data it generates, such as logs, metrics, and traces.
- Observability relies on telemetry derived from instrumentation that comes from the endpoints and services in your multi-cloud computing environments.
- Observability encompasses Monitoring

SE ZG583, Scalable Services Lecture No. 14

449

451

Monitoring



Monitoring

- Can use Monitoring to: **Diagnose issues** OR **Prevent them from arising**

Diagnosing Issues

- Broken services (network problem, internal service problem)
- Communication (connectivity lost between services, or between service and event queue)
- Alert on the abnormal decrease in throughput of a service

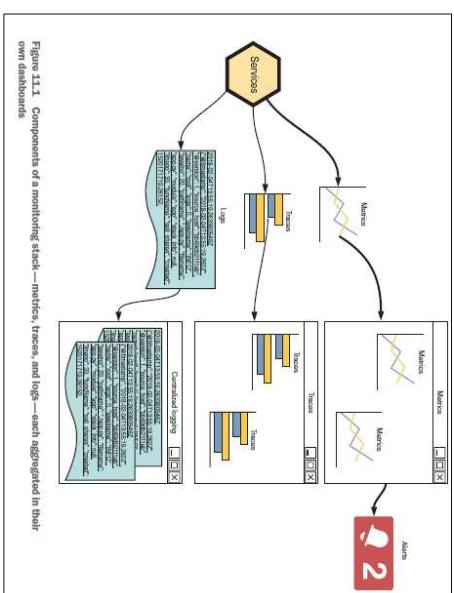


Figure 11.1 Components of a monitoring stack—metrics, traces, and logs—each aggregated in their own dashboards

BITS Pilani, Pilani Campus
454

Monitoring Stack



Monitoring - Example



- A monitoring stack comprises various components working together to collect, store, analyze, and visualize data about the health, performance, and behavior of a system.
- Here are the key components typically found in a monitoring stack:
 - Metrics
 - Logs
 - Traces

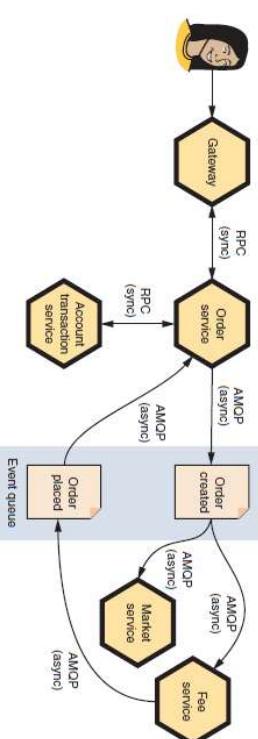


Figure 11.2 Services involved in placing orders and their communication protocols

BITS Pilani, Pilani Campus
456

Golden Signals



3. Traffic (Synonym for Load) - Amount of use of service per time unit

- This signal measures the demand placed on a system.
- It can vary depending on the type of system being observed, the number of requests per second, network I/O, and so on.
- Correlation to Throttling of the system?? Reliability? Steps to overcome?

4. Saturation - Consumption of System Resources

- How much resources (Memory, I/O) your application takes. Application performance degrades when the system reaches 100% utilisation.
- At a given point, this measures the capacity of the service.
- It mainly applies to resources that tend to be more constrained, like CPU, memory, and network.

You need to monitor these signals and keep them in control for making your service available throughout the time.

Golden Signals (Metrics)



• Focus on 4 golden signals while collecting metrics - Latency, Errors, Traffic and Saturation

1. Latency (Synonym for Lag) - Time it takes to serve a request

- Measures how much time passes between when you make a request to a given service and when the service completes the request. We can infer that the service is degrading if it shows increasing latency.

2. Errors - Rate of requests that fail

This signal determines the number of requests that don't result in a successful outcome.

- Explicit error - HTTP 500 (Unexpected Server Error)
- Implicit error - HTTP 200 but with wrong content

Types of Metrics



2. Gauges

Gauges are metrics representing single numerical values that can arbitrarily go up or down.

Some examples of metrics using gauges are:

- Number of connections to a database
- Current Memory used
- CPU used
- Load average
- Number of services operating abnormally
- Number of currently running processes

Types of Metrics



Visualising the output of Golden Signals

Four ways of Showcasing Metrics - Counters, Gauges, Histograms, Summaries

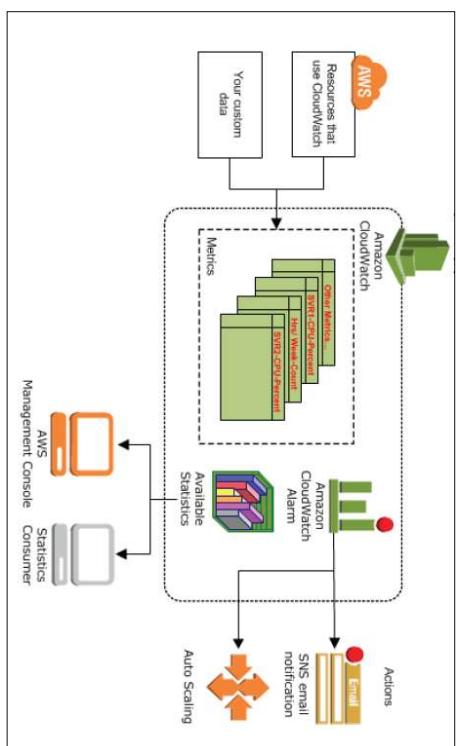
1. Counters

Counters are a cumulative metric representing a single numerical value that'll always increase or be reset to zero on restart. Examples of metrics using counters are:

- Number of requests
- Number of errors
- Number of each HTTP code received [HTTP 200, 400x, 500x etc]
- Bytes transmitted

You shouldn't use a counter if the metric it represents can also decrease. Ex: Number of currently running processes. For that, you should use a gauge instead.

AWS CloudWatch Monitoring Tool



BITS Pilani, Pilani Campus
462

Types of Metrics



3. Histograms

- You use histograms to sample observations and categorize them in configurable buckets
- per type, time, and so on.
- Metrics which samples observations (usually things like request durations or response sizes) and counts them in configurable buckets. It also provides a sum of all observed values.

Examples of metrics represented by histograms are:

- Latency of a request
- I/O latency
- Bytes per response

Key aspects of Alerts



1. **Thresholds:** Alerts are typically triggered when a monitored metric exceeds or falls below a specified threshold. For example, an alert might be set to trigger when CPU utilization exceeds 90% or when the response time of a service exceeds a certain duration.
2. **Conditions:** Alerts can be based on various conditions and combinations of metrics. They can include simple conditions like a single metric crossing a threshold or complex conditions that involve multiple metrics and their relationships. For instance, an alert might be triggered when both CPU utilization and Memory usage exceed certain levels simultaneously.
3. **Severity levels:** Alerts can be categorized into different severity levels to indicate the urgency or impact of the issue. For example, critical alerts indicate severe problems that require immediate attention, while warnings or informational alerts may represent less critical situations.

BITS Pilani, Pilani Campus
464

Alerts



Alerts in monitoring refer to notifications or warnings triggered by predefined conditions or thresholds to indicate potential issues or anomalies in a system.

- When monitoring systems detect that a specific metric or event has crossed a predefined threshold, an alert is generated and sent to system administrators, operators, or developers.
- Alerts play a crucial role in timely identification and response to problems, allowing for proactive troubleshooting and mitigation.

Alerts



Raising Sensible & Actionable Alerts

- Aim to have as few alerts as possible.
- Alerting on symptoms that are associated with end-user pain rather than trying to catch every possible way that pain could be caused.
- Alerts should link to relevant consoles and make it easy to figure out which component is at fault.
- For online servicing systems, the key metric is on high latency and error rates as high up in the stack as possible.
- For offline processing systems, the key metric is how long data takes to get through the system, so page if that gets high enough to cause user impact.
- For batch jobs it makes sense to page if the batch job has not succeeded recently enough, and this will cause user-visible problems.
- While not a problem causing immediate user impact, being close to capacity often requires human intervention to avoid an outage in the near future.
- It is important to have confidence that monitoring is working. Accordingly, have alerts to ensure that Prometheus servers, Alertmanagers, PushGateways, and other monitoring infrastructure are available and running correctly.

Key aspects of Alerts



4. **Notification channels:** When an alert is triggered, it needs to be communicated to the appropriate individuals or teams responsible for addressing the issue. Notification channels can include Email, SMS, Slack, PagerDuty, or other collaboration and incident management tools. [Or Cloud ones like Amazon SNS, SES etc.]
5. **Escalation policies:** In complex systems or larger organizations, alerts may follow predefined escalation policies. These policies determine who receives the initial alert and specify subsequent steps if the alert is not acknowledged or resolved within a certain timeframe.
6. **Alert suppression and de-duplication:** Monitoring systems often employ mechanisms to prevent excessive alerting or alert floods. This can include suppressing redundant alerts for the same underlying issue or consolidating similar alerts into a single notification.

Logs



Useful Info in Log Entries & Logging the Right Information

- Provide the logs wherever necessary.
- Use Tags to identify Modules, Methods, Requests / Response etc.,
- Provide the Warning / Error Code along with Description aptly based on the operation you perform.
- Make the logs user friendly with simple language.
- Don't put logs everywhere which will be too noisy when you are stuck on an error and debugging it.

Logs and Traces



Using Logs & Traces

- Logs provide information about what happened inside the service.
- Tracing tells you what happened between services/components and their relationships.
- This is extremely important for microservices, where many issues are caused due to the failed integration between components.

Log - Centralized

Single Application running on Single Machine – Single Application Log

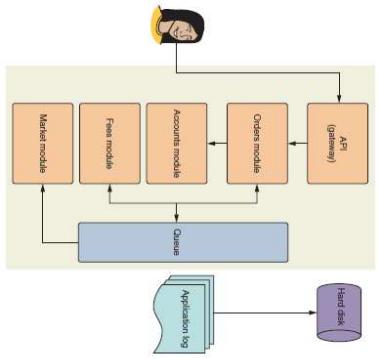


Figure 12.1. The sell order use case implemented in a single application

Situation in Microservice

- Multiple Services
- Each service multiple replicas?
- How to have centralized log?

BITS Pilani, Pilani Campus
470

Logs

Tools for Logging

- **Elk** - Log management platform that works by enabling you to gather massive amounts of log data from anywhere across your infrastructure into a single place, then search, analyse and visualize it in real time.
- **Splunk** - Based on event management and handles the data to give their enterprise with alerts and event logs.
- **Grafana** - Provides charts, graphs, and alerts for the web when connected to supported data sources
- **Fluentd** - Data collector for building the unified logging layer. Once installed on a server, it runs in the background to collect, parse, transform, analyse and store various types of data.

Log – Distributed – Difficulties in accessing?

1. Persistence - Make sure you persist log data so it survives through service restarts and scaling

2. Aggregation - Aggregate all log data from multiple services and instances of those services in a central location

3. Querying - Make the stored data usable, allowing for easy searches and further processing

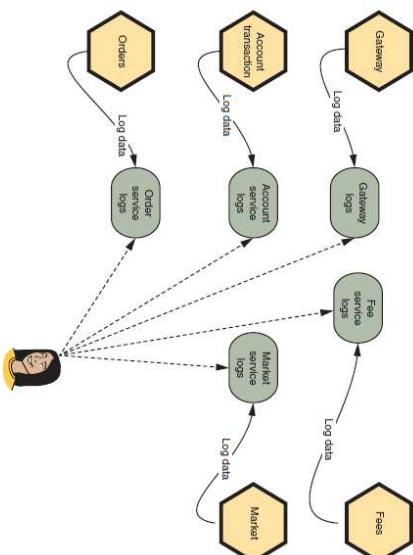


Figure 12.3 Accessing logs for each service in different running instances is challenging.

BITS Pilani, Pilani Campus
472

Log - Distributed

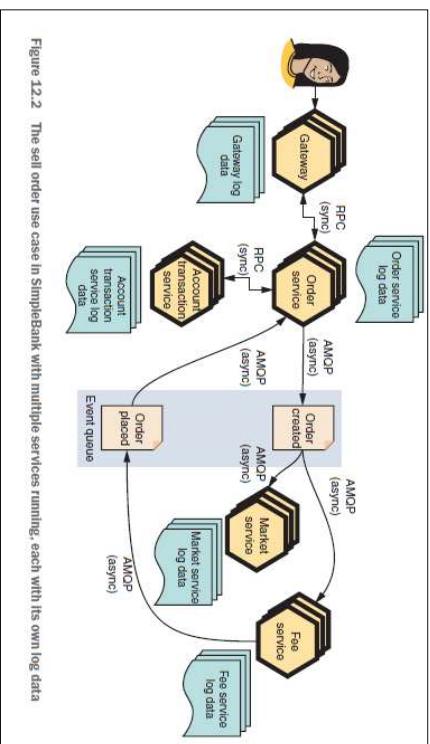


Figure 12.2 The sell order use case in SimpleBank with multiple services running, each with its own log data

Sample Log collected from Logstash library Python

Logstash - A tool to collect, process, and forward events and log messages from multiple sources. It provides multiple plugins to configure data collecting.

```
Information about the source: the host
running the application
{
    "source_host": {"@version": "7003131928a", "@type": "filebeat", "host": {"name": "elasticsearch-01", "os": {"name": "Ubuntu", "version": "18.04.1 LTS"}, "ips": [{"ip": "10.0.2.15"}], "mac": "56:84:7A:41:00:04", "cpu": {"model": "Intel(R) Core(TM) i5-8265U CPU @ 1.60GHz", "cores": 4, "frequency": 1600000000}, "mem": 8192, "swap": 0}, "offset": 20, "relative_offset": 38646125793457031, "process": 1, "processes": 1, "service": "orders-service", "axis": 1, "tags": ["runnables-orders"], "file": "runnables-orders.py", "function": "start", "line": 64, "module": "runnables", "stack": "File: /home/elasticsearch-01/.local/lib/python3.6/site-packages/elastic/runnables/orders.py:64 in start\n  mainProcess = Process(target=mainThread, args=(args,))\n  mainProcess.start()\n  mainProcess.join()\n  return mainProcess", "thread": 14061251794546, "timestramp": "2018-02-02T18:12:00.119Z", "version": 1, "message": "Starting service orders-service", "level_name": "INFO", "stack_info": null, "thread_name": "MainThread", "ts": 1520275324.189045, "created": 1520275324.189045}
}

Version of the formatter [logstash-formatter] 1
```



Traces [Distributed Tracing]

X-Ray Tracing

- X-Ray helps developers analyze and debug distributed applications, such as those built using a microservices architecture.
- With X-Ray, you can understand how your application and its underlying services are performing, and identify and troubleshoot the root cause of performance issues and errors.
- **Complete Tracing** - X-Ray provides an end-to-end, cross-service view of requests made to your application. It gives you a view of requests flowing through your application by aggregating the data gathered from individual services in your applications.

BITS Pilani, Pilani Campus
474

Types of Log

- Application logs
- Database logs
- Network logs
- Performance data collected from the underlying operating system



Traces [Distributed Tracing]

Tracing Interaction between Services

- Tracing is a fundamental process in software engineering, used by programmers along with other forms of logging, to gather information about an application's behaviour.
- Traditional tracing runs into problems when it is used to troubleshoot applications built on a distributed software architecture.
- Since microservices scale independently, it's common to have multiple iterations of a single service running across different servers, locations, and environments simultaneously, creating a complex web through which a request must travel.
- These requests are nearly impossible to track with traditional techniques designed for a single service application.
- **Distributed tracing solutions** solve this problem, and numerous other performance issues, because it can track requests through each service or module and provide an end-to-end narrative account of that request.
- **Distributed tracing, sometimes called distributed request tracing, is a method to monitor applications built on a microservices architecture.**
- IT and DevOps teams use distributed tracing to follow the course of a request or transaction as it travels through the application that is being monitored.
- Analysts, SREs, developers, and others can observe each iteration of a function, enabling them to conduct performance monitoring by seeing which instance of that function is causing the app to slow down or fail, and how to resolve it.

BITS Pilani, Pilani Campus
476

BITS Pilani, Pilani Campus
473

Monitoring in Kubernetes [Self Study – Exercise]



1. Kubernetes Monitoring - <https://logz.io/blog/kubernetes-monitoring/>

2. Prometheus and Grafana setup in Minikube -
<https://brain2life.hashnode.dev/prometheus-and-grafana-setup-in-minikube>

Thank You!



Logs and Traces

Logging	Tracing
Consumed primarily by system administrators	Consumed primarily by developers
Logs "high level" information (e.g. failed installation of a program)	Logs "low level" information (e.g. a thrown exception)
Must not be too "noisy" (containing many duplicate events or information that is not helpful for its intended audience)	Can be noisy
Event log messages are often localized	Localization is rarely a concern
Addition of new types of events, as well as new event messages, need not be agile	Addition of new tracing messages <i>must</i> be agile



References

- Book: Chapter 11, 12 Microservices in Action
- Explore more about Prometheus:
<https://www.youtube.com/@PromLabs>

478

480

BITS Pilani, Pilani Campus

BITS Pilani, Pilani Campus

479

What is Kubernetes?

- Container management tool?????
- Kubernetes is a container management tool
- How is it different from Docker?



SE ZG583, Scalable Services Lecture No. 15

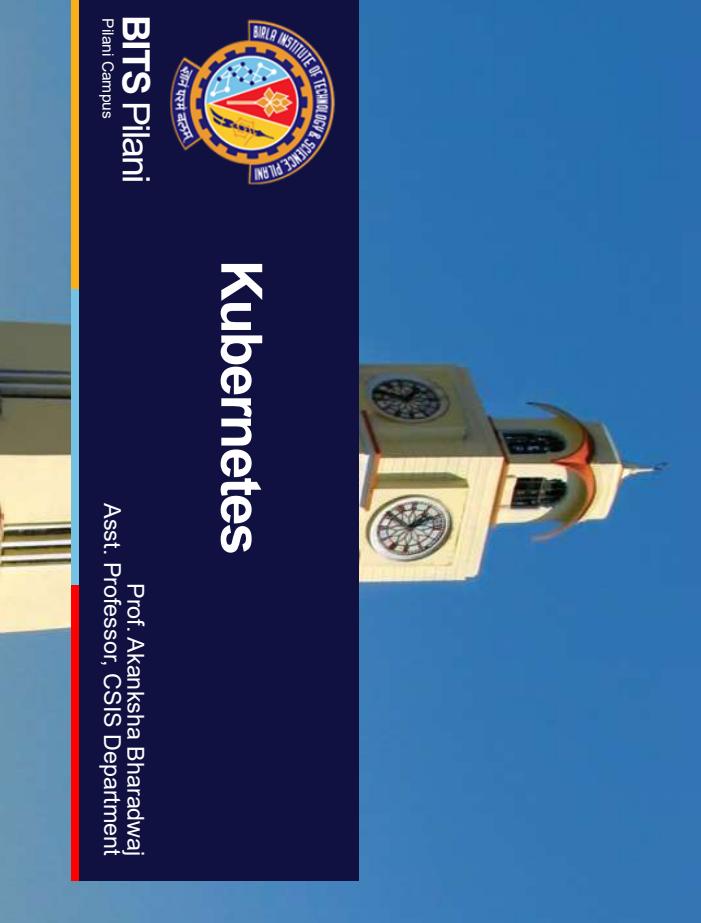
482

Agenda

Kubernetes

- What is Kubernetes?
- Deployment of Microservices using Kubernetes
- Scalability in Kubernetes
- Security in Kubernetes
- CI/CD using Kubernetes
- Kubernetes Dashboard

484



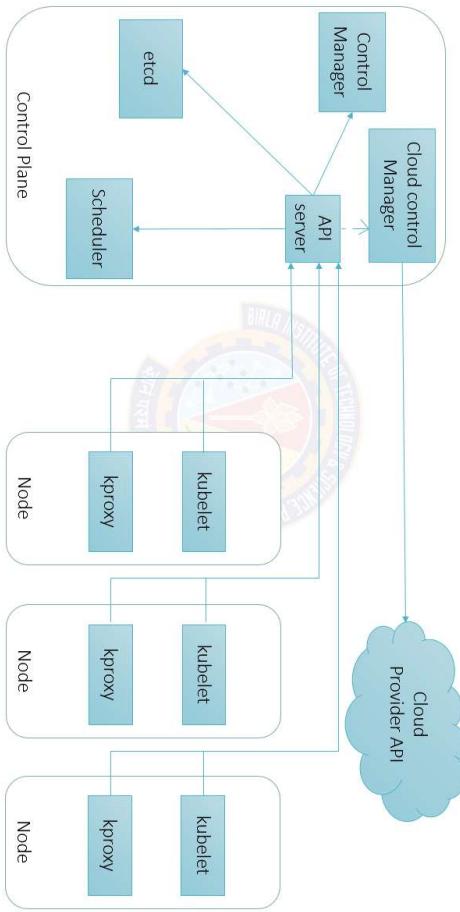
Kubernetes

Kubernetes design principles

A Kubernetes cluster should be:

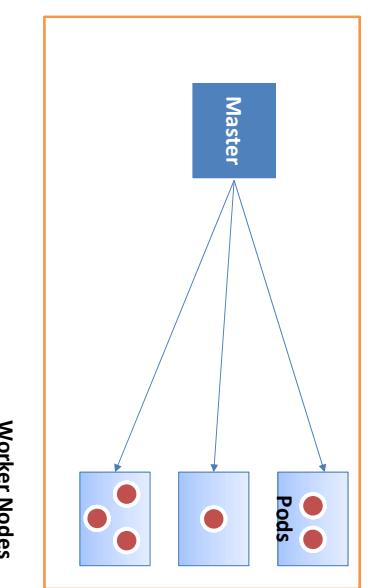
- Secure
- Easy to use
- Extendable

Kubernetes Architecture



BITS Pilani, Pilani Campus
486

Basic Kubernetes Architecture



Worker Nodes

What does K8s provides?

- Service discovery and load balancing

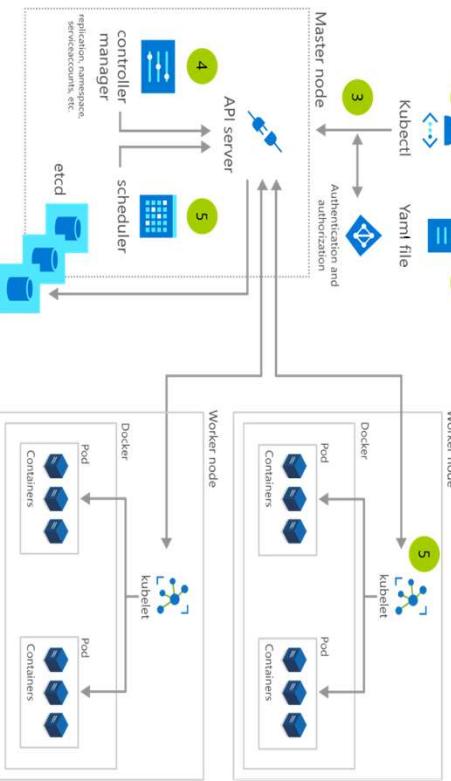
- Storage orchestration
- Automated rollouts and rollbacks
- Automatic bin packing
- Self-healing
- Secret and configuration management



BITS Pilani, Pilani Campus
485

488

How Deployment works in Kubernetes

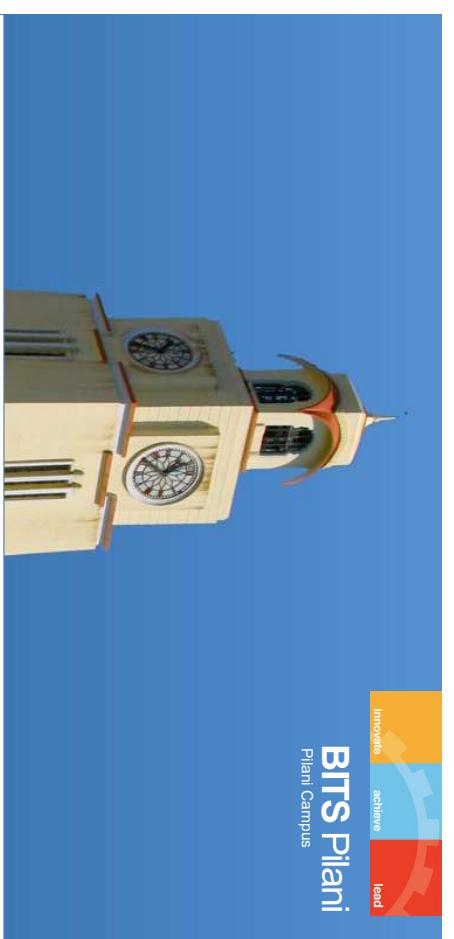


490

Important elements of a Kubernetes deployment

- YAML file
- Pods
- ReplicaSet
 - Kube-controller-manager
 - Kube-scheduler
- Rollout

Scalability in Kubernetes



492

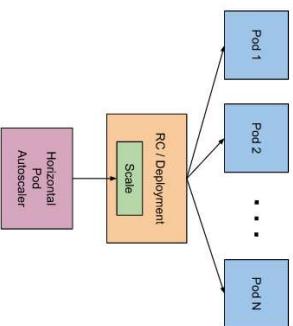
Kubernetes practice lab sources

<https://training.play-with-docker.com/>
<https://training.play-with-kubernetes.com/>
<https://labs.play-with-k8s.com/>



Horizontal Pod Autoscaler

- We can use the Kubernetes Horizontal Pod Autoscaler to automatically scale the number of pods in a deployment.



Cluster Autoscaler

- While HPA scales the number of running pods in a cluster, the cluster autoscaler can change the number of nodes in a cluster.

BITS Pilani, Pilani Campus
494



Vertical Pod Autoscaler

- We can use the Kubernetes Vertical Pod Autoscaler to automatically adjust the resource requests and limits for containers running in a deployment's pods.



Types of scaling IN Kubernetes:

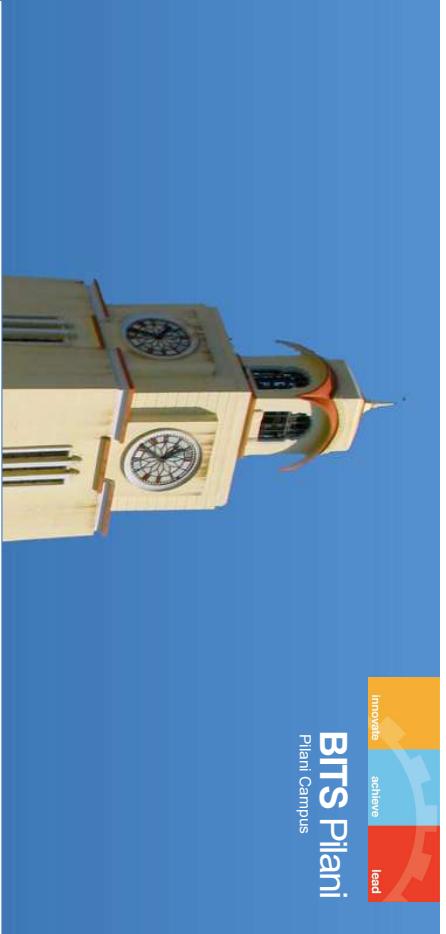
- Horizontal Pod Autoscaler (HPA)
- Vertical Pod Autoscaler (VPA)
- Cluster Autoscaler

With Auto scaling Kubernetes will automatically scale up your cluster as soon as you need it, and scale it back down to save your money

Introduction

- The Kubernetes storage architecture is based on Volumes as a central abstraction.
- Volumes can be persistent or non-persistent
- Kubernetes allows containers to request storage resources dynamically, using a mechanism called volume claims.

Storage in Kubernetes



BITS Pilani, Pilani Campus
498

Non-Persistent Storage

- By default, Kubernetes storage is temporary
 - Any storage defined as part of a container in a Kubernetes Pod, is held in the host's temporary storage space, which exists as long as the pod exists, and is then removed.
 - Container storage is portable, but not durable.

Volumes

- Volumes are the basic entity containers use to access storage in Kubernetes.
- Support for local storage devices, NFS and cloud storage services.
- Regular Volumes are deleted when the Pod hosting them shuts down.



BITS Pilani, Pilani Campus
500

Core Security Areas in Kubernetes

- API Access Control
- Workload Security
- Network Policies
- Secrets Management
- Cluster Hardening



Kubernetes Security

502

Why Security Matters in Kubernetes



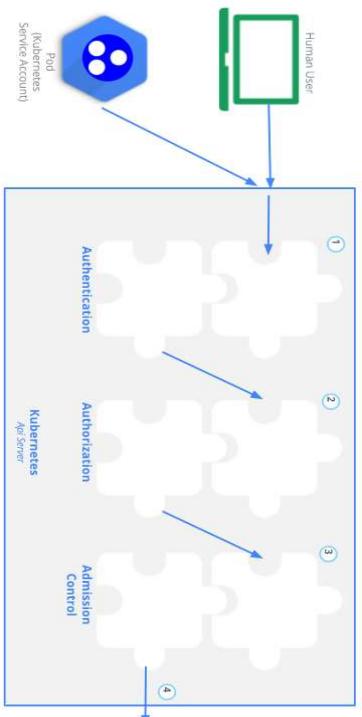
- Dynamic, distributed architecture
- Multiple attack surfaces: API server, etcd, network, containers
- Shared environments and multitenancy

- PVs are volume plugins like Volumes, but have a lifecycle independent of any individual Pod that uses the PV.

- In order for a Pod to start using these volumes, it must request a volume by issuing a persistent volume claim (PVC).

Security Tools & Ecosystem

- Kube-bench (CIS benchmark checks)
- Kube-hunter (penetration testing)
- Trivy or Grype (image scanning)
- OPA/Gatekeeper (policy enforcement)
- Falco (runtime security)

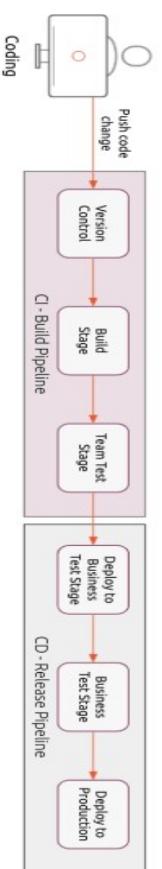


Access control in Kubernetes

BITS Pilani, Pilani Campus
506

What is CI/CD pipeline?

- CI/CD pipeline is a series of stages and automated steps software goes through, from code development to production deployment.



BITS Pilani, Pilani Campus
508

Other Security measures in Kubernetes

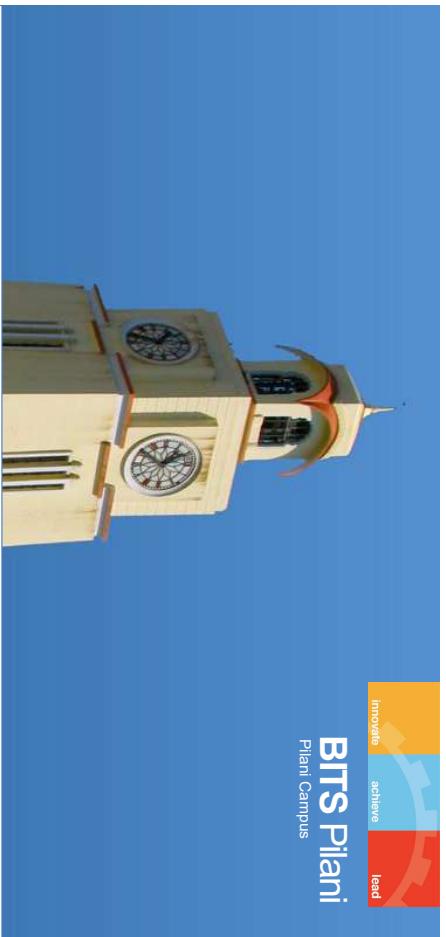
- Use namespaces to establish security boundaries

- Separate sensitive workloads
- Secure cloud metadata access
- Create and define cluster network policies
- Run a cluster-wide Pod Security Policy
- Harden node security
 - Turn on audit logging
 - Keep the Kubernetes version up to date



Why CI/CD in Kubernetes?

- Native support for automation and scaling
- Ideal for microservices and container-based apps
- Easily integrates with modern DevOps tools
- Supports rolling updates, blue-green, and canary deployments



BITS Pilani, Pilani Campus
510

CI/CD Workflow in Kubernetes

- Developer pushes code to Git
 - CI pipeline runs tests and builds image
 - Image is pushed to registry
 - CD tool pulls new image
 - Kubernetes deployment is updated
 - Monitoring and rollback mechanisms activated



BITS Pilani, Pilani Campus
512

Core CI/CD Components for Kubernetes

- Source control: GitHub, GitLab

- CI servers: Jenkins, GitHub Actions, GitLab CI, CircleCI

- Build tools: Docker, Kaniko, BuildKit

- Artifact storage: Docker Hub, Amazon ECR, Harbor

- CD tools: ArgoCD, Flux, Spinnaker

- Kubernetes manifests: Helm, Kustomize, YAML

CI/CD with Kubernetes

Testing in CI/CD Pipelines



- Unit & integration tests
- Linting (YAML, Helm)
- Security scans (Trivy, Snyk)
- Kubernetes manifest validation
- Dry runs (kubectl apply --dry-run)
- Canary tests in production

Security & Observability in CI/CD



- Image scanning in CI (e.g., Trivy)
- RBAC and Secrets Management in CD
- Monitoring with Prometheus/Grafana
- Logging with Fluentd, Loki
- Alerting with Alertmanager or Slack integration

GitOps Approach to CD



- Uses Git as the **single source of truth** for deployments
- Tools like **ArgoCD** and **Flux** sync cluster state with Git
- Benefits: Auditable, repeatable, easy rollbacks

Packaging & Deployment



- Use **Helm charts** or **Kustomize** to manage deployment logic
- Automate rollouts with **ArgoCD**, **Flux**, or native Kubernetes
- Support for **blue-green**, **canary**, and **rolling updates**

BITS Pilani, Pilani Campus

514

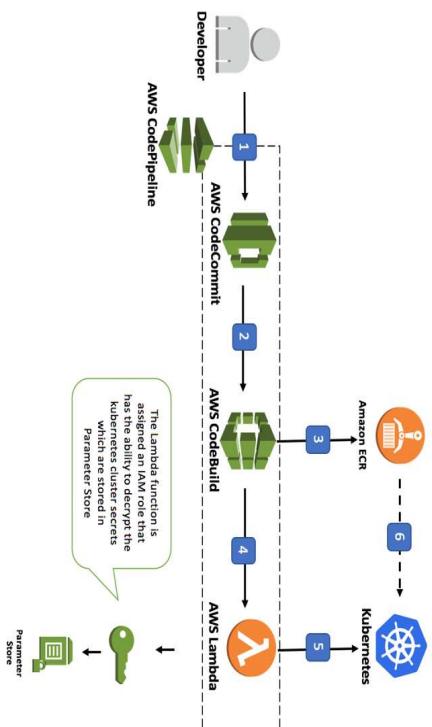
BITS Pilani, Pilani Campus

516

BITS Pilani, Pilani Campus

513

Example of Continuous Deployment to Kubernetes



BITS Pilani, Pilani Campus
518

Infrastructure as Code in Kubernetes



- Kubernetes infrastructure and configuration are managed as code using tools like Terraform or Kubernetes-specific tools (e.g., kops, kubectl apply) to provision and manage Kubernetes clusters and associated resources.

- Dashboard is a web-based Kubernetes user interface.
 - You can use Dashboard to get an overview of applications running on your cluster, as well as for creating or modifying individual Kubernetes resources
- Dashboard also provides information on the state of Kubernetes resources in your cluster and on any errors that may have occurred.

Deploying the Dashboard UI



- Command:
`kubectl apply -f https://raw.githubusercontent.com/kubernetes/dashboard/v2.3.1/aio/deploy/recommended.yaml`

BITS Pilani, Pilani Campus
520

Kubernetes dashboard



BITS Pilani, Pilani Campus
517

References

- 1. Link: <https://kubernetes.io/docs>
- 2. Book: Kubernetes in Action by Marko Lukša
- 3. Link: <https://azure.microsoft.com/en-in/overview/kubernetes-deployment-strategy/>
- 4. Link: <https://aws.amazon.com/blogs/devops/continuous-deployment-to-kubernetes-using-aws-codepipeline-aws-codecommit-aws-codebuild-amazon-ecr-and-aws-lambda/>



Agenda

- How to create Kubernetes Cluster on AWS
- Create a Kubernetes cluster on AWS using eksctl
- How to deploy an application on Kubernetes (cloud)



BITS Pilani, Pilani Campus
522

524

Self Study

- <https://www.magalix.com/blog/create-a-ci/cd-pipeline-with-kubernetes-and-jenkins>
- <https://ubuntu.com/tutorials/gitlab-ci-cd-pipelines-with-microk8s#1-overview>
- <https://www.cncf.io/blog/2021/03/22/kubernetes-security/>
- <https://www.katacoda.com/courses/kubernetes/launch-single-node-cluster>
- Article: <https://techyinurse.medium.com/deploying-microservices-to-kubernetes-minikube-b807994ee4bd>
- Article: <https://techdozo.dev/building-and-deploying-microservices-application-locally-in-the-kubernetes/>

Kubernetes Cluster

BITS Pilani
Pilani | Doha | Gea | Hyderabad

Prof. Akanksha Bharadwaj

Pre-requisites

- Assign minimum IAM policies required for eksctl usage

Attached directly	Policy type	Add inline policy
AmazonEC2FullAccess	AWS managed policy	x
IAMFullAccess	AWS managed policy	x
AmazonEC2ContainerRegistryFullAccess	AWS managed policy	x
FSECAutoScaling	Managed policy	x
FSECAmazon	Managed policy	x
FSECBNetworking	Managed policy	x
FSECEKS	Managed policy	x
FSESCloudFormation	Managed policy	x
AWSAppRunnerServicePolicyForECRAccess	AWS managed policy	x

526

Pre-requisites

- Create an AWS account
- Create a user in IAM management console of AWS and download csv file that contains the AWS Secret Access Key
- Create using AWS CLI

<https://docs.aws.amazon.com/eks/latest/userguide/create-cluster.html>

User name	Groups	Last activity	MFA	Password	Actions

Add users

Introducing the new Users list experience

We've redesigned the Users list experience to make it easier to use. Let us know what you think.

528

How to create an Amazon EKS cluster

- Create using eksctl
- Create using AWS management Console
- Create using AWS CLI

<https://docs.aws.amazon.com/eks/latest/userguide/create-cluster.html>



525

Pre-requisites

Install kubectl – A command line tool for working with Kubernetes clusters. Use the latest version. I used Kubernetes version 1.20

- Open a PowerShell terminal in Windows
- Download the Amazon EKS vended kubectl binary for your cluster's Kubernetes version from Amazon S3. Command for Windows:


```
curl -o kubectl.exe https://amazon-eks.s3.us-west-2.amazonaws.com/1.20.4/2021-04-12/bin/windows/amd64/kubectl.exe
```
- Download the SHA-256 sum for your cluster's Kubernetes version for Windows. Command on windows


```
curl -o kubectl.exe.sha256 https://amazon-eks.s3.us-west-2.amazonaws.com/1.20.4/2021-04-12/bin/windows/amd64/kubectl.exe.sha256
```
- Check the SHA-256 sum for your downloaded binary. Compare the generated SHA-256 sum in the command output against your downloaded SHA-256 file. The two should match, although the PowerShell output will be uppercase.
- Edit your system PATH environment variable to add the directory containing kubectl.exe binary
- Check the download by using:


```
PS C:\Users\admin> kubectl version --short --client
Client Version: v1.20.4-eks-607464
```

530

Create your Amazon EKS cluster and nodes

To create your cluster with Amazon EC2 Linux managed nodes

- Create key-pair with the following command: aws ec2 create-key-pair --region us-east-1 --key-name myKeyPair

```
PS C:\Users\admin\aws\kubernetes\k8s-postman-tutorial> aws ec2 create-key-pair --region us-east-1 --key-name myKeyPair
{
    "keyFingerprint": "ec:8c:22:ab:1b:1a:85:d5:b1:cc:69:8f:4d:64:57:3e:58:fa:70:e7",
    "keyName": "myKeyPair"
}

PS C:\Users\admin> aws --version
aws-cli/2.0.20 Python/3.8.8 Windows/10 exe/AMD64 prompt/off
```

Pre-requisites

- Download, Install and configure AWS CLI

• Check the AWS CLI version to confirm

```
PS C:\Users\admin> aws --version
aws-cli/2.0.20 Python/3.8.8 Windows/10 exe/AMD64 prompt/off
```

- To configure use the below-mentioned command and then specify AWS Access key id (get from IAM user), AWS Secret Access Key (get from IAM user), default region and default format

```
aws configure
aws configure
AWS Access Key ID [*****]: XXXXXXXXXXXXXXIOEV1:XXXXXXXXXXXXXXXXXXXXXX
AWS Secret Access Key [*****]: XXXXXXXXXXXXXXXXXXXXXX3ZVM:XXXXXXXXXXXXXXXXXXXXXX
Default region name [us-east-1]: us-east-1
Default output format [json]: json
```

Access key ID	Created	Last used	Status
XXXXXXXXXXXXXXIOEV1:XXXXXXXXXXXXXXXXXXXXXX	2021-07-27 15:59 UTC+0530	2021-07-28 16:12 UTC+0530 with sts in us-east-1	Active Make inactive ✘

532

Pre-requisites

Install eksctl – A command line tool for working with EKS clusters that automates many individual tasks. Use version 0.58 or later

- Install the binaries with the following command: choco install -y eksctl
- If they are already installed, run the following command to upgrade: choco upgrade -y eksctl
- Test that your installation was successful with the following command: eksctl version

```
PS C:\Users\admin> eksctl version
0.58.0
```



Create your Amazon EKS cluster and nodes

- Create cluster command: `eksctl create cluster --name test1-cluster --region us-east-1 --zones=us-east-1a,us-east-1b --with-oidc --ssh-access --ssh-public-key myKeyPair --managed`

```
PS C:\Users\admin\aws\kubernetes\eks-postman-tutorial> eksctl create cluster --name test1-cluster --region us-east-1 --zones=us-east-1a,us-east-1b --with-oidc --ssh-access --ssh-public-key myKeyPair --managed
2021-07-28 12:43:48 [0] waiting for CloudFormation stack "eksctl-test1-cluster-nodegroup-ing-e4f43d067" to finish
2021-07-28 12:44:02 [0] waiting for CloudFormation stack "eksctl-test1-cluster-nodegroup-ing-e4f43d067"
2021-07-28 12:44:21 [0] waiting for CloudFormation stack "eksctl-test1-cluster-nodegroup-ing-e4f43d067"
2021-07-28 12:44:39 [0] waiting for CloudFormation stack "eksctl-test1-cluster-nodegroup-ing-e4f43d067"
2021-07-28 12:45:01 [0] waiting for CloudFormation stack "eksctl-test1-cluster-nodegroup-ing-e4f43d067"
2021-07-28 12:45:21 [0] waiting for CloudFormation stack "eksctl-test1-cluster-nodegroup-ing-e4f43d067"
2021-07-28 12:45:24 [0] waiting for CloudFormation stack "eksctl-test1-cluster-nodegroup-ing-e4f43d067"
2021-07-28 12:45:24 [0] waiting for the control plane availability...
2021-07-28 12:45:24 [0] saved kubeconfig as "C:\\Users\\admin\\kube\\config"
2021-07-28 12:45:24 [0] no tasks
2021-07-28 12:45:25 [0] all EKS cluster resources for "test1-cluster" have been created
2021-07-28 12:45:26 [0] nodegroup "ng-e4f43d067" has 2 node(s)
2021-07-28 12:45:26 [0] node "ip-192-168-0-231.ec2.internal" is ready
2021-07-28 12:45:26 [0] node "ip-192-168-0-39.ec2.internal" is ready
2021-07-28 12:45:26 [0] waiting for at least 2 node(s) to become ready in "ng-e4f43d067"
2021-07-28 12:45:26 [0] nodegroup "ng-e4f43d067" has 2 node(s)
2021-07-28 12:45:26 [0] mode "[ip-192-168-0-231.ec2.internal]" is ready
2021-07-28 12:45:26 [0] mode "[ip-192-168-0-39.ec2.internal]" is ready
2021-07-28 12:47:46 [0] kubectl command should work with "C:\\Users\\admin\\kube\\config", try 'kubectl get nodes'
2021-07-28 12:47:46 [0] EKS cluster "test1-cluster" in "us-east-1" region is ready
```

534

Create your Amazon EKS cluster and nodes

- kubectl get nodes -o wide

```
PS C:\Users\admin\aws\kubernetes\eks-postman-tutorial> kubectl get nodes -o wide
NAME          STATUS   ROLES      AGE   VERSION   INTERNAL-IP    EXTERNAL-IP   OS-IMAGE
ip-192-168-0-231.ec2.internal   Ready    <none>   11m   v1.20.4-eks-5b7464   192.168.0.231   54.227.283.142   Amazon Linux
ip-192-168-0-39.ec2.internal   Ready    <none>   11m   v1.20.4-eks-5b7464   192.168.43.39   54.174.148.158   Amazon Linux
ip-192-168-0-231.ec2.internal   Ready    <none>   11m   v1.20.4-eks-5b7464   192.168.0.231   54.227.283.142   Amazon Linux
ip-192-168-0-39.ec2.internal   Ready    <none>   11m   v1.20.4-eks-5b7464   192.168.43.39   54.174.148.158   Amazon Linux
2021-07-28 12:18:47 [0] nodegroup "ng-e4f43d07" will use "" [amazonlinux2/1.20]
2021-07-28 12:18:49 [0] using E2 key pair %q(%string(mil))
2021-07-28 12:18:49 [0] using Kubernetes version 1.20
2021-07-28 12:18:49 [0] creating EKS cluster "test1-cluster" in "us-east-1" region with managed nodes
2021-07-28 12:18:49 [0] will create 2 separate CloudFormation stacks for cluster itself and the initial managed nodegroup
2021-07-28 12:18:49 [0] if you encounter any issues, check CloudFormation console or try 'eksctl utils describe-stacks --region=us-east-1 --cluster=test1-cluster'
2021-07-28 12:18:49 [0] CloudWatch Logging will not be enabled for cluster "test1-cluster" in "us-east-1"
2021-07-28 12:18:49 [0] you can enable it with eksctl utils update-cluster-logging --enable-types=[SPECIFY-YOUR-LOG-TYPE-S-HERE (e.g. all)] --region=us-east-1 --cluster=test1-cluster"
2021-07-28 12:18:49 [0] Kubernetes API endpoint access will use default of {publicAccess=true, privateAccess=false} for c
luster "test1-cluster" in "us-east-1"
```

536

View resources

- get cluster details: `eksctl get cluster --region us-east-1`

```
PS C:\Users\admin\aws\kubernetes\eks-postman-tutorial> eksctl get cluster --region us-east-1
2021-07-28 12:50:11 [0] eksctl version 0.58.0
2021-07-28 12:50:11 [0] using region us-east-1
NAME          REGION   EKSCLOUDCREATED
test1-cluster  us-east-1  True
```



View resources

- kubectl get pods --all-namespaces -o wide

```
PS C:\Users\admin\aws\kubernetes\eks-postman-tutorial> kubectl get pods - -all-namespaces -o wide
NAMESPACE     NAME          READY   STATUS    RESTARTS   AGE   IP           NODE
kube-system   kube-dns-vizx   1/1    Running   0          11m   192.168.0.231   ip-192-168-0-231.ec2.internal
kube-system   kube-dns-vizx   1/1    Running   0          11m   192.168.43.39   ip-192-168-43-39.ec2.internal
kube-system   coredns-5b645f-7z84q  1/1    Running   0          26m   192.168.27.57   ip-192-168-0-231.ec2.internal
kube-system   coredns-5b645f-7z84q  1/1    Running   0          26m   192.168.24.205  ip-192-168-0-231.ec2.internal
kube-system   kube-proxy-6fr7k   1/1    Running   0          11m   192.168.0.231   ip-192-168-0-231.ec2.internal
kube-system   kube-proxy-6fr7k   1/1    Running   0          11m   192.168.43.39   ip-192-168-43-39.ec2.internal
<none>
```

533

Deploy application

Building and containerizing the microservices

The first step of deploying to Kubernetes is to build your microservices and containerize them. So every microservice should be made into a docker image with the relevant ports that should be exposed configured.

Example:

```
docker build -t system:1.0-SNAPSHOT system/
```



```
C:\Users\ADMINI\awskubernetes\guide-cloud-aws\start-dns>sts get-caller-identity --output text --query "Account"
{
    "Account": "12345678901234567890"
}
```

- get AWS account_id: aws sts get-caller-identity --output text --query "Account"

```
C:\Users\ADMINI\awskubernetes\guide-cloud-aws\start-dns>sts get-caller-identity --output text --query "Account"
{
    "Account": "12345678901234567890"
}
```

- Authenticate your docker client: docker login -u AWS -p [password_string]

```
https://[aws_account_id].dkr.ecr.[region].amazonaws.com
```

```
C:\Users\ADMINI\awskubernetes\guide-cloud-aws\start-docker> docker login -u AWS -p [password_string]
The push refers to repository [12345678901234567890.dkr.ecr.ap-south-1.amazonaws.com/12345678901234567890]
12345678901234567890: Pulling from 12345678901234567890
Digest: sha256:40f32d8630a22a21773e259e070a534530a127606120973850f0a1024
Status: Image is up to date for 12345678901234567890:latest
Error response from daemon: Get "https://12345678901234567890.dkr.ecr.ap-south-1.amazonaws.com/v2/_catalog": dial tcp: lookup 12345678901234567890.dkr.ecr.ap-south-1.amazonaws.com on 53: no such host
```

```
aws ecr get-login-password --region ap-south-1
```

```
LogInSuccess
```

538

Deploy application

Building and containerizing the microservices

The first step of deploying to Kubernetes is to build your microservices and containerize them. So every microservice should be made into a docker image with the relevant ports that should be exposed configured.

Example:

```
docker build -t system:1.0-SNAPSHOT system/
```



```
C:\Users\ADMINI\awskubernetes\guide-cloud-aws\start-dns>sts get-caller-identity --output text --query "Account"
{
    "Account": "12345678901234567890"
}
```

- get AWS account_id: aws sts get-caller-identity --output text --query "Account"

```
C:\Users\ADMINI\awskubernetes\guide-cloud-aws\start-dns>sts get-caller-identity --output text --query "Account"
{
    "Account": "12345678901234567890"
}
```

- Authenticate your docker client: docker login -u AWS -p [password_string]

```
https://[aws_account_id].dkr.ecr.[region].amazonaws.com
```

```
C:\Users\ADMINI\awskubernetes\guide-cloud-aws\start-docker> docker login -u AWS -p [password_string]
The push refers to repository [12345678901234567890.dkr.ecr.ap-south-1.amazonaws.com/12345678901234567890]
12345678901234567890: Pulling from 12345678901234567890
Digest: sha256:40f32d8630a22a21773e259e070a534530a127606120973850f0a1024
Status: Image is up to date for 12345678901234567890:latest
Error response from daemon: Get "https://12345678901234567890.dkr.ecr.ap-south-1.amazonaws.com/v2/_catalog": dial tcp: lookup 12345678901234567890.dkr.ecr.ap-south-1.amazonaws.com on 53: no such host
```

```
aws ecr get-login-password --region ap-south-1
```

```
LogInSuccess
```

540

Deploy application

Pushing the images to a container registry

- First, you must authenticate your Docker client to your ECR registry. Start by running the get-login command:

```
aws ecr get-login-password --region ap-south-1
```

```
C:\Users\ADMINI\awskubernetes\guide-cloud-aws\start-dns>aws ecr get-login-password --region ap-south-1
The push refers to repository [12345678901234567890.dkr.ecr.ap-south-1.amazonaws.com/12345678901234567890]
12345678901234567890: Pulling from 12345678901234567890
Digest: sha256:40f32d8630a22a21773e259e070a534530a127606120973850f0a1024
Status: Image is up to date for 12345678901234567890:latest
Error response from daemon: Get "https://12345678901234567890.dkr.ecr.ap-south-1.amazonaws.com/v2/_catalog": dial tcp: lookup 12345678901234567890.dkr.ecr.ap-south-1.amazonaws.com on 53: no such host
```

```
aws ecr get-login-password --region ap-south-1
```

```
LogInSuccess
```

537

How to deploy an application on Kubernetes (cloud)

Deploy application

Tag your container images

Next, you need to tag your container images with the relevant data about your registry:

```
Example : docker tag system:1.0-SNAPSHOT 072302530057.dkr.ecr.us-east-1.amazonaws.com/awsguide/system:1.0-SNAPSHOT
```

```
C:\Users\ADMIN\aws-kubernetes\guide-cloud-aws>start-docker tag system:1.0-SNAPSHOT 072302530057.dkr.ecr.us-east-1.amazonaws.com/awsguide/system:1.0-SNAPSHOT
```

```
C:\Users\ADMIN\aws-kubernetes\guide-cloud-aws>start-docker tag inventory:1.0-SNAPSHOT 072302530057.dkr.ecr.us-east-1.amazonaws.com/awsguide/inventory:1.0-SNAPSHOT
```

Spec:
replicas: 3
selector:
matchLabels:
app: myapp
template:
metadata:
labels:
app: myapp

Spec:
contains:
- name: myapp
image: nginx

542

Creating Deployment Example file

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
name: myapp-deployment
```

```
labels:
```

```
app: myapp
```

```
spec:
```

```
replicas: 3
```

```
selector:  
matchLabels:  
app: myapp  
template:  
metadata:  
labels:  
app: myapp
```



Understanding Kubernetes Yaml file

Deploy application

Make a repository to store the images

```
Example: aws ecr create-repository --repository-name awsguide/system
```

```
C:\Users\ADMIN\aws-kubernetes\guide-cloud-aws>start-docker ecr create-repository --repository-name awsguide/system
```

```
{"repository": {"repositoryArn": "arn:aws:ecr:us-east-1:072302530057:repository/awsguide/system", "repositoryName": "awsguide/system", "repositoryUri": "https://072302530057.dkr.ecr.us-east-1.amazonaws.com/awsguide/system", "createdAt": "1992-07-28T13:58:37.495+00:00", "imageAcceptability": "PUBLISHABLE", "imageScanningOnBuild": false, "encryptionType": "AES256"}, {"repository": {"repositoryArn": "arn:aws:ecr:us-east-1:072302530057:repository/awsguide/inventory", "repositoryName": "awsguide/inventory", "repositoryUri": "https://072302530057.dkr.ecr.us-east-1.amazonaws.com/awsguide/inventory", "createdAt": "1992-07-28T15:51:30+00:00", "imageAcceptability": "PUBLISHABLE", "imageScanningOnBuild": false, "encryptionType": "AES256"}, {"repository": {"repositoryArn": "arn:aws:ecr:us-east-1:072302530057:repository/awsguide/inv", "repositoryName": "awsguide/inv", "repositoryUri": "https://072302530057.dkr.ecr.us-east-1.amazonaws.com/awsguide/inv", "createdAt": "1992-07-28T15:51:30+00:00", "imageAcceptability": "PUBLISHABLE", "imageScanningOnBuild": false, "encryptionType": "AES256"}}
```

Deploy application

Push your images to the registry

```
Finally, push your images to the registry:
```

```
Example:
```

```
docker push [system-repository-uri]:1.0-SNAPSHOT
```

```
C:\Users\ADMIN\aws-kubernetes\guide-cloud-aws>start-docker push 072302530057.dkr.ecr.us-east-1.amazonaws.com/awsguide/system:1.0-SNAPSHOT
```

```
The push refers to repository [072302530057.dkr.ecr.us-east-1.amazonaws.com/awsguide/system]
```

544

Deploying the microservices

- Prepare kubernetes.yaml file.
- `kubectl apply -f kubernetes.yaml`

```
C:\Users\admin\awskubernetes\guide-cLOUD-AWS\start>kubectl apply -f kubernetes.yaml
deployment.apps/system-deployment created
deployment.apps/inventory-deployment created
service/system-service created
service/inventory-service created
```

- `kubectl get pods`

```
C:\Users\admin\awskubernetes\guide-cLOUD-AWS\start>kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
inventory-deployment-599dfc69ad-29ndl  0/1     ImagePullBackoff   0          47s
system-deployment-8655bc87d5-27kq5      0/1     ImagePullBackoff   0          48s
```

Clean up

- When you no longer need your deployed microservices, you can delete all Kubernetes resources by running the `kubectl delete` command:
`kubectl delete -f kubernetes.yaml`

- Delete the ECR repositories used to store the images:

Example: `aws ecr delete-repository --repository-name awsguide/system --force`

- Remove your EKS cluster:

Example: `eksctl delete cluster --name guide-cluster`

Understanding Kubernetes Yaml file

Creating Services Example file

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-clusterip-sv
spec:
  selector:
    app: myapp
  ports:
    - name: myapp-clusterip
      protocol: TCP
      port: 4000
      targetPort: 5000
```

Deploying new version of an existing microservice

Creating Services Example file

- Build the new version of the container image
 - Example: `docker build -t system:2.0-SNAPSHOT system/`
- Tag your container image with the relevant data about your registry:
 - Example: `docker tag system:2.0-SNAPSHOT [system-repository-uri]:2.0-SNAPSHOT`
- Push your image to the registry
 - Example: `docker push [system-repository-uri]:2.0-SNAPSHOT`
- Update the deployment to use the new container image that you just pushed to the registry
 - Example: `kubectl set image deployment/system-deployment system-container=[system-repository-uri]:2.0-SNAPSHOT`
- Use the following command to find the name of the pod that is running the system microservice:
`kubectl get pods`

Amazon EKS security best practices

IAM administrators control who can be *authenticated* (signed in) and *authorized* (have no permissions) to use Amazon EKS resources. IAM is an AWS service that you can use with no additional charge.

Restricting access to the Amazon EC2 instance metadata service (IMDS)

Recommendation is to block pod access to IMDS to minimize the permissions available to your containers if:

- You've implemented IAM roles for service accounts and have assigned necessary permissions directly to all pods that require access to AWS services.
- No pods in your cluster require access to IMDS for other reasons, such as retrieving the current Region.

References

1. Link: <https://kubernetes.io/docs/home/>
2. Link: <https://docs.aws.amazon.com/eks/latest/userguide/>
3. Link: <https://docs.aws.amazon.com/cli/latest/userguide/cli-chap-install.html>
4. Link: <https://openliberty.io/>

Best Practices to use Kubernetes Cluster

• Use the latest version

• Namespaces

• Small container images

• Resource requests and limit

• Autoscale

• Monitoring

• Use readiness Probe

• Use liveness Probe

• Use Role Based Access Controls

• Git-based workflow

Self Study

• Exercise: <https://kubernetes.io/docs/tutorials/stateless-application/guestbook/>

• Exercise: <https://kubernetes.io/docs/tutorials/stateful-application/mysql-wordpress-persistent-volume/>

• Exercise: <https://documenter.getpostman.com/view/13059338/TVmLCyNjTutorial.html>

• Kubernetes dashboard: <https://docs.aws.amazon.com/eks/latest/userguide/dashboard-tutorial.html>

• Security in AWS EKS: <https://docs.aws.amazon.com/eks/latest/userguide/security.html>

• Horizontal pod autoscaler: <https://docs.aws.amazon.com/eks/latest/userguide/horizontal-pod-autoscaler.html>

• CI/CD: <https://aws.amazon.com/blogs/devops/ci-cd-on-amazon-eks-using-aws-codecommit-aws-codepipeline-aws-codedbuild-and-fluxcd/>

• <https://kubernetes.io/docs/setup/production-environment/>

Cloud Computing - Notes on the CAP Theorem

Prof. Douglas Thain, March 2016

Caution: These are high level notes that I use to organize my lectures. You main find them useful for reviewing main concepts, but they aren't a substitute for participating in class.

References

1. Eric Brewer, Towards robust distributed systems. In Proceedings of the 19th Annual ACM Symposium on Principles of Distributed Computing (July 16-19, Portland, Oregon): 7
2. Werner Vogels, Eventually Consistent, Communications of the ACM, Volume 52, Number 1, January 2009.
3. Eric Brewer, CAP Twelve Years Later: How the "Rules" Have Changed, IEEE Explore, Volume 45, Issue 2, 2012.
4. Chapter 7 of Tanenbaum and Steen, "Distributed Systems: Principles and Paradigms", Pearson, 2007.

The CAP Theorem

The CAP theorem¹ is an observation about the tradeoffs inherent in designing a distributed system for storing data. Simply put, the CAP theorem states that a given system design involves a tradeoff between the desirable properties of Consistency, Availability, and Partitionability. A given system cannot maximize all three of these properties simultaneously. The concept is generally attributed to Eric Brewer, who posed it as a conjecture at the PODC conference in 2000, building on ideas of consistency and performance that had been explored by systems and databases for some time.

The three properties may be loosely defined as follows:

Consistency – All clients of the system see the same data.

Availability – Clients are able to access and update data rapidly.

Partitionability – The system is able to operate even when the network fails.

You may sometimes hear this explained as "pick two of three" but that is an oversimplification. Each of these properties exists upon a continuum, and attempting to increase one of them requires decreasing another to a certain degree.

This is all very abstract, and best explained with some examples. To get the examples right, we will have to discuss them in some detail, and then come back to discuss the CAP theorem more generally.

¹ The CAP "theorem" isn't really a proper theorem, since it uses some rather fuzzy

Example System 1 – Direct Access

Let's start off with a model of a distributed file system. Suppose that we have a server S that contains some files named X, Y, and Z. The server is accessible via a network to clients A, B, and C. The clients occasionally wish to modify those files, so they can send messages like "read X" or "write Y" to change those files. As the server receives these requests, it sends a response message back to the client, indicating that the change is complete. The client waits for the response to come back before attempting another request.

(Sketch the system and some example interactions here.)

Now, keep in mind that these messages flow over a network. This has two effects on the users of the system:

- First, the network imposes some minimum latency (let's say 1ms) on each message. As a result, reading and writing a large amount of data from this file server is going to be very slow – much slower than accessing a local disk.
- Second, the network makes message delivery unreliable. Either a request or a response could be delayed arbitrarily, or completely dropped. In the case of a backhoe cutting through a network cable, there may be no communication at all along a particular link. If a client doesn't receive a response to a message, then it has no choice but to wait and try the request again. It could wait a very long time!

Let's evaluate this system according to the CAP criteria:

- Consistency – HIGH – Example 1 is highly consistent because every operation is applied in a known order, and all clients have an unfiltered view of the central server.
- Availability – LOW – Every single read or write requires a network operation, making this system much slower than accessing a local file system.
- Partitionability – MEDIUM – If a single client is partitioned from the file server, it cannot perform any operations. However, all other connected clients are able to continue.

(The careful reader will note that we have not yet defined precisely what constitutes higher or lower for each of these three properties. Bear with me and let that be vague for a while, and we will formalize it more below.)

Example System 2 – Write-Through Cache

Take the system from example 1, and modify it by adding a cache of finite capacity to each client node. Let's give each client some simple logic for managing these caches:

- Read – When a client attempts to read a file, it first looks in its cache to see if that file's data is already present. If so, the read is satisfied from that data. If not, then the client issues a read request to the server, waits for the response, and replaces the least recently used (LRU) item in the cache.
- Write – When a client attempts to write a file, it first issues a write command to the server, and waits for a response. If an older version of the file exists in the cache, it is updated to the new value. If not, then the client replaces the LRU item in the cache with the newly written value.

(Sketch the system and some example operations here.)

Example System 3 – Consistent Caches

The previous system had a very simple method of dealing with cached data that (obviously) resulted in some serious inconsistencies. Let's try to solve the problem in a different way, by making the system avoid inconsistencies in cached data.

Take the design from System 2, where each client has a cache of finite capacity. Now, when each client tries to read or write a file, do this:

- Read – (Same as System 2)
- Write – When a client attempts to modify a file, it first sends a message to every other client, instructing it to invalidate its cached copy of that file. Once those caches acknowledge that they have purged their copies of the file, the client sends the write request to the server, and updates its own cache.

(Sketch the system and some example operations here.)

Ok, now you evaluate this system according to the CAP criteria:

- Consistency?
- Availability?
- Partitionability?

Thought Experiment: Now you design a system that has a combination of CAP that we haven't seen so far.

Thought Experiment: Consider a modification of System 2 in which clients write data only to their local caches, and send data back to the server only when evicted from the cache. How would this affect the CAP properties?

How do we quantify CAP?

Now that we have considered several model systems, it should be clear that there is a real tradeoff between Consistency, Availability, and Partitionability. Dealing with the P is central to distributed computing: when we cannot communicate, should we optimistically try to make progress, or pessimistically wait, in order to achieve consistency? There is no single answer to this problem: different applications will require different solutions.

In the analysis above, we were a bit sloppy about stating exactly what is meant by "more" or "less" of each of the three properties.

Availability is probably the easiest to quantify. For a given system, one could measure every attempt to read or write a value, and then compute a statistic like the mean, median, or 99th percentile of latency for various operations. A system that provides a lower mean is providing "more" availability.

Partition tolerance might be measured by evaluating what set of clients and operations can continue in the presence of network outages. A system that allows all clients to perform reads during a network outage is providing "more" partition tolerance than a system that allows no clients to perform reads.

Consistency is the most complex property to describe. There exist a variety of consistency models that can be implemented by adjusting just how and when caches are updated, and whether clients can continue to operate during a partition.

Strong Consistency:

Once an update is complete, all clients will see that new value.

Causal Consistency:

If process A tells B that it has updated X, then B will see the latest value of X.

Read-Your-Writes:

If process A updates X, then A will never see an older value of X.

Monotonic Reads:

If process A reads a value from X, then it will never read back an older value.

Monotonic Writes:

All writes by process A are applied in the order they are given.

Eventual Consistency:

All updates will become visible to everyone, if you wait long enough.

Chaos:

No guarantees!

Storage Replication

Many distributed systems replicate data across a large number of storage devices. This is done to provide insurance against storage failures, but also to provide a high degree of availability for commonly used data.

But multiple copies results in a new kind of consistency problem. If I update 2 out of 5 copies, is the write "complete" or not? Systems answer this question in various ways, which are known as quorum protocols.

Vogels [2] gives a framework for thinking about this:

Suppose you have a system with N replicated storage units. To update an item, a client must write W of the replicas upfront. (The remainder will get updated eventually in the background.) To read an item, a client must read R of the replicas, in order to decide whether the most recent value has been read. (If the values differ, assume you can tell which one is the newest)

For example:

N=2, W=2, R=1 is a strongly consistent system: a writer must update both replicas, and a reader can read either one of them. (RAID 1)

N=2, W=1, R=2 is also a strongly consistent system: a writer can update either replica, and a reader must read both to obtain the latest.

N=2, W=1, R=1 is an eventually consistent system: the writer can update either replica, and the reader can read either replica, so you may not see consistent results.

And now some general observations:

(W + R) > N is strongly consistent ,while (W+R) <= N is weakly consistent.

W < (N+1)/2 means write conflicts can occur.

Monotonic read (and write) consistency is achieved by making clients "sticky" with respect to the replicas that they use.

Strong consistency models can only be achieved by delaying either reads or writes during partitions. In a large enough distributed system, partitions are omnipresent. Ergo, very large distributed systems almost always rely on weak or eventual consistency in order to achieve acceptable availability.